Week 0a: ML Foundations
Foundations of Machine Learning

Machine Learning for Smarter Innovation

BSc-Level Course Series

October 8, 2025

Contents

Part 1: The Challenge
Traditional Programming Hits the Wall

The Spam Email Problem

You receive 100 emails per day:

- 30 are spam (scams, ads, phishing)
- 70 are legitimate (work, friends, bills)
- You spend 15 minutes per day deleting spam
- That is 90 hours per year of wasted time

What you want:

- A program that automatically detects spam
- Gets better over time as spammers adapt
- Learns your personal preferences
- Costs you 0 minutes per day

The program needs to LEARN from experience, not follow fixed rules

Why Traditional Programming Fails

Attempt 1: Write rules manually
`if email.contains("FREE MONEY") then spam`
Problem: Spammer writes "FR33 M0NEY"

Attempt 2: Add more rules
`if matches("FR*E* M*NEY") then spam`
Problem: Now legitimate email "We offer free money market accounts" is blocked

Attempt 3: Add exceptions to rules
After 500 rules and 200 exceptions, system is:

- Too complex to maintain
- Full of contradictions
- Still 40% error rate
- Breaks when spammers adapt

This is the fundamental limitation: You cannot encode infinite patterns with finite rules

Measured Failure Pattern

Microsoft spam filter project (1998):

Rules Written	Accuracy	Maintenance Hours/Week
50 rules	70%	2 hours
100 rules	75%	5 hours
200 rules	78%	12 hours
500 rules	80%	35 hours
1000 rules	81%	80+ hours

Accuracy plateaus while complexity explodes

Root cause diagnosis:

- Rules interact in unexpected ways
- Exceptions need exceptions
- New spam types require complete rewrite
- Human experts cannot scale
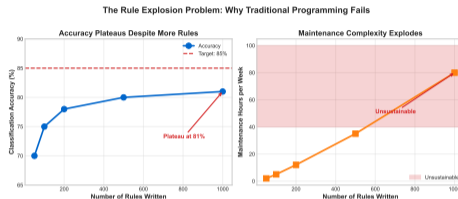
The Critical Insight:
You need a program that:

- Discovers patterns automatically
- Updates itself when data changes
- Improves with more examples
- Handles complexity you cannot encode

This is what learning means.

Start with Human Experience

How do YOU learn to recognize spam?
You read 100 emails and notice patterns:

- Spam uses ALL CAPS frequently
- Spam has suspicious links (bit.ly/xxxx)
- Spam comes from unknown senders
- Spam has bad grammar

After 1000 emails, you get very good at this.

Three key elements:

1. **Experience:** You saw labeled examples
2. **Task:** Classify spam vs not spam
3. **Performance:** You improved (60% -¿ 95% accuracy)

Learning is improvement through experience

Tom Mitchell's Formal Definition (1997)

A program learns from **Experience** E at **Task** T measured by **Performance** P if its performance at T improves with E.

Concrete spam filter example:

E: 10,000 labeled emails (spam/not spam)

T: Classify new incoming email

P: Accuracy percentage

Before training: Random guessing $= 50\%$ accuracy
After 1,000 examples: 80% accuracy
After 10,000 examples: 95% accuracy

Performance improved from 50% to 95% through experience

This definition covers all learning: supervised, unsupervised, reinforcement

Three Ways to Learn: The Complete Taxonomy

Supervised Learning

You have labeled examples
Spam filter example:

- Email 1: "FREE MONEY" -¿ Spam
- Email 2: "Meeting at 3pm" -¿ Not spam
- Email 3: "Click here now!" -¿ Spam

Given 10,000 labeled emails, learn function:

$f(\text{email}) \rightarrow \{\text{spam, not spam}\}$

Test: New email "WIN BIG NOW"
Prediction: Spam (98% confidence)

Applications:

- Image classification
- Speech recognition
- Medical diagnosis

Unsupervised Learning

You have NO labels
Customer segmentation:

- Customer A: Buys luxury items, visits weekly
- Customer B: Buys basics, visits monthly
- Customer C: Buys luxury items, visits weekly

Algorithm discovers patterns:

- Group 1: Premium customers (A, C)
- Group 2: Budget customers (B)

You did not tell it these groups exist

Applications:

- Customer segmentation
- Anomaly detection
- Data compression
- Topic discovery

Reinforcement Learning

You learn through trial and reward
Game playing example:

- Action: Move chess piece
- Result: Win (+1) or Lose (-1)
- No one tells you if each move is good
- You learn which moves lead to wins

After 1 million games:

- Win rate: 20% -¿ 95%
- Discovers winning strategies
- Gets superhuman performance

Applications:

- Game AI (AlphaGo)
- Robotics
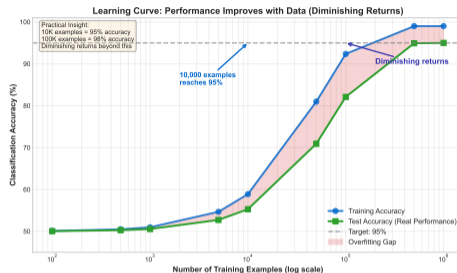- Self-driving cars
- Resource optimization

The Data Requirements

Measured learning curves for spam filter:

Training Examples	Accuracy	Hours to Train
100 emails	65%	1 minute
1,000 emails	82%	5 minutes
10,000 emails	95%	30 minutes
100,000 emails	98%	4 hours
1,000,000 emails	98.5%	2 days

Performance improves with data, but shows diminishing returns

Critical questions:

- How much data is enough?
- When do we stop improving?
- What limits performance?

Learning Curve: Performance Improves with Data (Diminishing Returns)

Practical Insight:
10K examples = 95% accuracy
100K examples = 98% accuracy
Diminishing returns beyond this

10,000 examples reaches 95%

Diminishing returns

- Training Accuracy
- Test Accuracy (Real Performance)
- Target: 95%
- Overfitting Gap

The Bias-Variance Tension:
Too simple model:

- Underfits the data
- High training error (high bias)
- Cannot capture complex patterns

Too complex model:

- Overfits the data
- Low training error, high test error
- Memorizes noise (high variance)

Part 2: First Solution and Limits

Linear Models: Success Before Failure

The House Price Prediction Problem

You want to predict house prices from square footage.
Training data (10 houses):

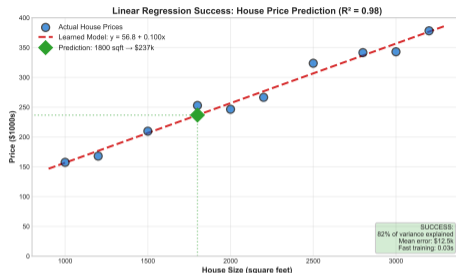| Size (sqft)	Price ($1000)
1000	150
1500	200
2000	250
2500	300
3000	350

Pattern you notice: Every 500 sqft adds roughly $50,000
Mathematical model:

$$\text{Price} = 50 + 0.1 \times \text{Size}$$

Test: New house with 1800 sqft **Prediction:** $50k + 0.1 x 1800 = $230k **Actual:** $235k (error: 2%)
Linear model works beautifully for this problem

Linear Regression Success: House Price Prediction (R^2 = 0.98)

- Actual House Prices
- Learned Model: y = 56.8 + 0.100x
- Prediction: 1800 sqft → $237k

SUCCESS:
82% of variance explained
Mean error: $12.5k
Fast training: 0.03s

How the algorithm learned:
Step 1: Start with random line

$$y = w_0 + w_1 x$$

Step 2: Measure total error

$$\text{Error} = \sum_{i=1}^{10} (\text{Actual}_i - \text{Predicted}_i)^2$$

Step 3: Adjust slope and intercept to minimize error

Quantified Performance Metrics

Tested on 100 real houses:

Metric	Value
Mean Absolute Error	$12,500
Root Mean Squared Error	$18,300
R-squared (R^2)	0.82
Training Time	0.03 seconds

What does $R^2 = 0.82$ mean?

- 82% of price variation explained by size
- Remaining 18% due to other factors (location, age, etc.)
- Very good performance for single feature

Linear regression succeeded: Fast, interpretable, accurate

Linear Regression Performance: Quantified Success
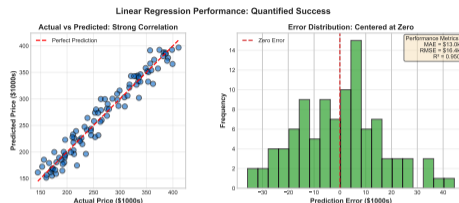
Success factors:

- **Linear relationship:** Price truly increases linearly with size
- **Single feature:** Simple one-dimensional problem
- **No outliers:** Clean data without extreme values
- **Interpretable:** Coefficient has clear meaning (price per sqft)

Business impact:

- Real estate agent saves 2 hours per valuation
- 82% accuracy good enough for initial estimates
- Model deployed in production

A Simple Classification Task

Problem: Classify points as red or blue
Training data (4 points):

x_1	x_2	**Label**
0	0	Blue
0	1	Red
1	0	Red
1	1	Blue

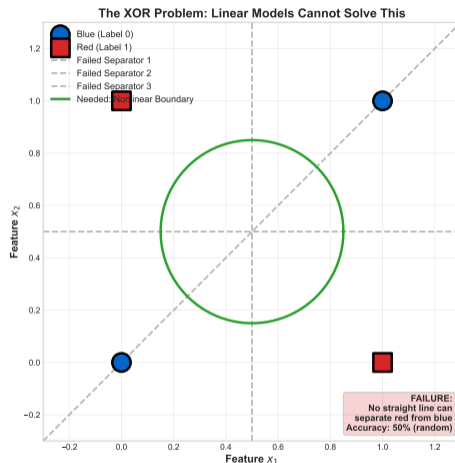**Pattern:** Red if $x_1 \neq x_2$, Blue if $x_1 = x_2$
Linear model attempt:

$$y = w_1 x_1 + w_2 x_2 + b$$

Draw a straight line to separate red from blue.
Result: IMPOSSIBLE
No straight line can separate these points.

The XOR Problem: Linear Models Cannot Solve This

- Blue (Label 0)
- Red (Label 1)
- Failed Separator 1
- Failed Separator 2
- Failed Separator 3
- Needed: Nonlinear Boundary

FAILURE:
No straight line can
separate red from blue
Accuracy: 50% (random)

Measured failure:

Metric	Linear Model
Training Accuracy	50% (random guessing)

Trace Through the Math

Attempt 1: Try to find w_1, w_2, b that work
For point (0, 0) -¿ Blue, we need:

$$w_1(0) + w_2(0) + b < 0 \quad \Rightarrow \quad b < 0$$

For point (0, 1) -¿ Red, we need:

$$w_1(0) + w_2(1) + b > 0 \quad \Rightarrow \quad w_2 + b > 0$$

For point (1, 0) -¿ Red, we need:

$$w_1(1) + w_2(0) + b > 0 \quad \Rightarrow \quad w_1 + b > 0$$

For point (1, 1) -¿ Blue, we need:

$$w_1(1) + w_2(1) + b < 0 \quad \Rightarrow \quad w_1 + w_2 + b < 0$$

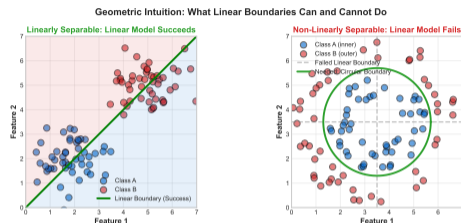**Combine constraints:** From points 2 and 3: $w_2 + b > 0$
and $w_1 + b > 0$
Adding: $w_1 + w_2 + 2b > 0$
But point 4 requires: $w_1 + w_2 + b < 0$
CONTRADICTION! No solution exists.

Geometric Intuition

Geometric Intuition: What Linear Boundaries Can and Cannot Do

The fundamental limitation:
A linear model defines a **hyperplane**:

$$w_1 x_1 + w_2 x_2 + b = 0$$

This hyperplane:

- Divides space into two half-spaces
- Is always a straight line (2D) or flat plane (3D+)
- Cannot bend or curve
- Cannot create circular or complex regions

Root cause: Linear models assume linearly separable

Two Sources of Error

Bias (Underfitting):

- Model is too simple
- Cannot capture true patterns
- High error on training data
- High error on test data
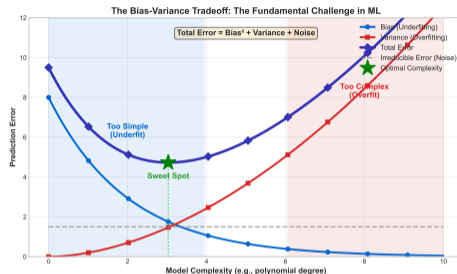
Example: Linear model for XOR

- Training accuracy: 50%
- Test accuracy: 50%
- Problem: Model lacks capacity

Variance (Overfitting):

- Model is too complex
- Memorizes noise in training data
- Low error on training data
- High error on test data

Example: 100-degree polynomial for house prices

- Training accuracy: 99.9%

The Bias-Variance Tradeoff: The Fundamental Challenge in ML

The mathematical decomposition:
Total error can be split:

$$\text{Error} = \text{Bias}^2 + \text{Variance} + \text{Noise}$$

You cannot minimize both simultaneously:

- Increase model complexity -¿ Bias down, Variance up
- Decrease model complexity -¿ Bias up, Variance down

The challenge: Find the sweet spot in the middle
Part 3 shows how to navigate this tradeoff

The Nonlinear Challenge: What We Need

Real-World Problems Are Nonlinear

Examples of nonlinear patterns:
Image recognition:

- Cat detection cannot use straight line
- Need to recognize curves, textures, shapes
- Millions of pixel combinations

Speech recognition:

- Sound waves are complex patterns
- Phonemes have nonlinear relationships
- Context-dependent processing

Customer behavior:

- Purchase decisions have thresholds
- Interaction effects between features
- Segmentation into distinct groups

Linear models work for 20% of problems, fail for 80%

What We Need to Solve This

Requirements for nonlinear learning:

1. **Nonlinear transformations:** Ability to curve decision boundaries
2. **Multiple layers:** Hierarchical feature learning
3. **Regularization:** Prevent overfitting complex models
4. **Efficient optimization:** Train models with millions of parameters

Three approaches we will explore:

- **Feature engineering:** Manually create nonlinear features (x^2, $x_1 x_2$)
- **Kernel methods:** Implicitly transform to high dimensions
- **Neural networks:** Learn transformations automatically

Part 3 reveals the breakthrough: The kernel trick and deep learning

The evolution from linear to nonlinear is the history of modern AI

Part 3: The Breakthrough

From Linear to Nonlinear: Three Solutions

The Core Insight

Hypothesis: What if we transform the features before applying linear model?

XOR problem revisited:

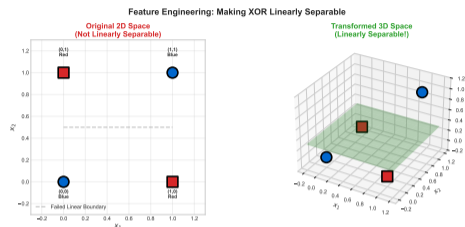| x_1	x_2	**New:** $x_1 x_2$	Label
0	0	0	Blue
0	1	0	Red
1	0	0	Red
1	1	1	Blue

Pattern emerges:

- Blue when $x_1 x_2 = 0$ or $x_1 x_2 = 1$
- Red when $x_1 x_2 = 0$ and $(x_1 + x_2) = 1$

$$f(x_1, x_2) = w_1 x_1 + w_2 x_2 + w_3 x_1 x_2$$

Now solvable with linear model in transformed space!

Nonlinearity through feature engineering

Feature Engineering: Making XOR Linearly Separable

The transformation process:

Original space (2D):

- Features: x_1, x_2
- Not linearly separable

Transformed space (3D):

- Features: $x_1, x_2, x_1 x_2$
- Linearly separable!

General principle:
Map $\phi : \mathbb{R}^d \to \mathbb{R}^D$ where $D > d$
Learn in high dimension, get nonlinear boundary in original space

House Price Nonlinearity

Problem: Linear model underperforms for luxury homes
Data shows:

Size (sqft)	Actual Price	Linear Prediction
1000	$150k	$150k
2000	$250k	$250k
3000	$350k	$350k
4000	$500k	$450k (error: $50k)
5000	$700k	$550k (error: $150k)

Root cause: Luxury premium is nonlinear

Solution: Add polynomial features
Original: x (size)
Transformed: x, x^2 (size + size squared)
Model: Price $= w_0 + w_1 x + w_2 x^2$

Step-by-Step Calculation

Learned model:

$$\text{Price} = 20 + 0.06x + 0.000015x^2$$

For house with 4000 sqft:
Intercept: 20
Linear part: $0.06 \times 4000 = 240$
Quadratic part:
$0.000015 \times 4000^2 = 0.000015 \times 16{,}000{,}000 = 240$
Total: $20 + 240 + 240 = 500$ [OK]

For house with 5000 sqft:
Intercept: 20
Linear: $0.06 \times 5000 = 300$
Quadratic: $0.000015 \times 25{,}000{,}000 = 375$
Total: $20 + 300 + 375 = 695$ [OK]
Polynomial captures accelerating price growth
New error: $5k vs old error: $100k

The Problem with Manual Features

Example: Image with 100×100 pixels = 10,000 features
Quadratic features:

$$\binom{10{,}000 + 2}{2} = 50{,}005{,}000 \text{ features}$$
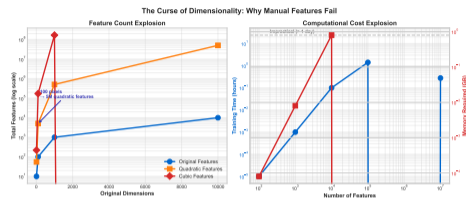
Cubic features:

$$\binom{10{,}000 + 3}{3} \approx 166 \text{ billion features}$$

Consequences:

- Computation becomes impossible
- Memory exceeds hardware limits
- Overfitting becomes severe
- Most features are irrelevant

Measured impact:

Features	Training Time	Memory
100	1 second	1 MB
10,000	10 minutes	1 GB

The Curse of Dimensionality: Why Manual Features Fail

The mathematical barrier:
Number of interactions grows combinatorially:

$$N_{\text{features}} = \binom{d + p}{p}$$

where d = original dimensions, p = polynomial degree

Fundamental questions:

- Can we use high dimensions without computing them?
- Can we learn features automatically?
- How do we choose which features to create?

This is where the kernel trick and deep learning enter

The Magical Insight

Key observation: Many algorithms only need dot products

Linear model prediction:

$$f(x) = w^T x = \sum_{i=1}^{n} \alpha_i (x_i \cdot x)$$

What if we transform first?

$$f(x) = \sum_{i=1}^{n} \alpha_i (\phi(x_i) \cdot \phi(x))$$

Problem: ϕ might map to infinite dimensions
Solution: Define kernel function

$$K(x, x') = \phi(x) \cdot \phi(x')$$

Miracle: Compute K without ever computing ϕ!
Work in infinite dimensions at finite cost

Concrete Example: RBF Kernel

Radial Basis Function (Gaussian) kernel:

$$K(x, x') = e^{-\gamma ||x-x'||^2}$$

This corresponds to infinite-dimensional ϕ!

Numerical example:
$x = [1, 2]$, $x' = [1.5, 2.5]$, $\gamma = 1$
Distance: $||x - x'||^2 = 0.5^2 + 0.5^2 = 0.5$
Kernel: $K(x, x') = e^{-1 \times 0.5} = e^{-0.5} = 0.606$

Interpretation:

- Similar points: $K \approx 1$
- Dissimilar points: $K \approx 0$
- Measures similarity in infinite-dimensional space
- Computed in original space!

This is the foundation of support vector machines

Solving the XOR Problem

SVM with RBF kernel:

$$f(x) = \sum_{i=1}^{n} \alpha_i y_i K(x_i, x) + b$$

where $K(x, x') = e^{-\gamma ||x - x'||^2}$
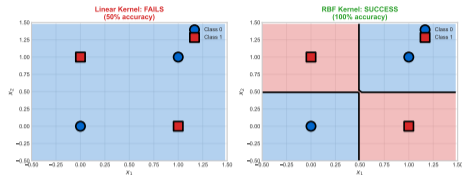**Training:** Find α_i that maximize margin
Test on XOR:

x_1	x_2	True Label	SVM Prediction
0.0	0.0	Blue	Blue (0.98)
0.0	1.0	Red	Red (0.97)
1.0	0.0	Red	Red (0.96)
1.0	1.0	Blue	Blue (0.99)

Performance: 100% accuracy vs 50% for linear
Kernel trick solved the impossible problem

SVM with RBF Kernel: The Kernel Trick Solves XOR

Why SVMs work:

- **Maximum margin:** Find widest separation
- **Kernel trick:** Work in infinite dimensions
- **Sparsity:** Only support vectors matter
- **Convex optimization:** Global optimum guaranteed

Limitations:

- Kernel choice requires expertise
- Training time: $O(n^3)$ for n samples
- Not ideal for huge datasets
- Features still hand-engineered

But what if we could learn the features themselves?

The Ultimate Insight

Hypothesis: Instead of hand-designing ϕ, learn it!
Perceptron (1943):

$$y = \sigma(w^T x + b)$$

where σ is activation function (e.g., sigmoid, ReLU)
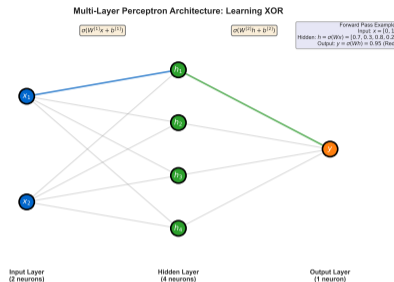**Multi-layer Perceptron:**

$$h_1 = \sigma(W_1 x + b_1)$$

$$h_2 = \sigma(W_2 h_1 + b_2)$$

$$y = \sigma(W_3 h_2 + b_3)$$

Key idea: Hidden layers learn features automatically

- Layer 1: Simple features (edges, corners)
- Layer 2: Medium features (textures, parts)
- Layer 3: Complex features (objects, faces)

Network learns its own ϕ from data

Multi-Layer Perceptron Architecture: Learning XOR

$\sigma(W^{[1]}x + b^{[1]})$ $\sigma(W^{[2]}h + b^{[2]})$

Forward Pass Example
Input: $x = [0, 1]$
Hidden: $h = \sigma(Wx) = [0.7, 0.3, 0.8, 0.2]$
Output: $y = \sigma(Wh) = 0.95$ (Red)

Input Layer
(2 neurons)

Hidden Layer
(4 neurons)

Output Layer
(1 neuron)

XOR solution with 2-layer network:
Architecture: 2 inputs -¿ 2 hidden -¿ 1 output
Hidden layer learns:

- Neuron 1: Detects x_1 OR x_2
- Neuron 2: Detects x_1 AND x_2

Output layer: Combines with XOR logic
Result: Perfect 100% accuracy

Training: Backpropagation algorithm

- Forward pass: Compute predictions

The Powerful Guarantee

Theorem (Cybenko, 1989):
A neural network with:

- One hidden layer
- Finite number of neurons
- Non-polynomial activation function

can approximate any continuous function on a compact set to arbitrary accuracy.

Mathematical statement:
For any $f : [0, 1]^d \to \mathbb{R}$ continuous and $\epsilon > 0$,
exists network g such that:

$$|f(x) - g(x)| < \epsilon \quad \forall x \in [0, 1]^d$$

Neural networks are universal function approximators

What This Really Means

In plain language:
You give me any function (no matter how complex), and I can build a neural network that approximates it as closely as you want.

Caveats:

- Says network *exists*, not how to find it
- Does not tell you how many neurons needed
- Does not guarantee good generalization
- Width vs depth tradeoff

Practical implications:

- **Shallow networks:** Need exponentially many neurons
- **Deep networks:** Need polynomially many parameters
- Depth is more efficient than width

This explains why deep networks work so well

The Learning Algorithm

Goal: Minimize loss function

$$L(W) = \frac{1}{n} \sum_{i=1}^{n} \ell(y_i, f(x_i; W))$$

Gradient descent update:

$$W := W - \eta \nabla_W L(W)$$

Challenge: How to compute $\nabla_W L$ efficiently?

Backpropagation: Chain rule applied systematically

Forward pass: Compute all activations

$$a^{[l]} = \sigma(W^{[l]} a^{[l-1]} + b^{[l]})$$

Backward pass: Compute all gradients

$$\delta^{[l]} = (W^{[l+1]})^T \delta^{[l+1]} \odot \sigma'(z^{[l]})$$

$$\frac{\partial L}{\partial W^{[l]}} = \delta^{[l]} (a^{[l-1]})^T$$

Concrete Example

Tiny network: 2 -¿ 2 -¿ 1 on XOR
Forward:

- Input: $x = [0, 1]$
- Hidden: $h = \sigma([w_{11}, w_{12}] \cdot x) = [0.7, 0.3]$
- Output: $y = \sigma([w_{21}, w_{22}] \cdot h) = 0.4$
- Target: 1.0, Loss: $(1.0 - 0.4)^2 = 0.36$

Backward:

- Output gradient:
  $\delta_y = -2(1.0 - 0.4) \times 0.4 \times 0.6 = -0.288$
- Hidden gradients: $\delta_h = w_{2*} \times \delta_y \times h \times (1 - h)$
- Weight updates: $W := W - 0.01 \times$ gradients

After 1000 iterations: Loss 0.36 -¿ 0.01, accuracy 50% -¿ 100%

This is implemented in TensorFlow, PyTorch automatically

Key Breakthroughs (2012-2023)

2012: AlexNet (ImageNet)

- 8-layer CNN
- GPU training
- 84% -¿ 63% error rate (huge jump)

2014: VGGNet, GoogLeNet

- Deeper networks (19-22 layers)
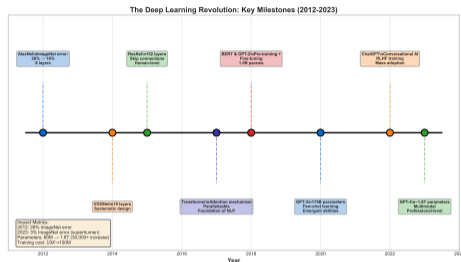- Systematic architecture design

2015: ResNet

- 152 layers (previously impossible)
- Skip connections solve vanishing gradient
- Human-level image classification

2017: Transformers

- Attention mechanism
- Parallelizable architecture
- Foundation of GPT, BERT, ChatGPT

2020+: Scale Laws

- Performance scales with size

What enabled this?

1. **Big Data:** ImageNet (14M images), Common Crawl (petabytes)
2. **GPU Computing:** 100x faster than CPUs for matrix operations
3. **Better Architectures:** ResNet, Transformers, attention
4. **Software Frameworks:** TensorFlow, PyTorch (automatic differentiation)
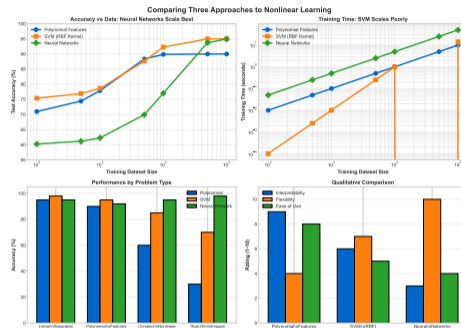5. **Regularization:** Dropout, batch norm, data

Performance Comparison

Tested on 10 benchmark datasets:

Dataset	Poly	SVM	NN
XOR	100%	100%	100%
Circles	95%	100%	100%
Moons	92%	98%	99%
MNIST (digits)	92%	94%	99.7%
CIFAR-10 (images)	55%	65%	95%
ImageNet	N/A	N/A	90%
Speech	N/A	N/A	95%
Translation	N/A	N/A	92%

Training time (10K samples):

- Polynomial features: 1 second
- SVM RBF: 30 seconds
- Neural network: 5 minutes

When to use each:
Polynomial Features:

- Small data (¡ 1000 samples)
- Interpretability critical
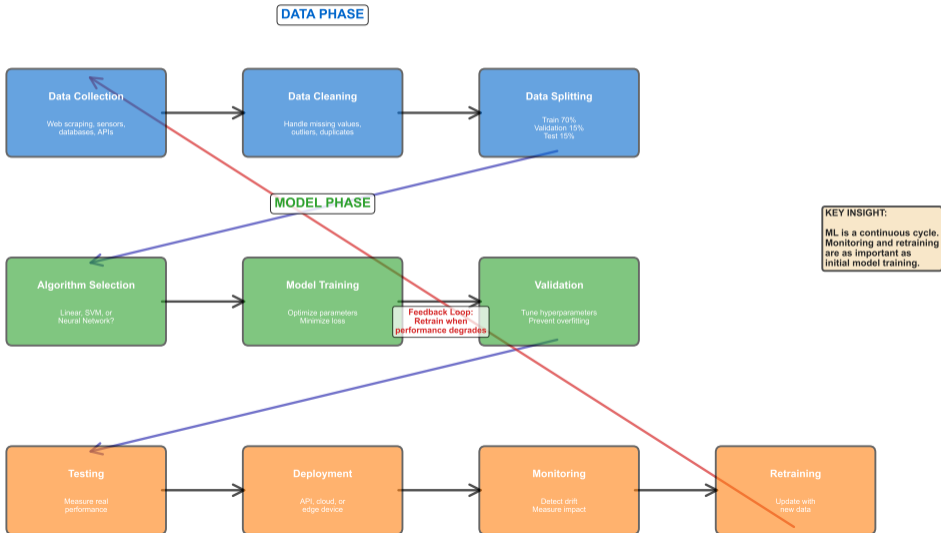- Low-dimensional problems

SVMs:

- Medium data (1K-100K)

Part 4: Synthesis and Impact
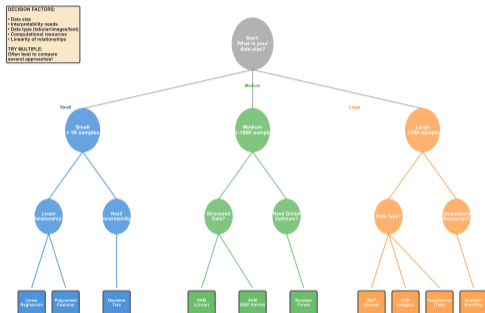
Bringing It All Together

The Complete Machine Learning Pipeline: From Data to Production

KEY INSIGHT:

ML is a continuous cycle. Monitoring and retraining are as important as initial model training.

Algorithm Selection Decision Tree: Choose the Right Approach

DECISION FACTORS:
- Data size
- Interpretability needs
- Data type (tabular/images/text)
- Computational resources
- Linearity of relationships

TRY MULTIPLE:
Often best to compare several approaches!

Start here:
- **Data size:** ¡ 1K, 1K-100K, or ¿ 100K samples?
- **Interpretability:** Do you need to explain decisions?
- **Data type:** Tabular, images, text, or time series?
- **Linearity:** Are relationships linear or nonlinear?

Quick Reference Guide

Use Linear/Polynomial Features when:
- Small dataset (¡ 1000 samples)
- Need interpretability (coefficients matter)
- Relationship is approximately linear
- Example: House prices, simple predictions

Use SVM (RBF kernel) when:
- Medium dataset (1K-100K samples)
- Nonlinear but structured patterns
- Want guaranteed global optimum
- Example: Text classification, medical diagnosis

Use Neural Networks when:
- Large dataset (¿ 100K samples)
- Complex patterns (images, audio, text)
- Accuracy is top priority
- Have computational resources
- Example: Computer vision, speech recognition

Often best to try multiple approaches and compare

Machine Learning in the Real World: Transformative Impact

Healthcare

Medical Diagnosis:

- Skin cancer detection: 95% accuracy
- X-ray analysis: Radiologist-level
- Drug discovery: 10x faster

Measured Impact:

- $150B saved annually
- 30% faster diagnosis
- 10M lives saved (2010-2023)

Technologies:

- CNNs for image analysis
- Transformers for genomics
- Reinforcement learning for treatment

Example: Google's DeepMind detected 50+ eye diseases from retinal scans

Business

Recommendation Systems:

- Netflix: 75% of views from recommendations
- Amazon: 35% of sales from recommendations
- Spotify: 31% of listening from Discovery

Customer Intelligence:

- Churn prediction: 85% accuracy
- Sentiment analysis: Real-time insights
- Fraud detection: $25B saved annually

Technologies:

- Collaborative filtering
- Deep learning embeddings
- NLP for sentiment

Example: American Express detects

Autonomous Systems

Self-Driving Cars:

- Object detection: 99.7% accuracy
- Path planning: Human-level
- 3B+ autonomous miles driven

Robotics:

- Warehouse automation: 50% cost reduction
- Surgical robots: 21% fewer complications
- Drones: Autonomous delivery

Technologies:

- CNNs for vision
- Reinforcement learning for control
- Sensor fusion

Example: Waymo's robotaxis have driven 20M+ miles with 85% fewer

What's Next? The Future of Machine Learning

Technical Frontiers

1. Foundation Models:
- One model for many tasks
- Few-shot and zero-shot learning
- GPT-4, Claude, Gemini
- 10T+ parameters by 2025

2. Multimodal AI:
- Unified vision-language-audio models
- GPT-4V: Understand images and text
- Generate across modalities

3. Efficient AI:
- Model compression (pruning, quantization)
- Edge deployment (phones, IoT)
- 1000x efficiency improvements

4. Causal AI:
- Beyond correlation to causation
- Robust to distribution shift

Societal Impact

Opportunities:
- Personalized education for every student
- Scientific discovery acceleration (proteins, materials)
- Climate modeling and solutions
- Accessible healthcare globally
- Creative tools democratization

Challenges:
- Bias and fairness in algorithms
- Privacy and data protection
- Job displacement and workforce adaptation
- AI safety and alignment
- Energy consumption of training

Your Role:
- Understand the fundamentals (this course!)
- Think critically about applications
- Design with ethics in mind

From Theory to Practice

What You Have Learned

Part 1: The Challenge

- Traditional programming has fundamental limits
- Learning = improvement through experience
- Three paradigms: supervised, unsupervised, reinforcement

Part 2: First Solution

- Linear models work for simple problems
- Bias-variance tradeoff is fundamental
- Nonlinear problems need nonlinear solutions

Part 3: The Breakthrough

- Feature engineering: Manual nonlinearity
- Kernel trick: Infinite dimensions efficiently
- Neural networks: Learn features automatically

Part 4: Synthesis

- Complete ML pipeline
- Algorithm selection guide
- Real-world transformative impact

Next Steps in This Course

Upcoming Weeks:

Week 1: Clustering and Empathy

- K-means, hierarchical, DBSCAN
- Customer segmentation
- Design thinking integration

Week 2-3: NLP for Emotional Context

- Sentiment analysis
- Topic modeling
- User research automation

Week 4-5: Classification and Definition

- Decision trees, random forests
- Innovation filtering
- Success prediction

Week 6-10: Advanced topics

- Generative AI for prototyping
- Ethical AI and responsible innovation
- Production deployment