

# Supervised Learning

## Essential Guide to Prediction & Classification

Machine Learning for Smarter Innovation

BSc Innovation & Design Thinking

# What is Supervised Learning?

## Definition

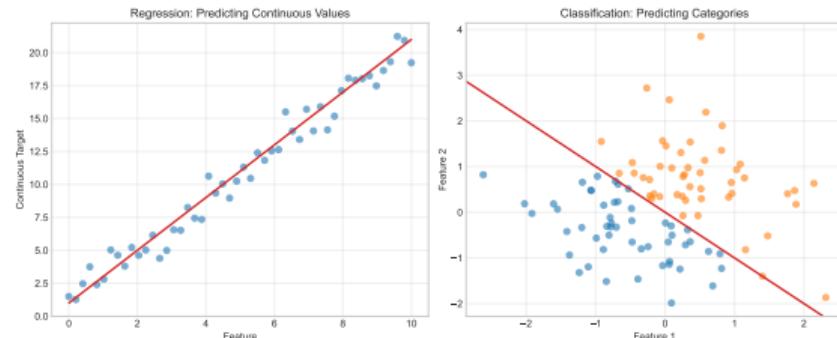
- Learning from **labeled examples**
- Input features  $X \rightarrow$  Output labels  $y$
- Algorithm learns mapping:  $f(X) \approx y$

## Two Main Tasks

- ① **Regression:** Predict continuous values
  - House prices, sales forecasts, temperatures
- ② **Classification:** Predict discrete categories
  - Spam detection, medical diagnosis, fraud

## Key Insight

Model learns patterns from historical data to predict future outcomes



## Real-World Examples

- Real estate pricing (regression)
- Email spam filtering (classification)
- Customer churn prediction (classification)
- Sales forecasting (regression)

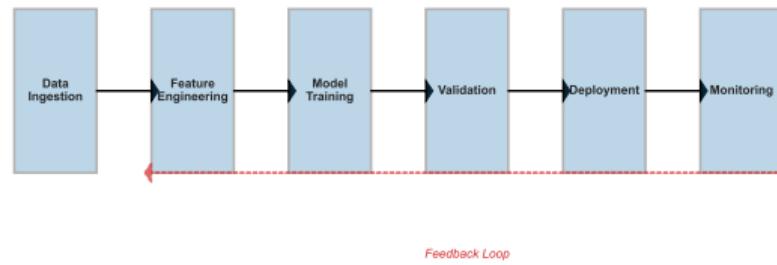
Supervised learning transforms labeled historical data into predictive models - the algorithm discovers patterns from examples rather than explicit programming

# The Supervised Learning Pipeline

Production ML Pipeline: End-to-End System

## 1. Training Phase

- Collect labeled data:  $(X_1, y_1), (X_2, y_2), \dots, (X_n, y_n)$
- Split: 70-80% training, 20-30% testing
- Algorithm learns pattern:  $\hat{f}(X)$
- Minimize error on training set



## 2. Prediction Phase

- Receive new unlabeled input:  $X_{new}$
- Apply learned model:  $\hat{y} = \hat{f}(X_{new})$
- Generate prediction

## 3. Evaluation Phase

- Test on held-out data
- Measure: Accuracy, RMSE, F1-score
- Validate generalization

### Critical Rule

**NEVER test on training data!**

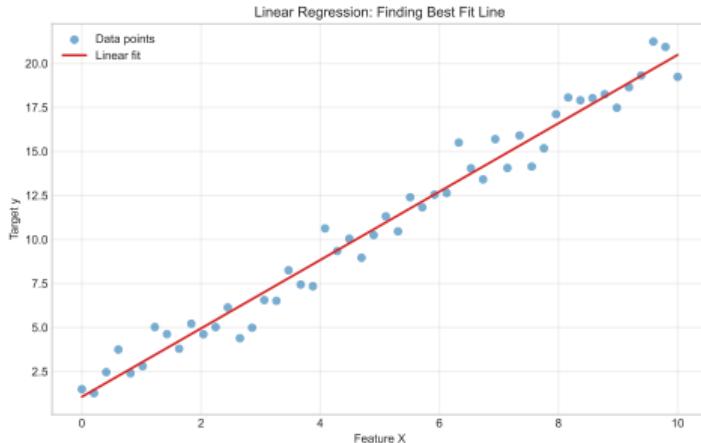
### Why?

- Model has already seen training data
- Cannot measure true generalization
- Leads to overestimation of performance

Train-test split prevents overfitting evaluation - testing on unseen data reveals true generalization performance rather than mere memorization

## Ordinary Least Squares (OLS)

- Model:  $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \epsilon$
- Goal: Minimize squared errors
- Solution:  $\hat{\beta} = (X^T X)^{-1} X^T y$
- Fast, interpretable



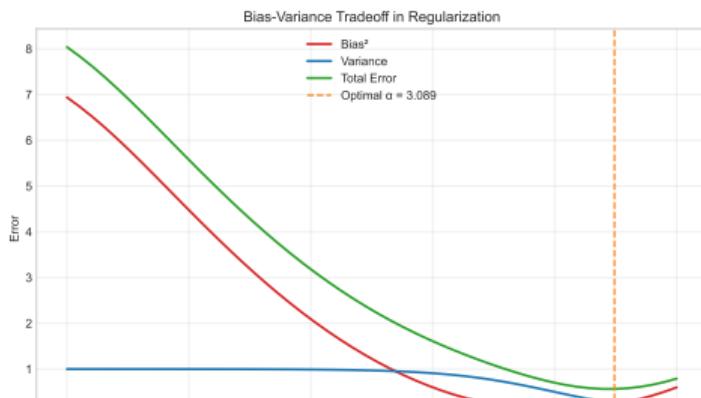
## Regularization Methods

### Ridge (L2):

- Penalizes large coefficients:  $\lambda \|\beta\|_2^2$
- Shrinks all coefficients smoothly
- Prevents overfitting

### Lasso (L1):

- Sparse penalty:  $\lambda \|\beta\|_1$



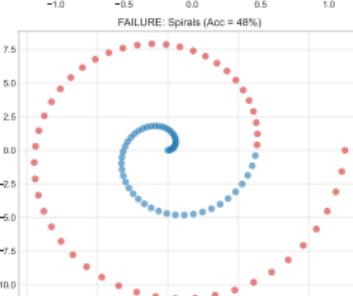
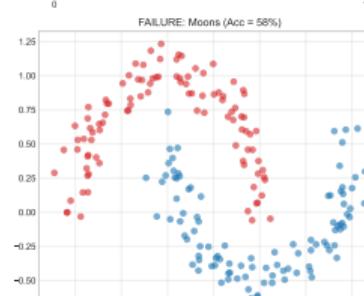
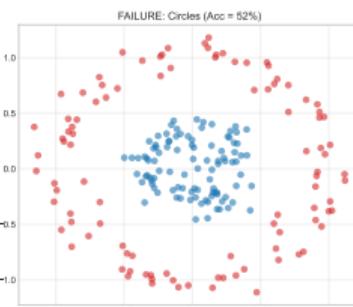
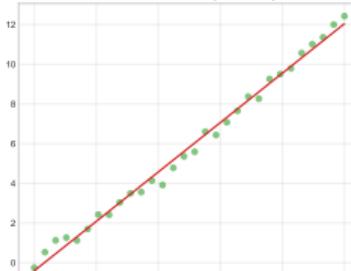
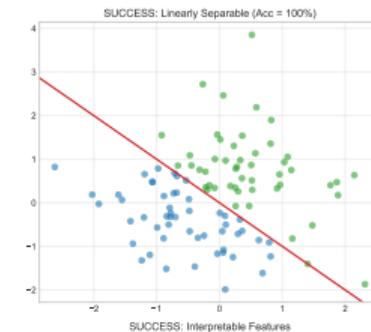
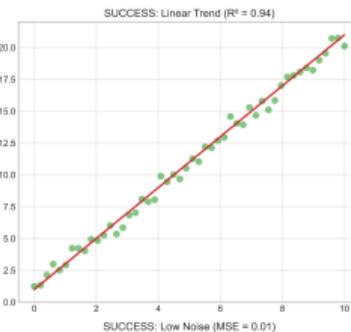
# When Linear Methods Work (and When They Don't)

## Success Cases

- **Simple relationships:** Monotonic, few features
- **Interpretability critical:** Regulatory requirements
- **Fast prediction needed:** Real-time systems
- **Limited data:** Few training examples

## Failure Cases

- **Complex patterns:** Nonlinear relationships
- **Feature interactions:** XOR-like problems
- **High-dimensional:** Many correlated features
- **Mixed data types:** Categorical + continuous



## How It Works

- Recursive binary splits
- Like playing “20 questions”
- At each node: Choose best feature + threshold
- Leaves contain predictions

## Splitting Criterion

Regression (minimize variance):

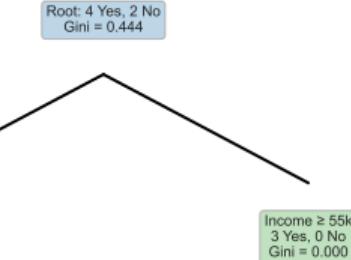
$$\text{Gain} = \text{Var}(\text{parent}) - \sum \text{Var}(\text{children})$$

Classification (Gini or Entropy):

$$\text{Gini} = 1 - \sum p_i^2$$

## Advantages

- Highly interpretable (extract rules)
- Handles mixed data types



Information Gain  
= 0.222

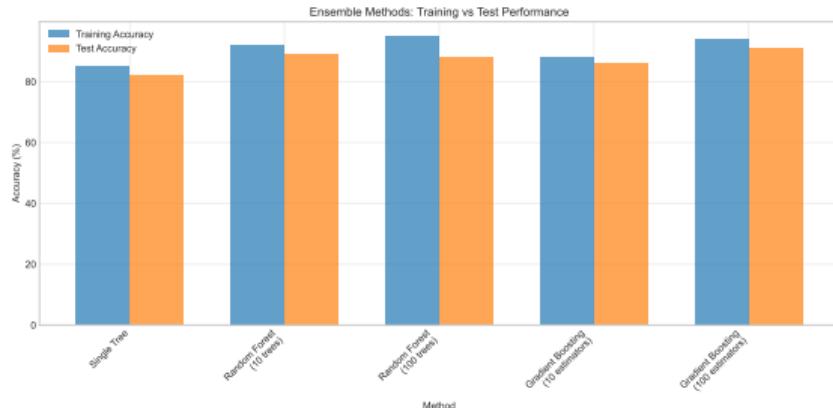
## Disadvantages

- Prone to overfitting (memorizes noise)
- High variance (unstable to data changes)

# Ensemble Methods: Power Through Combination

## Random Forest

- Build  $T$  trees (e.g., 100-500)
- Each tree sees:
  - Random subset of data (bootstrap)
  - Random subset of features
- Prediction: Average (regression) or vote (classification)
- **Effect:** Reduces variance



## Gradient Boosting

- Build trees sequentially
- Each tree corrects errors of previous
- Final: Weighted sum of all trees
- **Effect:** Reduces bias

## Modern Implementations

- XGBoost, LightGBM, CatBoost
- State-of-art for tabular data

**Trade-offs**  
**Advantages:**

- Best accuracy for tabular data
- Robust to overfitting
- Handles mixed data types
- Minimal feature engineering

**Disadvantages:**

- Less interpretable (black box)
- Slower training and prediction

# Algorithm Selection Guide

## Decision Flowchart

### 1. Is interpretability critical?

- YES → Linear methods or single tree
- NO → Consider accuracy needs

### 2. Is the relationship linear?

- YES → OLS, Ridge, Lasso, Logistic
- NO → Nonlinear methods needed

### 3. Do you need highest accuracy?

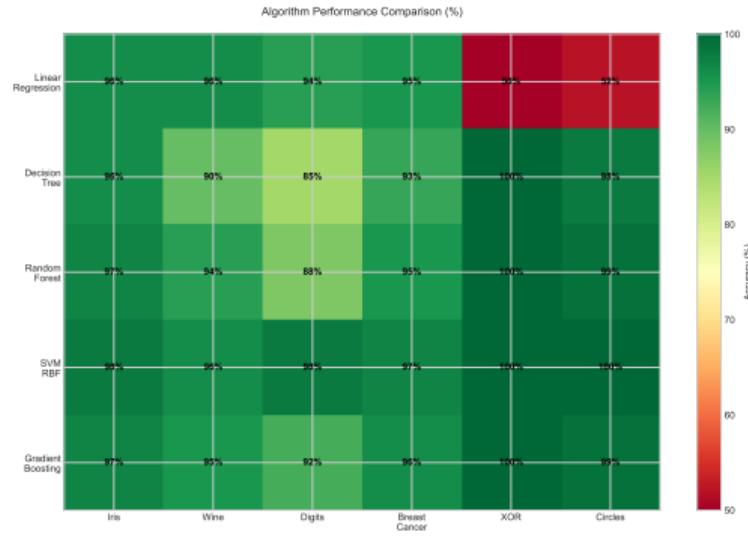
- YES → Random Forest, XGBoost
- NO → Single tree or linear

### 4. How much data do you have?

- Small (<1000) → Linear, regularization
- Medium (1000-100k) → Trees, forests
- Large (>100k) → Gradient boosting, deep learning

### 5. What are your feature types?

- All numeric → Any method
- Mixed (numeric + categorical) → Trees preferred



## General Strategy

- ① Start simple: Linear baseline
- ② Add complexity: If performance insufficient
- ③ Monitor trade-off: Interpretability vs accuracy
- ④ Cross-validate: Always test generalization

# Common Pitfalls to Avoid

## 1. Overfitting

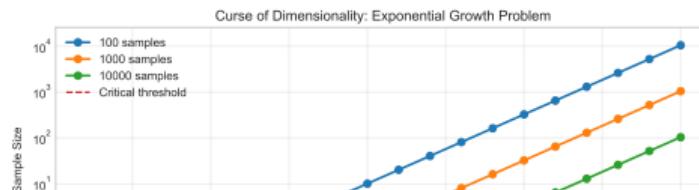
- Model memorizes training data
- High training accuracy, poor test accuracy
- **Fix:** Regularization, cross-validation, more data

## 2. Underfitting

- Model too simple for patterns
- Poor performance on both train and test
- **Fix:** Add complexity, more features, nonlinear methods

## 3. Curse of Dimensionality

- Too many features vs samples
- Distance metrics become meaningless
- **Fix:** Feature selection, dimensionality reduction, regularization



## 4. Data Leakage

- Test information leaks into training
- Examples:
  - Using future data to predict past
  - Including target in features
  - Normalizing before splitting
- **Fix:** Strict train-test separation, careful preprocessing

## 5. Imbalanced Classes

- Rare positive class (e.g., 1% fraud)
- Model predicts majority class always
- **Fix:** Resampling, cost-sensitive learning, different metrics (F1, AUC)

## 6. Not Testing Generalization

- Evaluating only on training data
- **Fix:** Always use held-out test set or cross-validation

# Best Practices & Summary

## Essential Best Practices

### 1. Data Splitting

- Always split train (70-80%) vs test (20-30%)
- Use cross-validation for hyperparameter tuning
- Never touch test set until final evaluation

### 2. Start Simple

- Begin with linear baseline
- Understand performance ceiling
- Add complexity incrementally

### 3. Feature Engineering

- Domain knowledge matters
- Handle missing values
- Encode categorical variables
- Scale/normalize features

### 4. Model Validation

- Monitor train vs test performance
- Use appropriate metrics (RMSE, accuracy, F1, AUC)
- Check predictions make sense

## Algorithm Summary

| Method                     | Strengths                    | Weaknesses            |
|----------------------------|------------------------------|-----------------------|
| Linear (OLS, Ridge, Lasso) | Fast, interpretable, stable  | Assumes linearity     |
| Logistic Regression        | Probabilistic, interpretable | Linear boundaries     |
| Decision Tree              | Interpretable, non-linear    | High variance         |
| Random Forest              | Accurate, robust             | Less interpretable    |
| Gradient Boosting          | Highest accuracy             | Slow, many parameters |

## Key Takeaways

- ① Supervised learning requires **labeled data**
- ② Two tasks: **regression** (continuous) vs **classification** (discrete)
- ③ Always use **train-test split**
- ④ Start **simple**, add complexity only if needed
- ⑤ Monitor **interpretability-accuracy** tradeoff
- ⑥ **Ensembles** achieve best accuracy for tabular data