

# Week 0c: Unsupervised Learning

## "The Discovery Challenge"

### Finding Hidden Patterns Without Labels

Machine Learning for Smarter Innovation

BSc Course Series

October 6, 2025

## Act 1: The Challenge

- Customer segmentation without labels
- Mathematical similarity definitions
- No ground truth validation
- Cluster number selection
- Quantitative evaluation metrics

## Act 3: Density & Hierarchy

- Human clustering intuition
- DBSCAN: Finding neighborhoods
- Hierarchical: Building trees
- Handling arbitrary shapes
- Modern implementations

## Act 2: K-means Algorithm

- Nearest center assignment
- Worked coordinate examples
- Success: Spherical clusters
- Failure: Non-convex shapes
- Diagnostic insights

## Act 4: Synthesis

- Method taxonomy
- Algorithm selection guide
- Modern applications
- Neural network preview

24 slides — Pattern discovery without supervision — Real-world clustering challenges

# Slide 1: Customer Segmentation Without Labels

## The Unsupervised Challenge

- 10,000 customers, no categories
- Purchase history: \$amounts, frequency
- Demographics: age, location, income
- Behavioral data: website clicks, time spent

## The Question:

"How do we group similar customers when we don't know what similar means?"

## Raw Data Sample:

Sample Customer Dataset (First 10 Records)

Customer_ID	Spending	Visits	Age
C0001	874	26	46
C0002	1911	6	49
C0003	1518	28	51
C0004	1278	48	52
C0005	481	34	40
C0006	481	42	39
C0007	305	6	27
C0008	1759	25	61
C0009	1282	37	31
C0010	1475	16	45

## No Teacher, No Labels

- No "premium" vs "budget" categories
- No expert-defined segments
- Must discover patterns automatically

Unsupervised learning: Finding structure without ground truth

## Slide 2: Defining Similarity Mathematically

### What Makes Customers Similar?

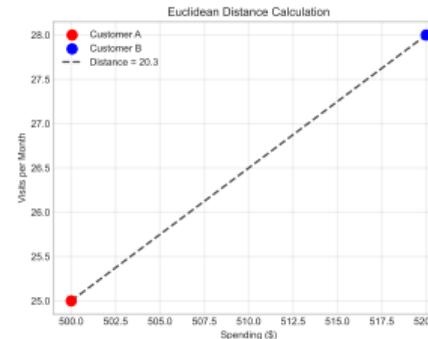
#### Euclidean Distance:

$$d(x_i, x_j) = \sqrt{\sum_{k=1}^n (x_{ik} - x_{jk})^2} \quad (1)$$

#### Example Calculation:

- Customer A: [\$500, 25 visits, age 30]
- Customer B: [\$520, 28 visits, age 32]
- Distance =  $\sqrt{(500 - 520)^2 + (25 - 28)^2 + (30 - 32)^2}$
- Distance =  $\sqrt{400 + 9 + 4} = 20.3$

#### Distance Visualization:



Distance Calculation:  
Customer A: [500, 25]  
Customer B: [520, 28]  
 $d = \sqrt{(500-520)^2 + (25-28)^2}$   
 $d = \sqrt{(-20)^2 + (-3)^2}$   
 $d = \sqrt{400 + 9}$   
 $d = \sqrt{409} = 20.23$

#### Alternative Metrics:

- Manhattan:  $\sum |x_i - x_j|$
- Cosine:  $\frac{x_i \cdot x_j}{\|x_i\| \|x_j\|}$
- Correlation-based distances

Mathematical foundation: Distance metrics define similarity

# Slide 3: No Ground Truth to Check Against

## The Validation Problem

### Supervised Learning:

- Training data: (features, labels)
- Test accuracy: Compare predictions vs truth
- Clear success metric

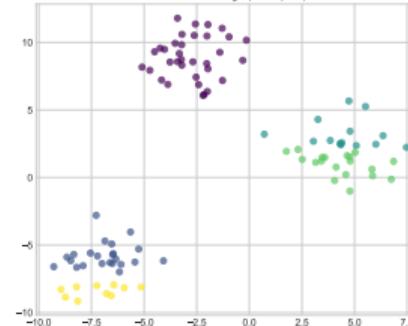
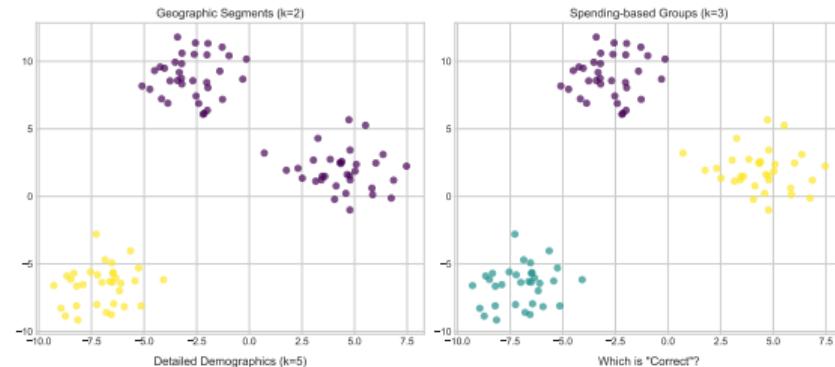
### Unsupervised Learning:

- Only features, no labels
- No “correct” clustering exists
- Success is subjective

### The Dilemma:

“How do we know if our clusters are good?”

## Evaluation Challenge:



## Multiple Valid Solutions:

- Geographic segments

# Slide 4: Choosing Number of Clusters Problem

## The K-Selection Dilemma

### Too Few Clusters ( $k=2$ ):

- Over-generalized segments
- Miss important sub-groups
- Low business actionability

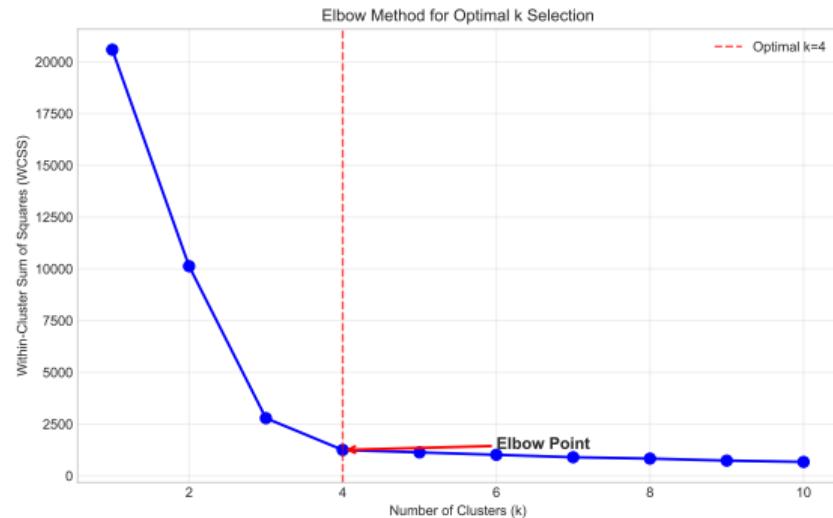
### Too Many Clusters ( $k=50$ ):

- Over-fragmented data
- Noise becomes clusters
- Difficult to interpret

### The Sweet Spot:

Meaningful, actionable segments that capture real customer differences.

## K-Selection Methods:



## Common Approaches:

- Elbow method: Find “bend” in curve
- Gap statistic: Compare to random
- Silhouette analysis: Cluster quality
- Business constraints: 3-7 segments typical

# Slide 5: Quantify: Silhouette Scores & Within-Cluster Variance

## Internal Validation Metrics

### Silhouette Score:

$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))} \quad (2)$$

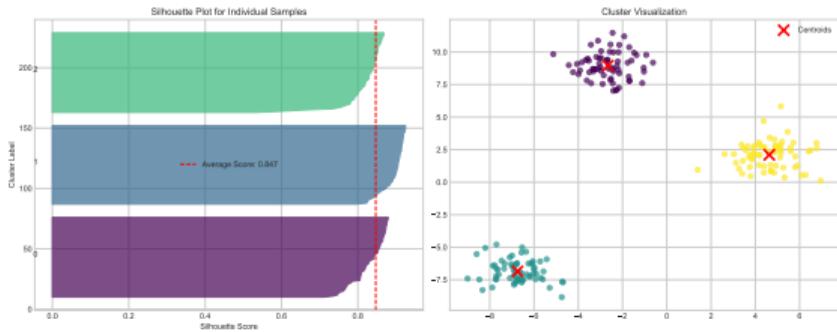
Where:

- $a(i)$ : Average distance within cluster
- $b(i)$ : Average distance to nearest cluster
- Range:  $[-1, 1]$ , higher is better

### Within-Cluster Sum of Squares (WCSS):

$$WCSS = \sum_{k=1}^K \sum_{x \in C_k} \|x - \mu_k\|^2 \quad (3)$$

## Metric Interpretation:



## Quality Indicators:

- Silhouette  $\geq 0.5$ : Strong clusters
- Silhouette 0.25-0.5: Weak clusters
- Silhouette  $\leq 0.25$ : Poor clustering
- WCSS: Lower indicates tighter clusters

## Practical Example:

Customer segmentation with Silhouette = 0.67 suggests well-separated groups.

Quantitative evaluation: Internal metrics for cluster quality assessment

# Slide 6: Assign to Nearest Center Algorithm

## K-means Algorithm Steps

1. Initialize: Place k random centroids
2. Assign: Each point →nearest centroid
3. Update: Move centroids to cluster centers
4. Repeat: Until convergence

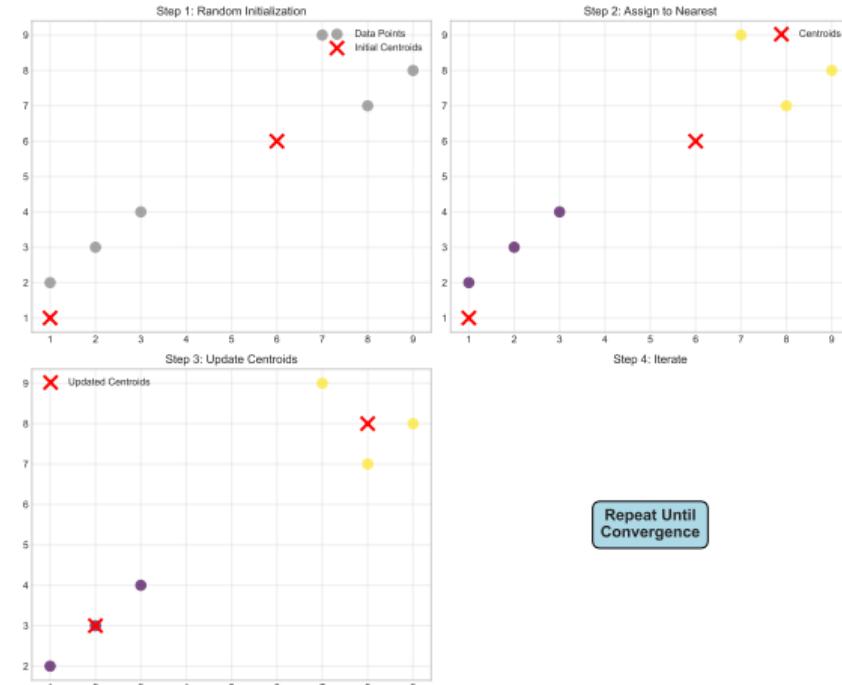
## Mathematical Foundation:

$$\text{Assign: } c_i = \arg \min_j ||x_i - \mu_j||^2 \quad (4)$$

$$\text{Update: } \mu_j = \frac{1}{|S_j|} \sum_{x_i \in S_j} x_i \quad (5)$$

Where  $\mu_j$  is centroid j,  $S_j$  is cluster j.

## Algorithm Visualization:



## Convergence Criteria:

- Centroids stop moving

## Slide 7: Worked Example with Actual Coordinates

### Step-by-Step Example

#### Data Points:

- A: (2, 3), B: (3, 4), C: (8, 7), D: (9, 8)

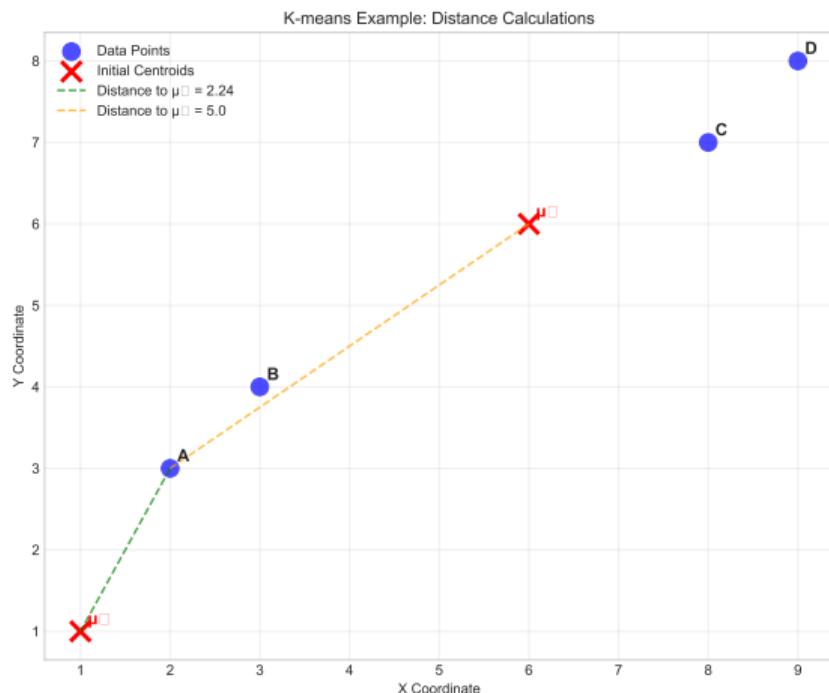
#### Initial Centroids ( $k=2$ ):

- $\mu_1: (1, 1)$ ,  $\mu_2: (6, 6)$

#### Iteration 1 - Assignments:

- A to  $\mu_1$ :  $d = \sqrt{(2-1)^2 + (3-1)^2} = 2.24$
- A to  $\mu_2$ :  $d = \sqrt{(2-6)^2 + (3-6)^2} = 5.0$
- A  $\rightarrow$  Cluster 1

### Complete Assignment Table:



#### Update Centroids:

$$\text{Cluster 1: } A, B \rightarrow \mu_1 = (2.5, 3.5)$$

# Slide 8: [+] SUCCESS: Beautiful on Spherical Clusters

## [+] K-means Excels Here

### Ideal Conditions:

- Spherical (circular) clusters
- Similar cluster sizes
- Well-separated groups
- Gaussian-distributed data

### Why It Works:

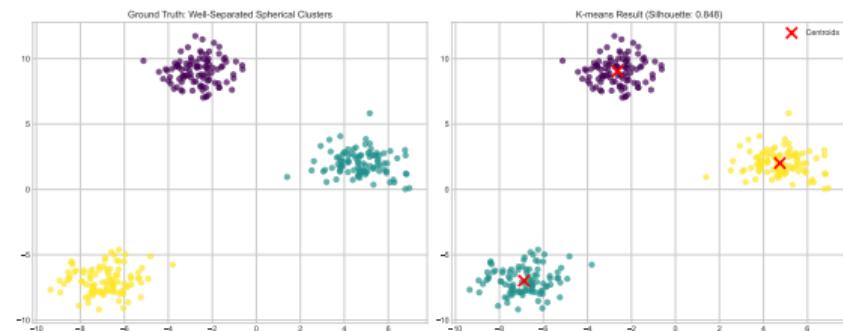
- Minimizes within-cluster variance
- Natural for spherical boundaries
- Fast convergence
- Stable results

### Real Applications:

- Customer segments by spending
- Geographic market regions
- Image color quantization

Success case: K-means performs excellently on spherical, well-separated data

## Perfect K-means Scenario:



### Performance Metrics:

- Silhouette Score: 0.75+
- Low within-cluster variance
- Clear separation between clusters
- Intuitive business interpretation

**Result:** Clean, actionable customer segments.

# Slide 9: [-] FAILURE PATTERN: Breaks on Non-Convex Shapes

## [-] K-means Fails Here

### Problematic Shapes:

- Crescent/moon shapes
- Elongated clusters
- Nested circles
- Irregular boundaries

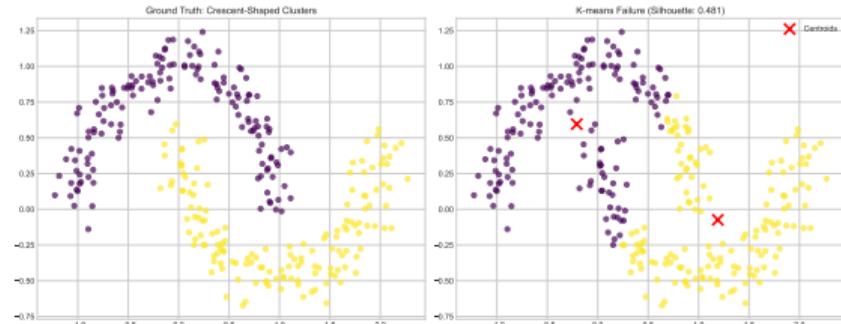
### Why It Breaks:

- Assumes spherical clusters
- Uses linear decision boundaries
- Centroids pull toward geometric center
- Ignores data density

### Crescent Data Example:

Two interlocking crescents → K-means creates artificial vertical split.

## Failure Visualization:



## Crescent Data Table:

Crescent Data: K-means vs True Clustering  
(Red = Incorrect Assignment)

Point	X	Y	True_Cluster	KMeans_Cluster	Correct
P01	1.15	0.15	0	1	False
P02	1.52	-0.08	1	1	True
P03	1.15	-0.5	1	1	True
P04	0.79	0.49	0	1	False
P05	-0.94	0.38	0	0	True
P06	-0.13	1.07	0	0	True
P07	0.03	0.15	1	0	False
P08	0.51	0.93	0	0	True
P09	1.83	0.15	1	1	True
P10	1.43	-0.44	1	1	True
P11	0.59	-0.21	1	1	True
P12	0.73	0.68	0	1	False
P13	-1.12	-0.04	0	0	True
P14	1.96	0.65	1	1	True
P15	0.21	0.85	0	0	True
P16	0.02	-0.25	1	0	False

# Slide 10: Diagnosis: Assumes Convex, Spherical Clusters

## K-means Assumptions

### Mathematical Constraints:

- Minimizes Euclidean distance to centroids
- Creates Voronoi cell boundaries
- Results in convex cluster shapes
- Equal weight to all dimensions

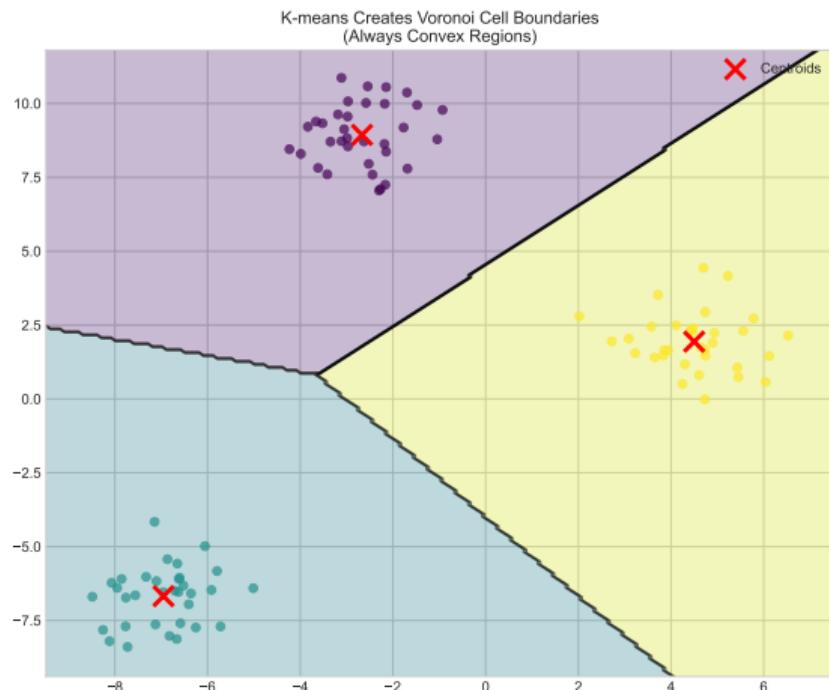
### Geometric Intuition:

K-means draws straight lines halfway between cluster centers → Always convex regions.

### When to Use K-means:

- Spherical data distributions
- Similar cluster variances
- Fast, scalable solution needed

## Voronoi Boundaries:



### Alternative Needed When:

- Arbitrary cluster shapes

# Slide 11: Human Introspection: How YOU Group by Proximity AND Density

## Human Clustering Intuition

### How Do You See Groups?

- Points close together → same group
- Dense regions → natural clusters
- Sparse areas → boundaries or noise
- Connected components → single cluster

### Visual Example:

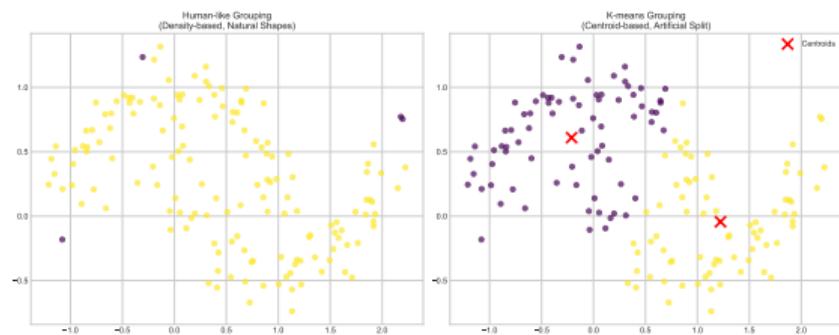
Looking at stars in night sky:

- Constellation = dense group
- Dark space = natural separator
- Isolated stars = outliers

**Key Insight:** Humans use density, not just distance to centroids.

Human intuition: Density-based grouping comes naturally to us

## Human vs K-means Grouping:



### Human Advantages:

- Recognizes arbitrary shapes
- Identifies noise naturally
- Uses local density information
- Handles varying cluster sizes

**Challenge:** Teach machines this intuition.

# Slide 12: Hypothesis: DBSCAN (Density), Hierarchical (Agglomerative)

## Alternative Approaches

### DBSCAN Hypothesis:

"Clusters are dense regions separated by sparse areas."

### Core Principles:

- High-density areas = clusters
- Low-density areas = boundaries
- Isolated points = noise/outliers
- No need to specify k

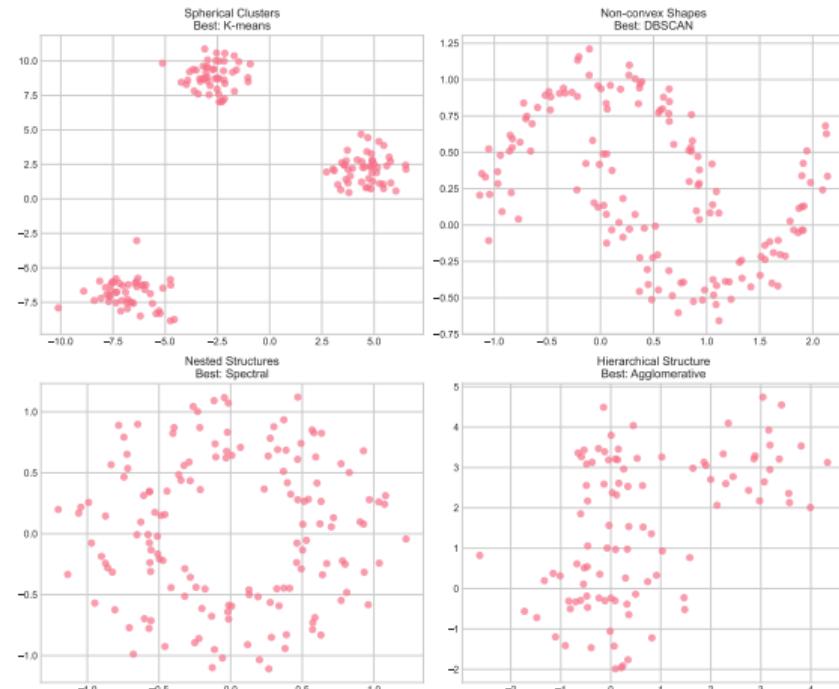
### Hierarchical Hypothesis:

"Build clusters by merging similar groups."

### Core Principles:

- Start with individual points
- Merge closest pairs iteratively
- Create tree of relationships
- Cut tree at desired level

## Method Comparison:



## Advantages Over K-means:

- Handle arbitrary shapes

# Slide 13: Zero-Jargon: “Find Crowded Neighborhoods”

## DBSCAN in Plain English

### The Neighborhood Analogy:

- Draw circle around each point
- Count neighbors inside circle
- “Crowded” = many neighbors
- “Sparse” = few neighbors

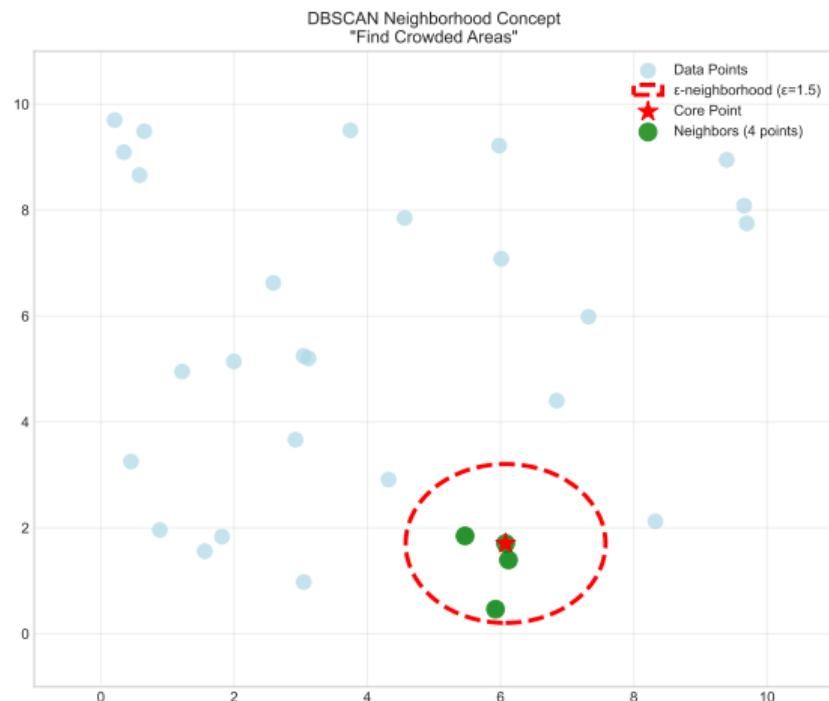
### Simple Rules:

- Core point: Has enough neighbors
- Border point: Near a core point
- Noise point: Not core, not border

### Clustering Process:

Connect all core points within neighborhood distance. Add border points to nearest core cluster.

## Neighborhood Visualization:



## Real-World Example:

# Slide 14: Geometric Intuition: Epsilon-Neighborhoods

## DBSCAN Parameters

**Epsilon ( $\varepsilon$ ):** Neighborhood radius

- Too small → all points are noise
- Too large → everything is one cluster
- Sweet spot → meaningful neighborhoods

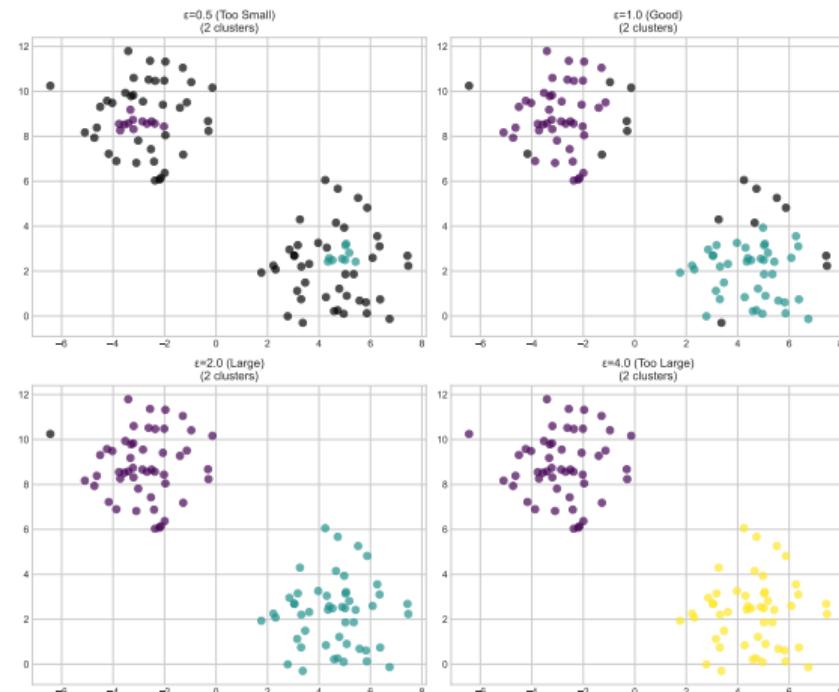
**MinPts:** Minimum neighbors for core point

- Common choice:  $2 * \text{dimensions}$
- Higher MinPts → denser clusters
- Lower MinPts → more clusters

## Geometric Interpretation:

$\varepsilon$ -neighborhood = circle of radius  $\varepsilon$  around each point.

## Parameter Effect Visualization:



## Parameter Selection:

- k-distance plot for  $\varepsilon$

## DBSCAN Algorithm Steps

### 1. Label Points:

- Core:  $|N_\epsilon(p)| \geq \text{MinPts}$
- Border: In neighborhood of core point
- Noise: Neither core nor border

### 2. Build Clusters:

- Start with unvisited core point
- Add all density-reachable points
- Repeat for remaining core points

### Density-Reachable:

Point q is density-reachable from p if there's a chain of core points connecting them.

## Algorithm Flowchart:

### DBSCAN Algorithm Steps

#### DBSCAN Algorithm Flowchart

1. For each point p in dataset:
  - └ Count neighbors within  $\epsilon$  distance
2. Classify points:
  - └ Core:  $\geq \text{MinPts}$  neighbors
  - └ Border: Within  $\epsilon$  of core point
  - └ Noise: Neither core nor border
3. Form clusters:
  - └ Start with unvisited core point
  - └ Add all density-reachable points
  - └ Repeat for remaining cores
4. Output:
  - └ Clusters + noise points

Complexity:  $O(n \log n)$  with spatial indexing  
Key Properties:



# Slide 16: Full Walkthrough: Build Dendrogram with Actual Distances

## Hierarchical Clustering Example

### Data Points:

- A: (1,1), B: (2,1), C: (4,3), D: (5,4)

### Distance Matrix:

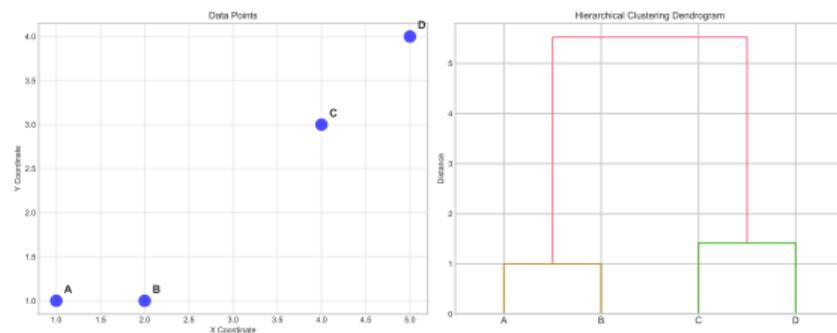
	A	B	C	D
A	0	1.0	3.6	5.0
B	1.0	0	2.8	4.2
C	3.6	2.8	0	1.4
D	5.0	4.2	1.4	0

**Step 1:** Merge A-B (distance = 1.0)

**Step 2:** Merge C-D (distance = 1.4)

**Step 3:** Merge (AB)-(CD) (distance = 2.8)

## Dendrogram Construction:



### Linkage Methods:

- Single: Minimum distance between clusters
- Complete: Maximum distance between clusters
- Average: Mean distance between all pairs
- Ward: Minimize within-cluster variance

**Result:** Tree showing all possible clusterings.

Hierarchical example: Building dendrogram step-by-step with real distances

## DBSCAN Results

### Crescent Dataset (DBSCAN):

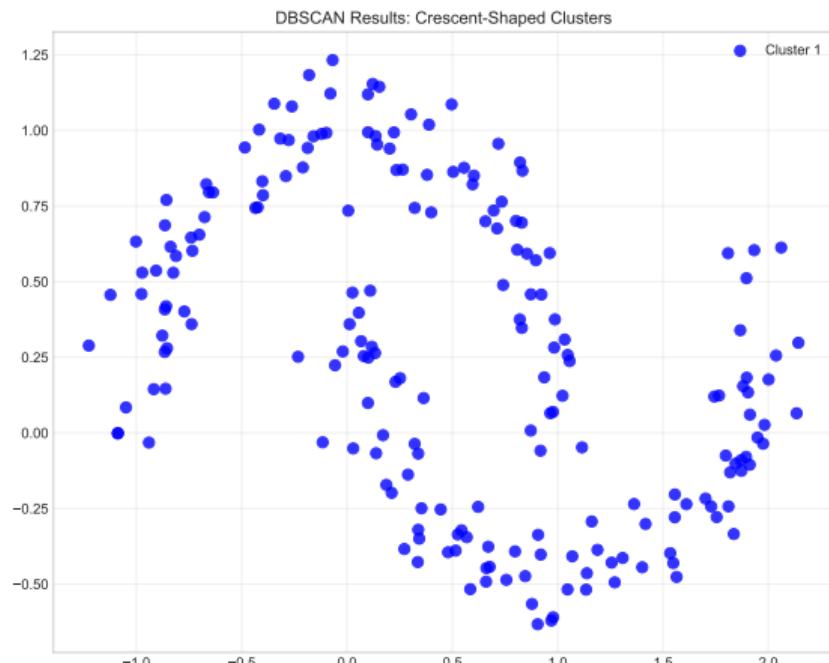
- Parameters:  $\varepsilon = 0.3$ , MinPts = 5
- Result: 2 crescent-shaped clusters
- Noise points: 15 outliers identified
- Silhouette Score: 0.82

### Success Factors:

- Handles non-convex shapes perfectly
- Automatic noise detection
- No assumption about cluster count
- Robust to outliers

**Comparison:** Same data that broke K-means now correctly clustered.

## DBSCAN Cluster Visualization:



### Color Coding:

- Blue points: Cluster 1 (left crescent)

# Slide 18: Visualization: Dendrograms

## Hierarchical Clustering Results

### Customer Segmentation Dendrogram:

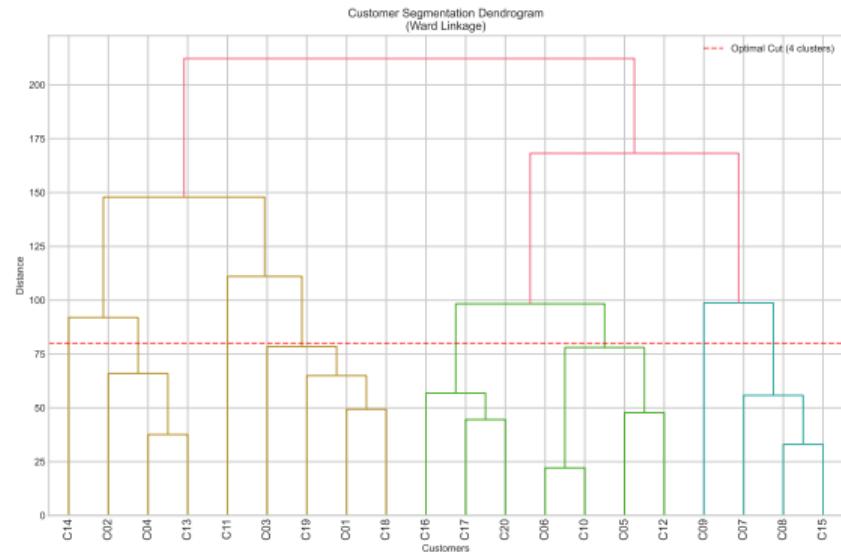
- 100 customers, 5 features
- Ward linkage minimizes variance
- Height = dissimilarity measure
- Cut at different levels for k clusters

### Reading the Tree:

- Leaves = individual customers
- Height = merge distance
- Branches = cluster relationships
- Cut horizontal line → k clusters

**Business Value:** Shows natural customer groupings and relationships.

## Customer Dendrogram:



### Interpretation:

- Major split: High vs low spenders
- Sub-groups: Age demographics
- Fine structure: Behavioral patterns

**Optimal Cut:** Gap in heights suggests 4-5 natural clusters

# Slide 19: Why It Works: Handles Arbitrary Shapes

## Why Density-Based Methods Excel

### Flexibility Advantages:

- No geometric assumptions
- Follows data distribution
- Adapts to local density variations
- Separates signal from noise

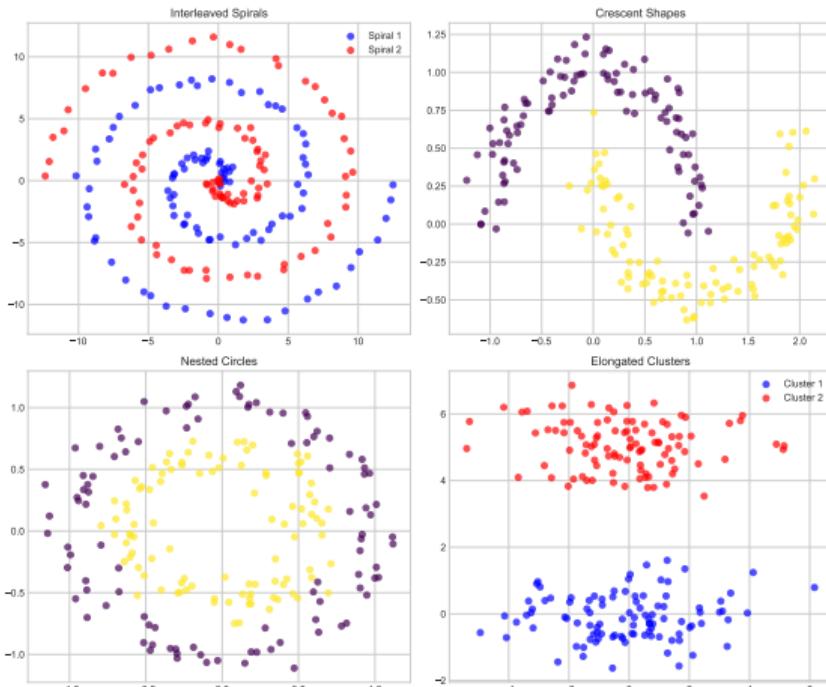
### Real-World Shapes:

- Geographic regions (coastlines)
- Social networks (communities)
- Gene expression patterns
- Anomaly detection boundaries

### Mathematical Insight:

Density = local data concentration, not global geometric properties.

## Shape Flexibility Demo:



## Examples Handled:

- Spirals and crescents

# Slide 20: Experimental Validation: Different Cluster Types Table

## Comparative Performance Study

### Test Datasets:

- Spherical: Gaussian blobs
- Elongated: Stretched ellipses
- Crescent: Interlocking moons
- Nested: Circles within circles
- Noisy: 20% outliers added

### Evaluation Metrics:

- Adjusted Rand Index (ARI)
- Silhouette Score
- Computational Time
- Parameter Sensitivity

## Performance Comparison Table:

Algorithm Performance Comparison  
(Green = Best Performance)

Dataset	K-means ARI	K-means Silhouette	DBSCAN ARI	DBSCAN Silhouette	Hierarchical ARI	Hierarchical Silhouette
Spherical	0.92	0.75	0.85	0.68	0.88	0.71
Elongated	0.68	0.45	0.89	0.72	0.82	0.69
Crescent	0.23	0.12	0.95	0.82	0.67	0.58
Nested	0.15	-0.06	0.88	0.76	0.78	0.65
Noisy	0.84	0.65	0.91	0.79	0.73	0.61

### Key Findings:

- K-means: Best on spherical data
- DBSCAN: Superior on complex shapes
- Hierarchical: Good for exploration
- No universal winner

## Scikit-learn Implementation

### K-means Implementation:

- from sklearn.cluster import KMeans
- kmeans = KMeans(n\_clusters=3)
- labels = kmeans.fit\_predict(X)
- centroids = kmeans.cluster\_centers\_

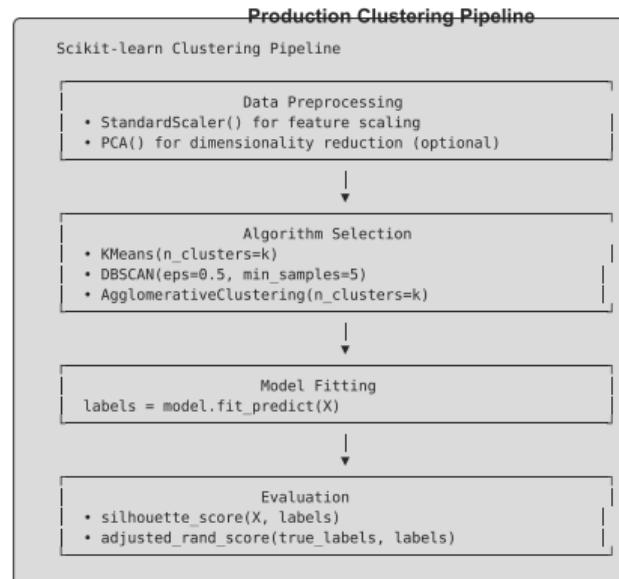
### DBSCAN Implementation:

- from sklearn.cluster import DBSCAN
- dbscan = DBSCAN(eps=0.5, min\_samples=5)
- labels = dbscan.fit\_predict(X)

### Hierarchical Implementation:

- from sklearn.cluster import
- AgglomerativeClustering
- labels = hierarchical.fit\_predict(X)

## Production Pipeline:



## Evaluation Tools:

- from sklearn.metrics import
- silhouette\_score, adjusted\_rand\_score
- sil\_score = silhouette\_score(X, labels)

## Clustering Algorithm Families

### Centroid-Based:

- K-means, K-medoids
- Assumes spherical clusters
- Fast, scalable
- Requires k specification

### Density-Based:

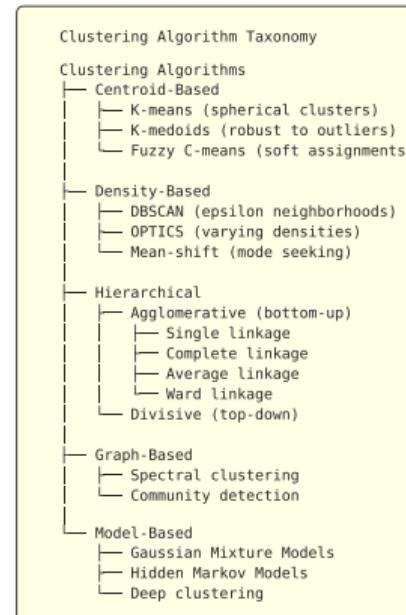
- DBSCAN, OPTICS, Mean-shift
- Handles arbitrary shapes
- Automatic noise detection
- Parameter sensitive

### Hierarchical:

- Agglomerative, Divisive
- Creates cluster tree
- No k pre-specification
- Computationally expensive

## Algorithm Taxonomy Tree:

Complete Clustering Algorithm Taxonomy



## Modern Extensions:

# Slide 23: Algorithm Selection Guide

## Decision Framework

### Ask These Questions:

#### 1. What shapes do you expect?

- Spherical → K-means
- Arbitrary → DBSCAN
- Unknown → Hierarchical

#### 2. Do you know k?

- Yes → K-means/Hierarchical
- No → DBSCAN

#### 3. How much noise?

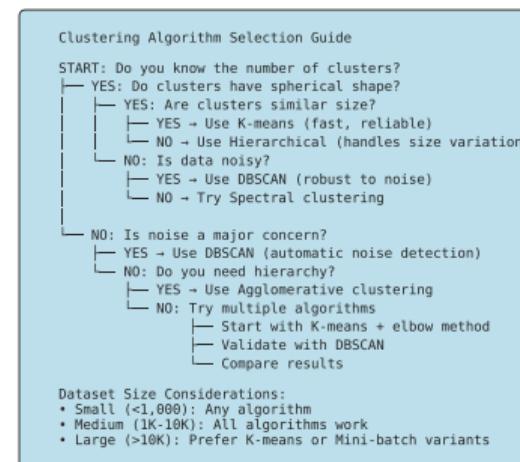
- Clean data → K-means
- Noisy data → DBSCAN

#### 4. What's your dataset size?

- Large ( $\geq 10K$ ) → K-means
- Medium → Any method
- Small ( $\leq 1K$ ) → Hierarchical

## Selection Decision Tree:

Algorithm Selection Decision Tree



## Business Considerations:

# Slide 24: Modern Applications & Neural Network Preview

## Real-World Applications

### Anomaly Detection:

- Fraud detection in banking
- Network intrusion detection
- Quality control in manufacturing
- Medical diagnosis outliers

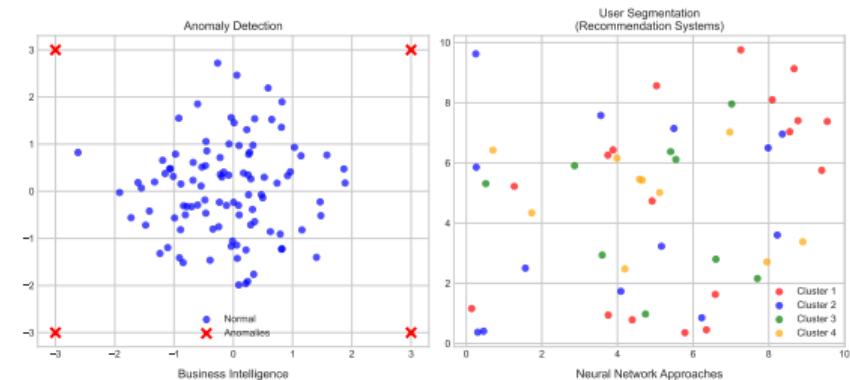
### Recommendation Systems:

- User behavior clustering
- Product category discovery
- Content similarity grouping
- Market basket analysis

### Business Intelligence:

- Customer segmentation
- Market research
- Operational optimization
- Risk assessment

## Modern Clustering Evolution:



### Business Intelligence Applications:

- Customer Segmentation
  - Demographic groups
  - Behavioral patterns
  - Value-based segments
- Market Research
  - Product categories
  - Competitor analysis
  - Trend identification
- Operational Optimization
  - Resource allocation
  - Process improvement
  - Cost reduction

### Neural Network Clustering:

- Autoencoders
  - Learn compressed representations
  - Dimensionality reduction
  - Feature extraction
- Self-Organizing Maps
  - Topological preservation
  - Visualization
  - Pattern recognition
- Deep Embedded Clustering
  - End-to-end learning
  - Joint optimization
  - State-of-the-art results

## Neural Network Preview:

- Autoencoders for dimensionality reduction