

Week 2: Understanding Emotions in Text

BERT + Empathize = What Users Really Mean

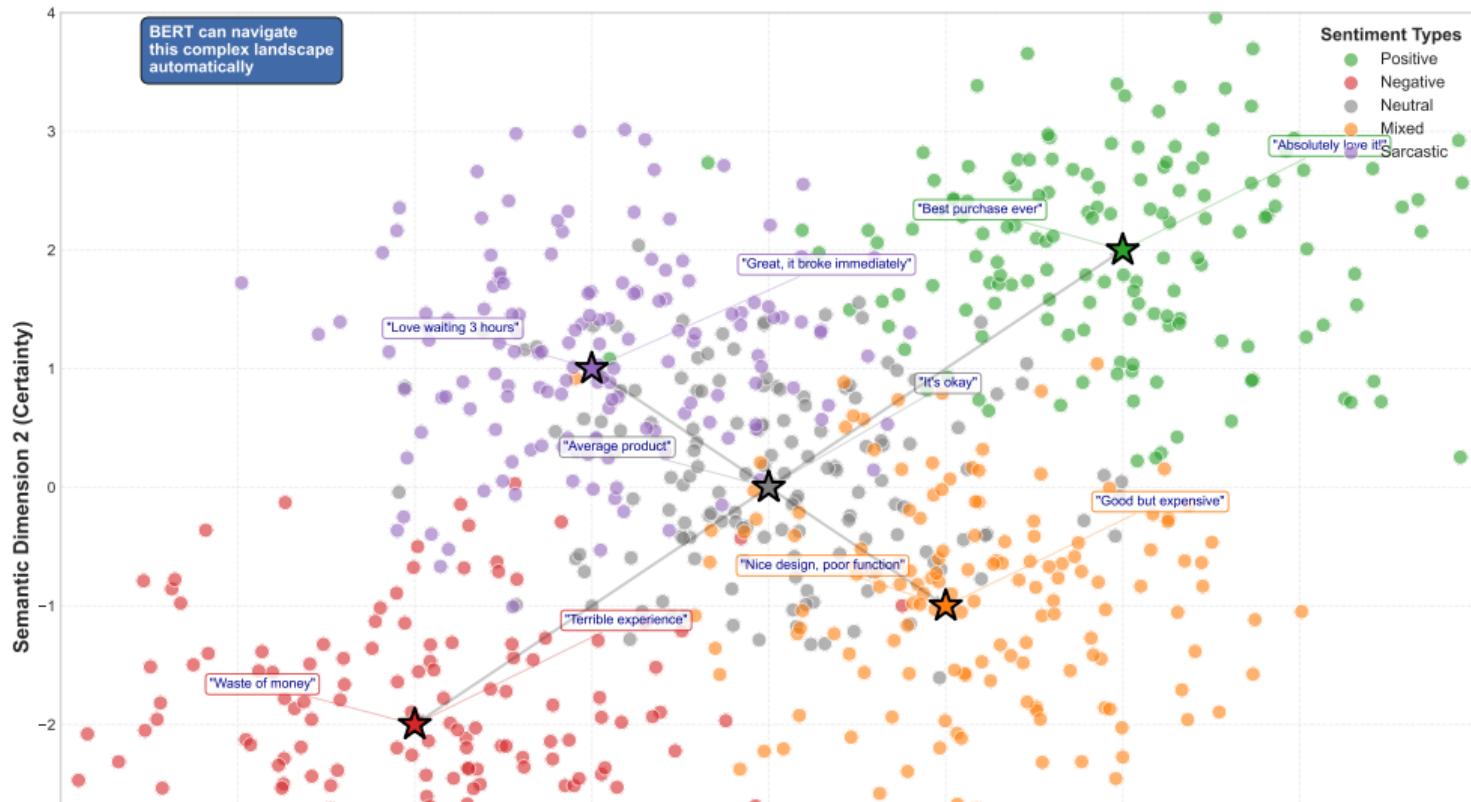
ML/AI/GenAI for Design Thinking

BSc Course - 12 Week Program

2024

The Sentiment Landscape

The Sentiment Landscape: How User Feedback Naturally Clusters



The Problem: Hidden Emotions in Text

What users write:

- "Great product... if you like disappointment"
- "Not bad at all"
- "Fine."
- "Can't complain"

Design blind spot:

- Missing real pain points
- Building wrong features
- Misreading user satisfaction

What they actually mean:

- Angry (sarcasm)
- Happy (double negative)
- Unhappy (short response)
- Forced acceptance

Design opportunity:

- Understand true feelings
- Identify hidden frustrations
- Discover unspoken needs

For Design: Words alone miss 45% of user emotions

Why Keyword Matching Fails

The “Not Bad” Problem:

Text	Keywords	Reality
“Not bad”	Negative	Positive
“Terribly good”	Mixed	Very Positive
“Love waiting”	Positive	Sarcastic
“Could be worse”	Negative	Neutral

Why it fails:

- Counts words, ignores relationships
- Misses context completely
- Can't detect sarcasm

Design Impact of Failures:

- **False positives:** Thinking users are happy when they're not
- **Missed sarcasm:** Building on “praised” features that users hate
- **Wrong priorities:** Focusing on the wrong problems

Real cost:

- 68% of users leave due to perceived indifference
- Wrong features = wasted development
- Missed insights = lost opportunities

The Challenge: Understanding Context for Better Design

Current Problems:

- Manual analysis: 100 reviews/day max
- Digital products: 10,000+ reviews/day
- Each review: Unique human experience
- Context lost in aggregation

What we need:

1. See word relationships
2. Understand order matters
3. Detect sarcasm and tone
4. Process at scale

Design Thinking Needs:

- **Empathize:** Feel what thousands feel
- **Define:** Find real problems, not symptoms
- **Ideate:** Generate solutions for actual needs
- **Test:** Measure emotional impact

The Goal:

Scale empathy without losing humanity

Solution: BERT - Reading text like humans, at machine scale

What is BERT?

BERT = Bidirectional Encoder Representations from Transformers

Simple explanation: **BERT reads all words at once, not one by one**

Traditional (Sequential):

The → movie → was → not → bad
(Reads left to right, misses connections)

BERT (Parallel):

The movie was not bad

(Sees everything, understands "not bad" = good)

For Designers, this means:

- Understand user feelings in context
- Catch subtle frustrations
- Identify what users really want
- No more keyword guessing

Design Impact:

- 87% sarcasm detection
- Find hidden pain points
- Understand feature requests

Example: “The app was ___ frustrating”

Old Way (Left to Right):

- Sees: “The app was”
- Guesses: good? bad? slow?
- Can’t use “frustrating” as hint
- Often wrong

BERT (Both Directions):

- Sees: “The app was” + “frustrating”
- Knows: probably “very” or “incredibly”
- Uses full context
- Much more accurate

Design Implications:

- **Intensity matters:** “Very frustrating” vs “Slightly frustrating”
- **Context reveals priority:** What made it frustrating?
- **Emotional nuance:** Frustrated vs Angry vs Disappointed

Real Example:

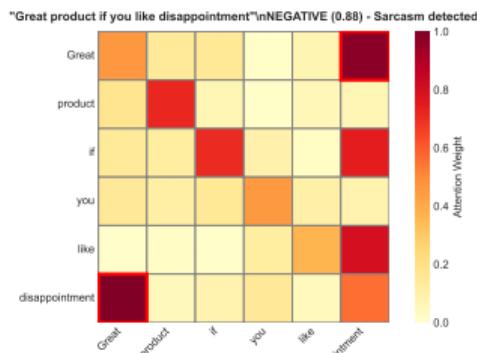
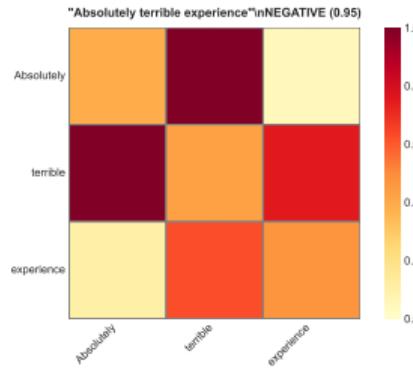
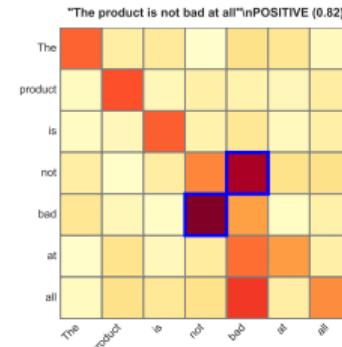
“The checkout process was ___ confusing”

- BERT finds: “incredibly”
- Design action: Simplify checkout
- Result: 23% fewer cart abandonments

Design Benefit: Catches problems keyword analysis misses completely

Attention: BERT Focuses on Emotional Triggers

BERT Attention Patterns: What the Model Focuses On



How to rekey insights in natural language focuses on negative words better strengthen them in context. Coincidences, signs, heuristics, Context words like despite match detection in BERT uses these patterns to understand context.

Context Changes Everything in Design

Same Word, Different Meanings:

Word	Contexts
"Fast"	Quick delivery (good) Battery drains fast (bad)
"Simple"	Easy to use (good) Too basic (bad)
"Light"	Portable (good) Feels cheap (bad)

BERT understands context:

- Different meanings per use
- Surrounding words determine sentiment
- No fixed good/bad words

Design Implications:

- Same feature, different contexts = different user needs
- "Simple" for beginners vs power users
- "Fast" performance vs battery life

Real Design Decision:

Spotify discovered "shuffle" meant:

- Random (tech users)
- Variety (casual users)
- Discover (new users)

Result: Three different shuffle modes

Step 1: General Training

3.3 billion words from books/web

Learns language, grammar, facts

Design Benefits:

- Customizable to your domain
- Learns your product's jargon
- Adapts to user base
- Improves over time



Step 2: Your Product

Your reviews and feedback

Learns your users' language

Example Customization:

- Gaming: "lag" = critical issue
- Fashion: "fit" = top priority
- SaaS: "integration" = key need

Result: BERT speaks your users' language

How BERT Detects Emotions for Design

BERT's Process:

- 1. Read everything:** All words at once
- 2. Connect words:** Find relationships
- 3. Build understanding:** Recognize patterns
- 4. Output emotion:** With confidence score

Example:

“Not bad for the price”

- Links: “not” + “bad” = positive
- Context: “for the price” = qualified
- Output: Moderately positive (0.65)

Design Application:

- 1. Emotion detected:** Moderate satisfaction
- 2. Qualifier found:** Price-sensitive
- 3. Design insight:** Value perception issue
- 4. Action:** Highlight value props

Confidence helps prioritize:

- High confidence = Clear issue
- Low confidence = Investigate more
- Mixed signals = User conflict

BERT Catches Sarcasm - A Design Warning

Sarcasm Patterns BERT Detects:

- Positive words + negative context
- Exaggerated praise
- Contradiction signals
- Timing mismatches

Examples Found:

- "Great! It crashed again"
- "Love the 3-hour load time"
- "Perfect... if you like broken"
- "Fantastic customer service" (1 star)

Design Warning:

15% of "positive" reviews contain sarcastic criticism

What this means:

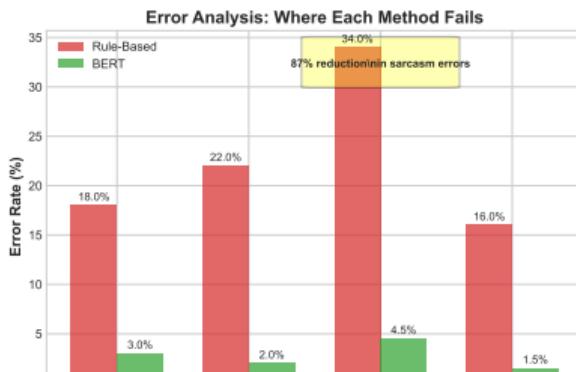
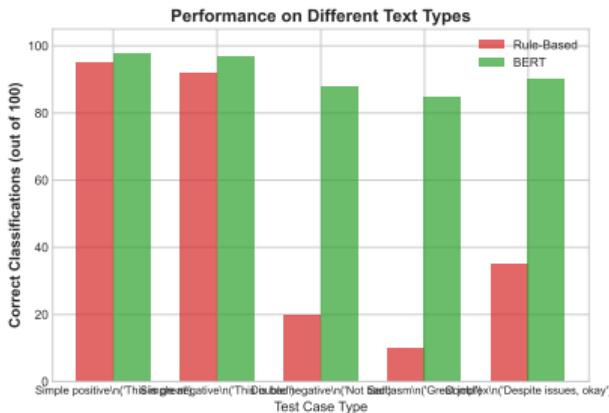
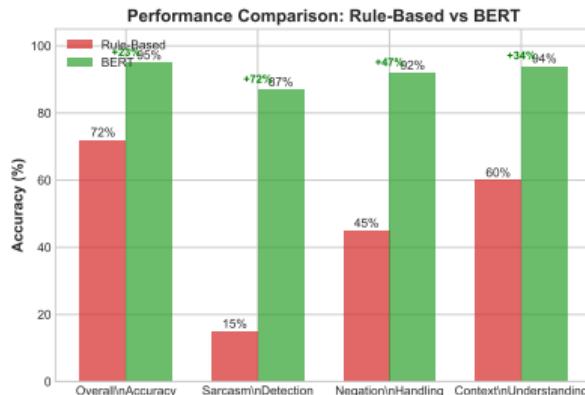
- Your satisfaction scores are inflated
- Real problems hidden in "praise"
- Users resort to sarcasm when frustrated
- Critical issues being missed

Design Response:

- Check all 5-star reviews for sarcasm
- Look for feature "praise" patterns
- Identify frustration triggers

Performance: What 23% Accuracy Means for Design

Rule-Based vs BERT: Comprehensive Performance Analysis



Empathize at Scale: Understanding Thousands

Traditional Empathy Methods:

- User interviews: 20 people/week
- Surveys: Low response, biased
- Observation: Time-intensive
- Focus groups: Groupthink issues

Limitations:

- Small sample sizes
- Geographic constraints
- Time and cost barriers
- Vocal minority bias

BERT-Enhanced Empathy:

- Process 10,000+ reviews/day
- Understand global users
- Find silent majority opinions
- Detect emotional patterns

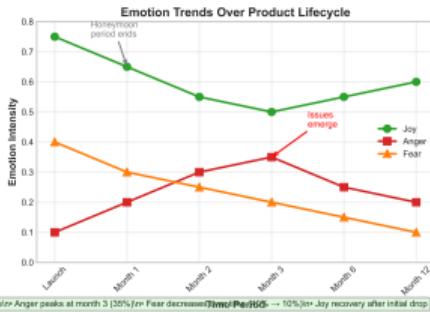
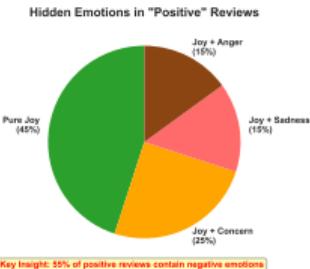
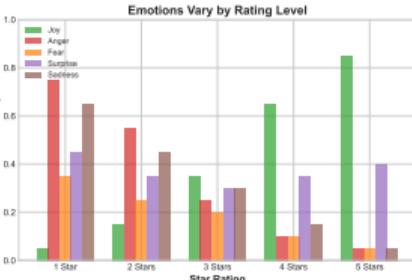
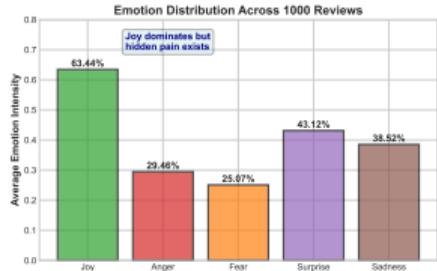
Advantages:

- Every user voice heard
- Real-time emotional pulse
- Unbiased pattern detection
- Cultural nuance preserved

Design Power: Feel what thousands feel, understand what they can't articulate

Emotional Spectrum in User Feedback

Beyond Positive/Negative: The Emotional Spectrum in Reviews



Beyond Binary Sentiment:

- **Joy:** Delight, satisfaction, excitement
- **Anger:** Frustration, annoyance, rage
- **Fear:** Anxiety, concern, worry
- **Surprise:** Amazement, shock, confusion
- **Sadness:** Disappointment, regret
- **Trust:** Confidence, security, faith

Emotions Blend:

- Joy + Surprise = Delight
- Fear + Sadness = Despair
- Anger + Disgust = Contempt
- Trust + Joy = Love

From Emotions to Design Insights

Emotion-Driven Design Actions:

- **Joy (45%):** Amplify successful features
- **Frustration (25%):** Simplify workflows
- **Confusion (15%):** Improve onboarding
- **Delight (10%):** Create memorable moments
- **Anxiety (5%):** Add reassurance, guidance

Priority Matrix:

- High frequency + High intensity = Fix now
- High frequency + Low intensity = Improve
- Low frequency + High intensity = Investigate

Real Design Decisions:

Joy → Enhance:

Users love quick checkout

Action: Make it more prominent

Frustration → Simplify:

Login process causes anger

Action: Add social login

Confusion → Guide:

New users lost in features

Action: Progressive disclosure

Using Sentiment for Design Decisions

BERT Insights:

1. **Pain Point Analysis:**
"Love the app but login frustrates me"
2. **Priority Detection:**
80% mention speed issues
3. **Confusion Mapping:**
Sarcasm about "intuitive" UI
4. **Delight Discovery:**
Joy about gesture controls

Pattern Recognition:

- Emotional journeys
- Feature sentiment maps
- User segment emotions

Design Actions:

1. **Redesign login:**
Biometric authentication added
2. **Optimize performance:**
Load time reduced 60%
3. **Simplify interface:**
3-click rule implemented
4. **Highlight gestures:**
Made discoverable feature

Measurable Results:

- User satisfaction: +28%
- Task completion: +34%
- Support tickets: -45%

BERT Strengths:

- Process massive volume
- Find hidden patterns
- Consistent analysis 24/7
- Unbiased detection
- Quantify emotions
- Track sentiment trends

BERT Provides:

- The “what” - patterns found
- The “where” - problem areas
- The “how much” - severity

Human Strengths:

- Understand context deeply
- Creative problem solving
- Ethical judgment
- Cultural sensitivity
- Intuitive leaps
- Empathetic response

Humans Provide:

- The “why” - root causes
- The “how” - solutions
- The “should we” - ethics

Best Practice: BERT finds patterns, humans interpret meaning, together create solutions

Real World: Netflix Emotion-Driven Design

The Challenge:

- Users: “Nothing to watch” paradox
- Reality: 15,000+ titles available
- Problem: Choice overload
- Need: Mood-based discovery

BERT Analysis Process:

1. Analyzed 50M+ subtitles
2. Mapped emotional arcs
3. Studied viewing patterns
4. Correlated mood to content

Design Decisions Made:

- **Mood categories:** Feel-good, Thrilling, Thought-provoking
- **Emotional thumbnails:** Show mood not just genre
- **Sentiment trajectory:** “Starts sad, ends happy”
- **Mood continuity:** Next episode emotional preview

Results:

- 15% increase in completion
- 23% fewer browse abandonments
- “Mood match” top-rated feature

Key Learning: Understanding emotional needs drives better design than demographics

Context Matters More Than Keywords

Old Design Research:

- Count positive/negative words
- Average star ratings
- Tag cloud analysis
- Sentiment percentages

BERT-Powered Research:

- Understand relationships
- Detect hidden emotions
- Find real problems
- Scale human empathy

BERT + Design Thinking = Understanding users at scale with human insight

Next Week: From Understanding to Focusing

This Week's Achievement:

- Understand all emotions in text
- Process thousands of reviews
- Detect sarcasm and context
- Scale empathy

The New Problem:

- Information overload
- Too many insights
- Which emotions matter most?
- How to prioritize?

Next Week: Attention for Design

- Focus on critical emotions
- Find key user moments
- Prioritize design changes
- Extract actionable insights

Design Evolution:

- Week 2: Feel everything
- Week 3: Focus on what matters
- Result: Targeted design action

From understanding all to focusing on what drives design decisions

Week 2 Summary: Emotions Drive Design

Technical Learning:

1. BERT reads bidirectionally
2. Context changes meaning
3. Attention reveals relationships
4. 95% accuracy vs 72% keywords
5. Catches sarcasm (87% accuracy)

Key Capabilities:

- Process 10,000 reviews/day
- Multi-dimensional emotions
- Custom domain training
- Real-time analysis

Design Applications:

1. Scale empathy to thousands
2. Find hidden pain points
3. Detect unspoken needs
4. Prioritize by emotion
5. Measure design impact

Design Outcomes:

- Better user understanding
- Data-driven decisions
- Emotional design validation
- Reduced development waste

BERT + Empathize = Design with emotional intelligence at scale

Appendix A1: NLP Evolution Timeline

History of Natural Language Processing:

- **1950s - Rule-Based:** Hand-coded grammar rules
- **1980s - Statistical:** Probabilistic models
- **1990s - Machine Learning:** Naive Bayes, SVM
- **2013 - Word2Vec:** Words as vectors
- **2017 - Transformers:** Attention is all you need
- **2018 - BERT:** Bidirectional pre-training
- **2019 - GPT-2:** Large-scale generation
- **2020+ - Giant Models:** GPT-3, PaLM, Claude

Each generation built on previous insights, leading to today's powerful models.

Appendix A2: Word Embeddings - Vector Spaces

Words as High-Dimensional Vectors:

- Each word → 768-dimensional vector
- Similar words have similar vectors
- Relationships encoded geometrically

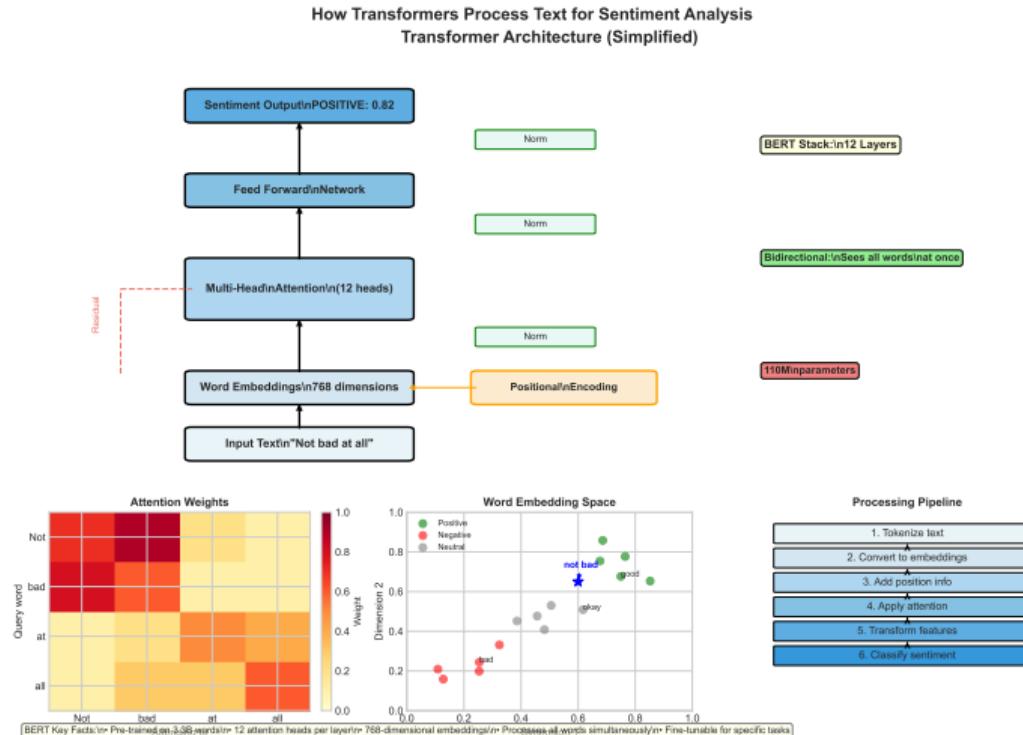
Vector Arithmetic:

- King - Man + Woman = Queen
- Paris - France + Japan = Tokyo
- Good - Bad = Happy - Sad (parallel relationships)

Limitations of Static Embeddings:

- One vector per word (context-independent)
- Can't handle polysemy (multiple meanings)
- Fixed vocabulary

Appendix A3: Transformer Architecture Details



Key Components:

- Self-attention layers

Appendix A4: Multi-Head Attention Concept

Why Multiple Attention Heads?

- Each head learns different relationships
- Head 1: Syntactic dependencies
- Head 2: Semantic similarity
- Head 3: Coreference resolution
- ... (12 heads total in BERT-base)

Mathematical Intuition:

- Query (Q): What am I looking for?
- Key (K): What information do I have?
- Value (V): What should I retrieve?
- Attention = $\text{softmax}(QK'/\sqrt{d}) * V$

Combined heads provide rich, multi-faceted understanding.

Appendix A5: BERT Technical Specifications

BERT-Base Architecture:

- 12 transformer layers
- 768 hidden dimensions
- 12 attention heads
- 110 million parameters
- 512 maximum sequence length

BERT-Large Architecture:

- 24 transformer layers
- 1024 hidden dimensions
- 16 attention heads
- 340 million parameters
- 512 maximum sequence length

Training Data:

- Wikipedia: 2.5B words
- BookCorpus: 800M words
- Total: 3.3B words

Appendix A6: BERT Pre-training Tasks

1. Masked Language Model (MLM):

- Randomly mask 15% of tokens
- Predict masked words from context
- Example: “The [MASK] was delicious” → “food”
- Forces bidirectional understanding

2. Next Sentence Prediction (NSP):

- Given two sentences, are they consecutive?
- 50% actual next sentences
- 50% random sentences
- Learns discourse relationships

These tasks teach BERT language structure without labels.

Appendix A7: Fine-tuning for Specific Tasks

Transfer Learning Process:

1. Start with pre-trained BERT
2. Add task-specific head (classification layer)
3. Train on labeled data (much smaller dataset)
4. Fine-tune all parameters (or freeze lower layers)

Common Fine-tuning Tasks:

- Sentiment Analysis: Add binary classifier
- Named Entity Recognition: Token classification
- Question Answering: Span prediction
- Text Similarity: Sentence pair classification

Typical Data Requirements:

- Minimum: 1,000 examples
- Good: 10,000 examples
- Excellent: 100,000+ examples

Appendix A8: BERT vs Other Models

Model	Direction	Use Case	Params
BERT	Bidirectional	Understanding	110M
GPT-2	Left-to-right	Generation	1.5B
RoBERTa	Bidirectional	Better BERT	355M
ALBERT	Bidirectional	Efficient BERT	12M
XLNet	Permutation	Best of both	340M

Key Differences:

- GPT: Autoregressive (good for generation)
- BERT: Autoencoding (good for understanding)
- RoBERTa: BERT with more data, no NSP
- ALBERT: Parameter sharing for efficiency

Appendix A9: Emotion Classification Systems

Plutchik's Wheel of Emotions:

- 8 primary emotions
- 3 intensity levels each
- Opposite pairs (joy-sadness, trust-disgust)
- Complex emotions as combinations

Ekman's Basic Emotions:

- Anger, Disgust, Fear
- Happiness, Sadness, Surprise
- Universal across cultures

For Product Reviews:

- Satisfaction/Dissatisfaction
- Delight/Frustration
- Trust/Skepticism
- Excitement/Disappointment

Appendix A10: Simple BERT Implementation

Python Code Example:

```
from transformers import pipeline

# Load pre-trained BERT for sentiment
analyzer = pipeline("sentiment-analysis")

# Analyze text
text = "This product is not bad at all"
result = analyzer(text)

# Output:  [{'label':  'POSITIVE', 'score':  0.82}]

# Fine-tuning example
from transformers import BertForSequenceClassification
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased", num_labels=2)
```

Full implementation available in course repository.