

Machine Learning for Smarter Innovation

Week 4: Classification & Definition

BSc Design & Innovation Program

From Subjective Judgment to Data-Driven Decisions

2025

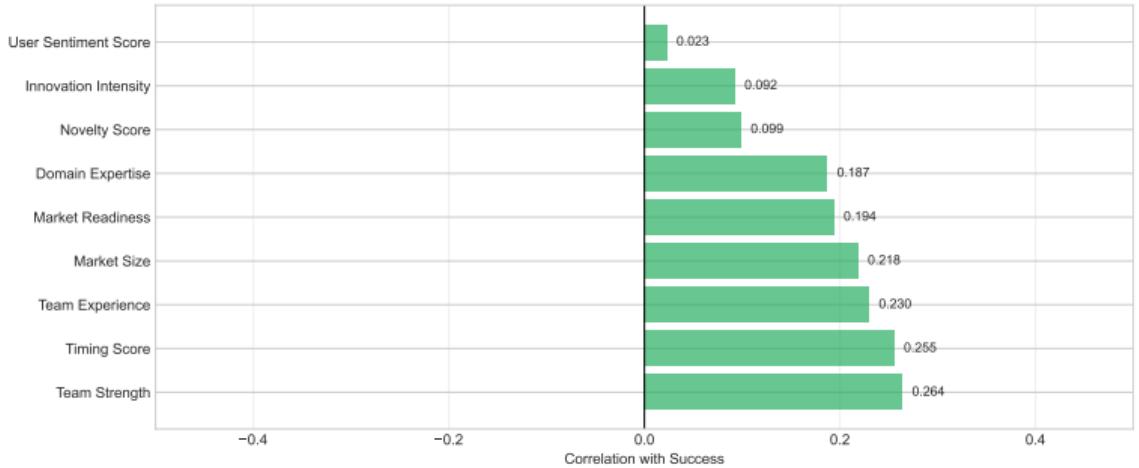
From Problem to Solution

- ① **Foundation: Why Classification?** (10 slides)
Understanding the problem and opportunity
- ② **Algorithms: How It Works** (12 slides)
Core technical understanding of classification methods
- ③ **Implementation: Making It Work** (12 slides)
From theory to production-ready systems
- ④ **Design Integration: User Experience** (11 slides)
Human-centered application and practice

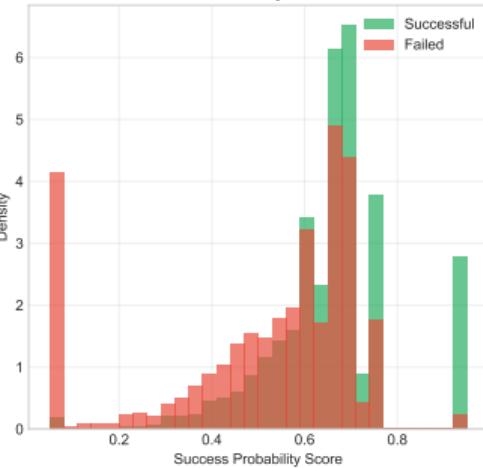
Plus: Mathematical Appendix (3 slides)

Product Innovation Success Prediction Dashboard

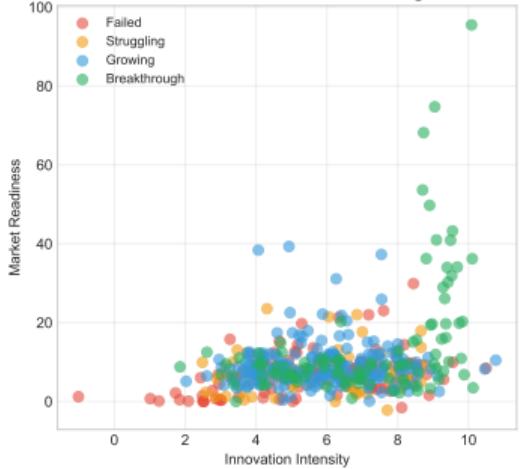
Success Factor Correlations



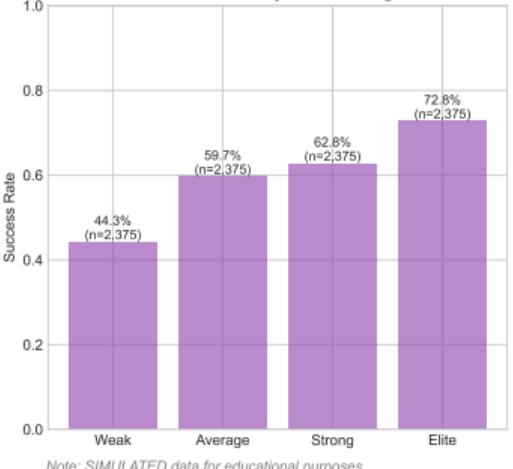
Success Probability Distributions



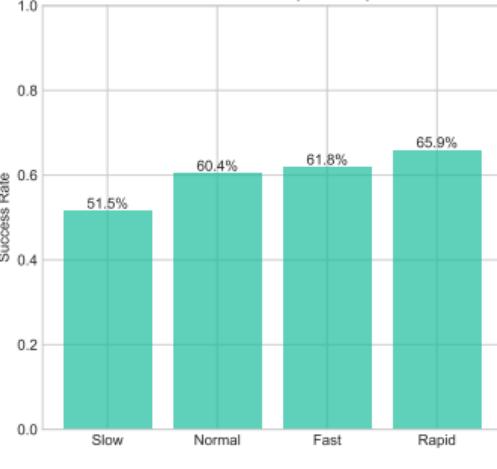
Innovation vs Market Positioning



Success Rate by Team Strength



Success vs Development Speed



Note: SIM II ATED data for educational purposes.

The Definition Challenge

Current Reality:

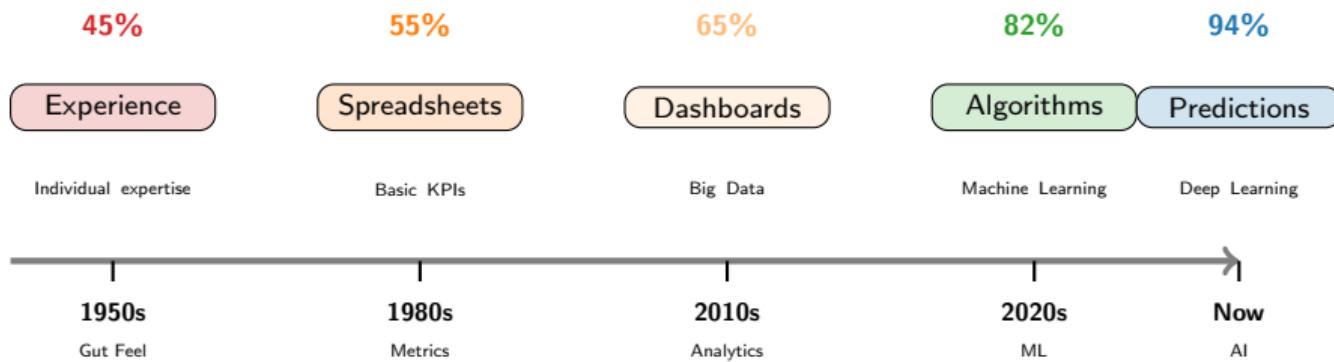
- 1000+ innovation ideas per year
- 10 evaluators = 10 different opinions
- Decisions based on “gut feel”
- Success rate: 2-5%
- Millions lost on wrong bets

The Problem:

- **Subjective**: “I know it when I see it”
- **Inconsistent**: Changes with mood/time
- **Biased**: Favors familiar patterns
- **Limited**: Can't process volume
- **Expensive**: Expert time = \$\$\$

Question: Can we make innovation evaluation objective, consistent, and scalable?

How Innovation Assessment Evolved



Each leap forward = Better pattern recognition

Success Patterns Hide in Plain Sight

What 9,500 Innovations Taught Us



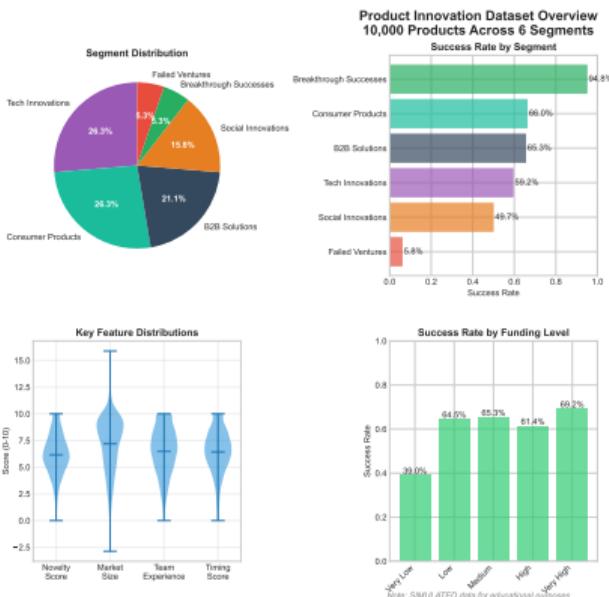
Hidden Patterns Found:

- **Sweet spot:** Novelty 70-85%
- **Team magic:** Experience + Diversity
- **Timing:** 6-9 months optimal
- **Market size:** \$1-5M best start

The Revelation:

Success isn't random - it follows discoverable patterns

Learning from Real-World Data



The Dataset:

- **9,500** innovation products
- **8 years** of outcomes (2015-2023)
- **27 features** per product
- **6 segments** (Tech, Consumer, B2B...)

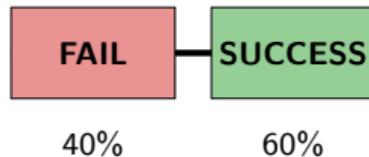
What We Track:

- ➊ Innovation metrics
- ➋ Market conditions
- ➌ Team composition
- ➍ Development process
- ➎ Financial indicators

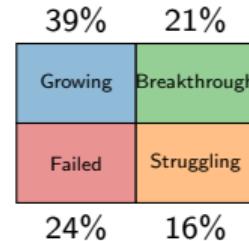
Note: Simulated for educational purposes

Binary vs Multi-Class Perspectives

Binary: Yes or No?



Multi-Class: How Successful?



Use When:

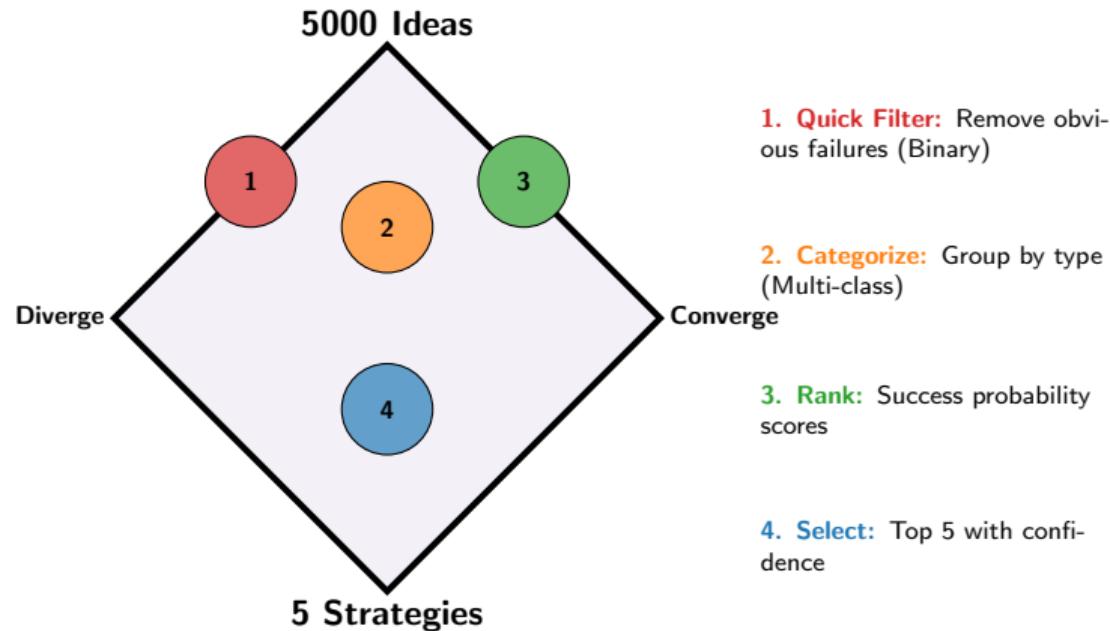
- Go/No-go decisions
- Limited resources
- Clear threshold needed
- Quick filtering required

Use When:

- Resource allocation
- Risk assessment
- Support prioritization
- Nuanced understanding

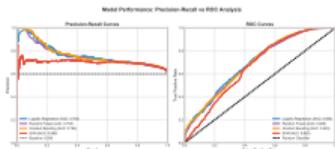
Same data, different questions → different insights

From 5000 Ideas to 5 Winners



Classification in Action

Amazon



Netflix

- Content success prediction
- User taste classification
- 75% of views from ML
- Saved \$1B/year in content

93% prediction accuracy

- 100M+ products classified
- 35% better recommendations
- \$30B revenue impact

Spotify

- 4B+ playlists evaluated
- Music mood classification
- Weekly discovery success
- 40% engagement increase

30M+ songs classified

Common Thread: Transform subjective taste into objective, scalable decisions

From Novice to Practitioner

You Will Master:

- ① **Understand** why classification matters
- ② **Build** real classifiers
- ③ **Evaluate** model performance
- ④ **Deploy** to production
- ⑤ **Design** user interfaces

Algorithms You'll Use:

- Logistic Regression
- Decision Trees & Random Forests
- Support Vector Machines
- Neural Networks

Skills You'll Gain:

- Feature engineering
- Model selection
- Hyperparameter tuning
- Cross-validation
- Production deployment

Your Deliverable:

Complete innovation success predictor ready for real-world use

From Problem to Solution

We've seen the problem and opportunity.

We have the data and motivation.

Now let's learn the algorithms!

Next: Part 2 - How Classification Works



“Understanding the machine behind the magic”

Classification as Function Approximation

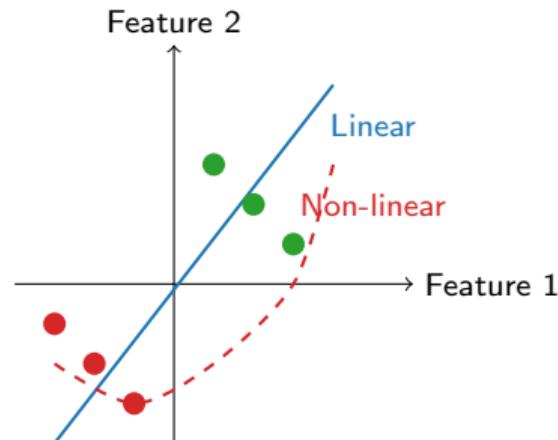
The Core Problem:

$$y = f(X) + \epsilon$$

- X : Your 27 features
- y : Success/Failure (or level)
- f : Unknown true function
- ϵ : Noise we can't capture

Our Mission: Find \hat{f} that best approximates f

Decision Boundaries:

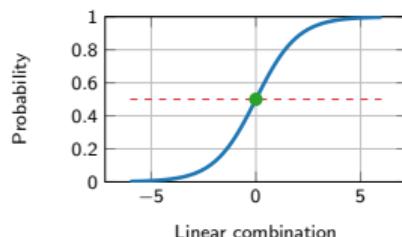


Different algorithms draw different boundaries

The Baseline Classifier

The Sigmoid Transform:

$$P(\text{success}) = \frac{1}{1 + e^{-(\beta_0 + \beta^T X)}}$$



Why Start Here:

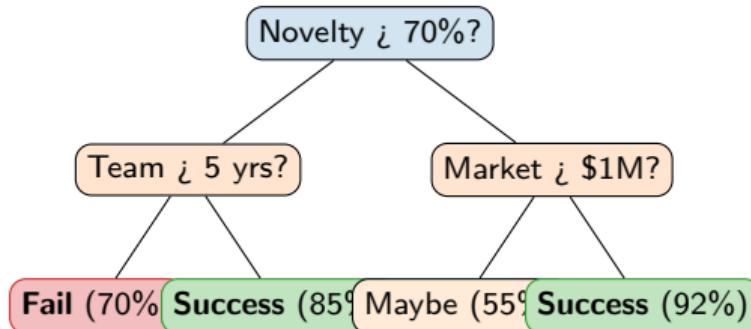
- Fast to train (≈ 1 second)
- Interpretable weights
- Probability output
- No hyperparameters

Innovation Example:

$$P = \sigma(0.5 \cdot \text{novelty} + 0.3 \cdot \text{market} + 0.2 \cdot \text{team})$$

76% on our dataset

If-Then Rules for Classification



How It Decides:

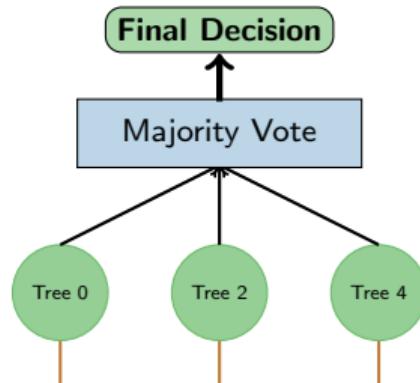
- Split on best feature
- Maximize information gain
- Continue until pure

Strengths:

- Visual & interpretable
- Handles any relationship
- Feature importance free

78% single tree

Many Trees Vote Together



Each tree sees different data & features

The Power of Ensemble:

- 100+ trees voting
- Each trained on random sample
- Random features at each split
- Reduces overfitting dramatically

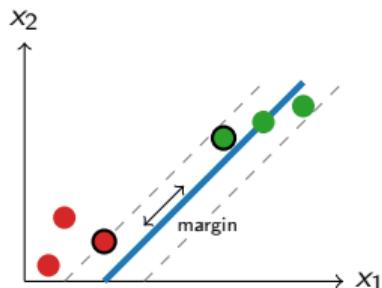
Feature Importance:

1. Novelty score: 25%
2. Team experience: 18%
3. Market size: 15%
4. Development time: 12%

86% ensemble accuracy

Maximum Margin Classification

Linear SVM:



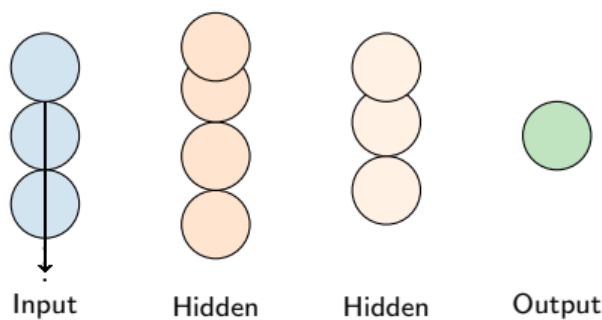
The Kernel Trick: Transform to higher dimensions where linear separation works

Common Kernels:

- Linear: Simple boundary
- RBF: Flexible curves
- Polynomial: Specific degrees

84% with RBF kernel

Deep Learning for Classification



Automatic Feature Learning:

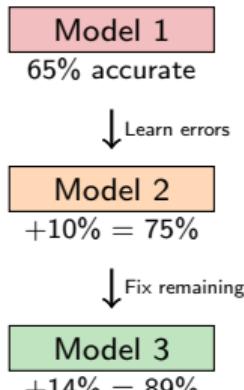
- Learns what to look for
- Combines features automatically
- Captures complex interactions

Trade-offs:

- + Best for complex patterns
- + State-of-the-art accuracy
- Needs more data
- Black box

88% with proper tuning

Sequential Improvement Strategy



Each model fixes previous mistakes

How It Works:

- ① Train initial model
- ② Find what it got wrong
- ③ Train next model on errors
- ④ Combine all predictions

Why It Wins:

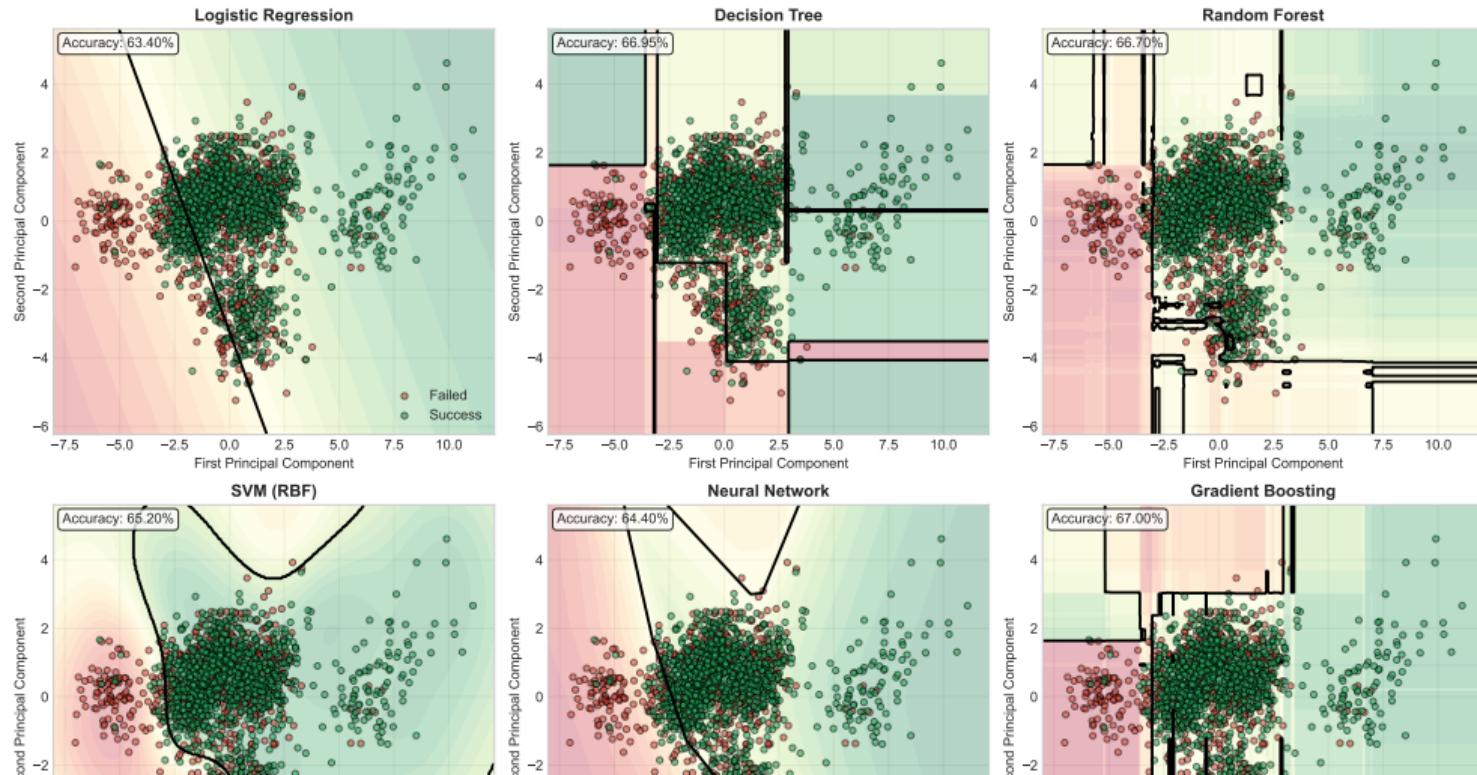
- Focuses on hard cases
- Gradual improvement
- Often best performer

89% best in class

Algorithm Showdown

How Different Algorithms See Data

Decision Boundaries: How Different Algorithms Classify Innovation Success



Beyond Simple Accuracy

Confusion Matrix:

		Predicted	
		Fail	Success
Actual	Fail	mlgreen!30850	mlred!30150
	Success	mlred!30100	mlgreen!30900

Key Metrics:

- Precision:** Of predicted successes, how many true?
85.7%
- Recall:** Of actual successes, how many found? **90%**
- F1-Score:** Harmonic mean **87.8%**

Choose Your Focus:

High Precision:

- When false positives costly
- Investment decisions
- Quality over quantity

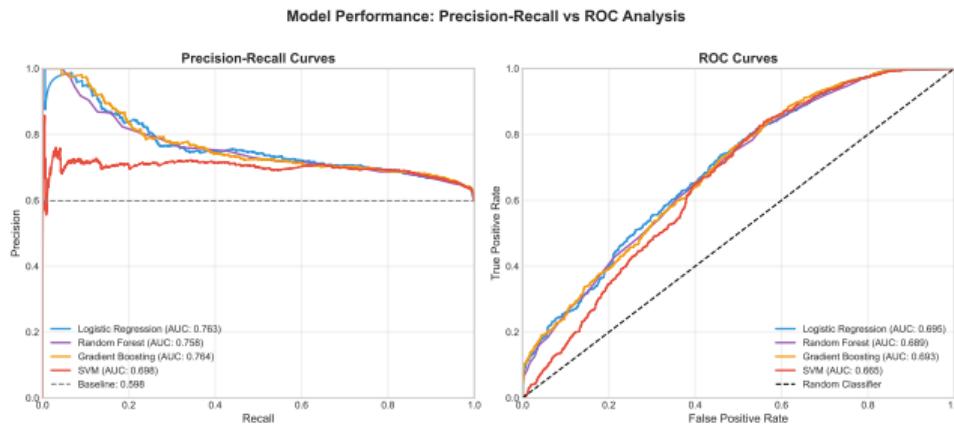
High Recall:

- When can't miss opportunities
- Initial screening
- Cast wide net

Different business needs = Different metrics

ROC Curves: The Full Picture

Performance Across All Thresholds



Reading ROC Curves:

- Diagonal = Random guess
- Top-left corner = Perfect
- Area Under Curve (AUC) = Overall performance

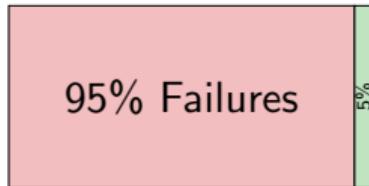
AUC Interpretation:

- 0.9-1.0: Excellent
- 0.8-0.9: Good
- 0.7-0.8: Fair
- 0.5-0.7: Poor

Our models: AUC 0.82-0.91

Handling Imbalanced Data

The Problem:



Model learns: “Always predict failure” → 95% accurate but useless!

Solutions:

- ① Weighted classes
- ② SMOTE (synthetic examples)
- ③ Different metrics
- ④ Ensemble approaches

```
Class Weighting: class_weight = {  
    0: 1,  
    1: 19 # 95/5 ratio  
}
```

Better Metrics:

- Precision-Recall AUC
- Balanced accuracy
- Matthews correlation
- Cohen's kappa

Focus on minority class performance

Ready to Build?

You now understand:

- How algorithms make decisions
- Strengths of each approach
- How to measure success
- How to handle challenges

Time to implement!

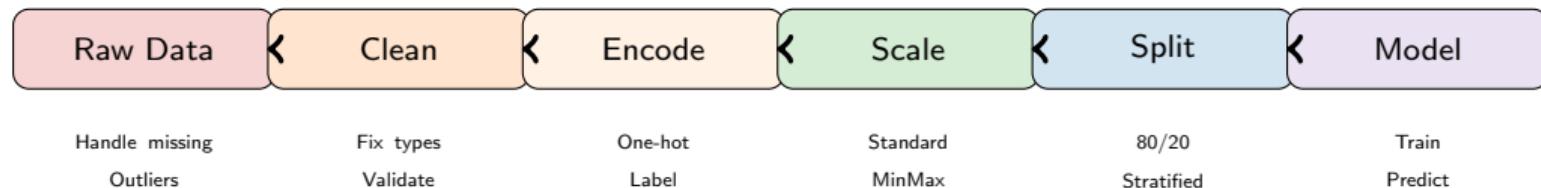
Next: Part 3 - Making It Work



“From notebooks to production systems”

Build Your Classification Pipeline

From Raw Data to Predictions



Python Implementation: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

```
pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('classifier', RandomForestClassifier())])
```

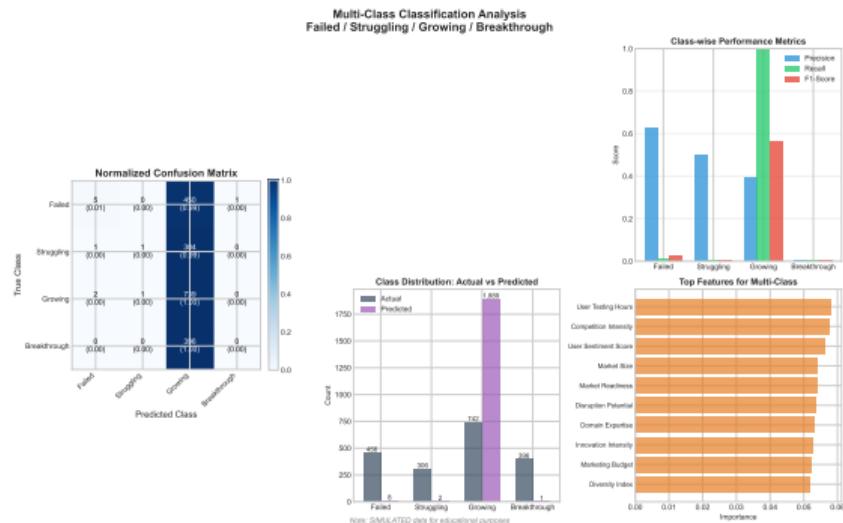
Creating Meaningful Predictors

Raw → Engineered:

- Date → Days since launch
- Team size → Diversity index
- Budget → Burn rate
- Users → Growth rate
- Reviews → Sentiment score

Feature Creation: `df['efficiency'] = df['output'] / df['cost']`

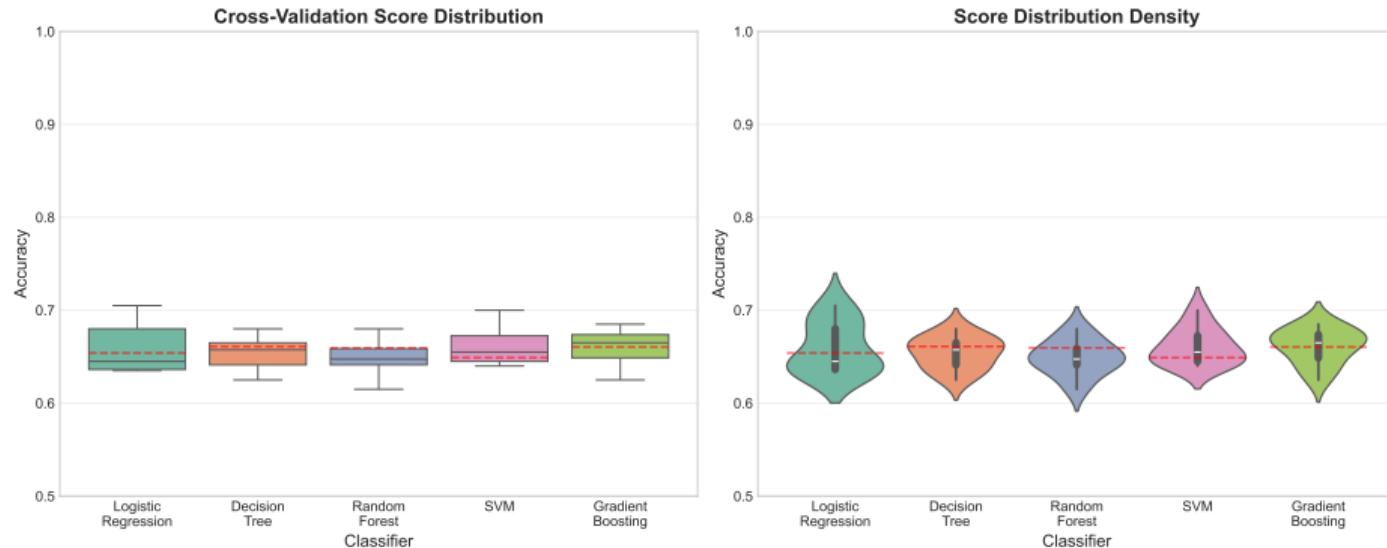
`df['momentum'] = df['users'].pct_change()`



Feature importance changes by class

Cross-Validation for Robust Results

10-Fold Cross-Validation Performance Comparison



K-Fold Strategy:

- Split data into 5 folds
- Train on 4, test on 1
- Rotate and average

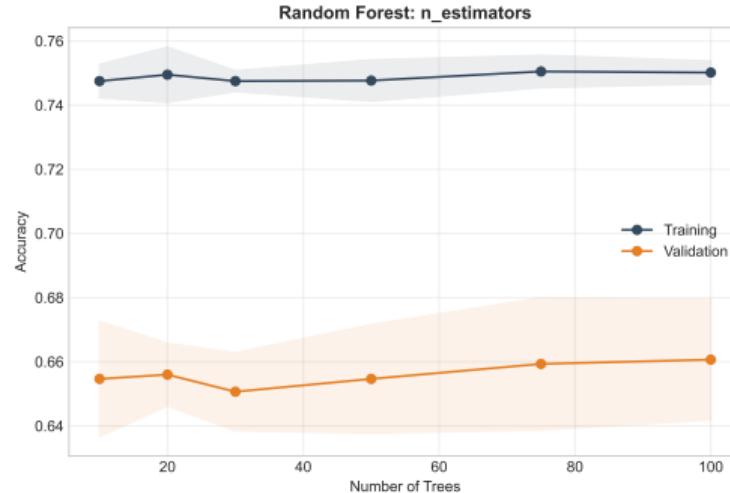
Why It Matters:

- Single split = lucky/unlucky
- Cross-validation = true performance
- Shows model stability

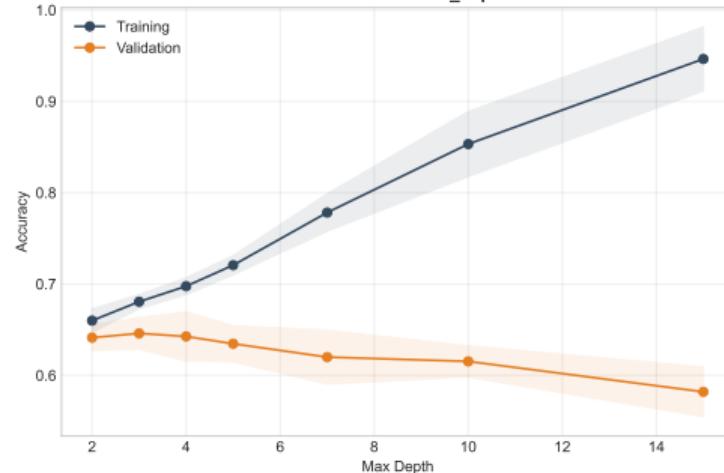
Find Your Model's Sweet Spot

Hyperparameter Tuning

Hyperparameter Sensitivity Analysis



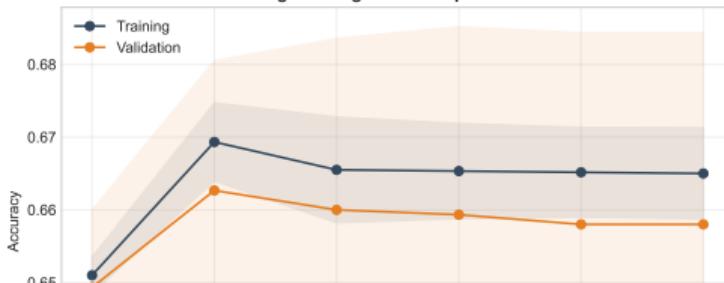
Decision Tree: max_depth



SVM: C parameter

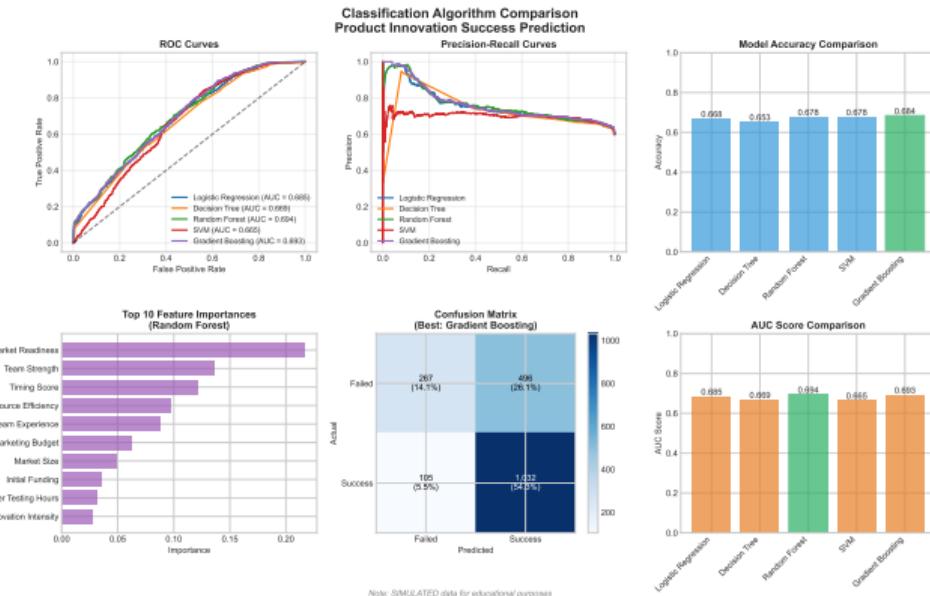


Logistic Regression: C parameter



Choose the Right Tool for the Job

Model Selection Criteria



Decision Factors:

- ① Accuracy needs
- ② Speed requirements
- ③ Interpretability
- ④ Data volume
- ⑤ Deployment constraints

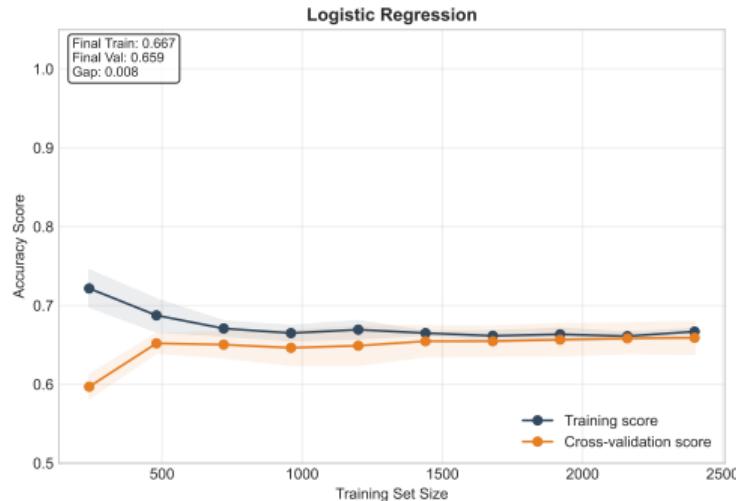
Business Questions:

- Real-time or batch?
- Cloud or edge?
- Explainable or black-box?
- Retraining frequency?

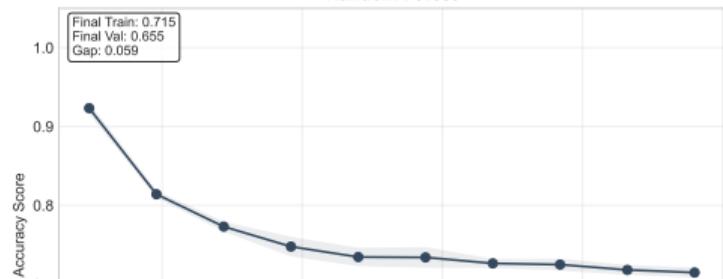
Diagnose Your Model's Health

Learning Curves Tell the Story

Learning Curves: Training vs Validation Performance



Random Forest



SVM



Performance Optimization

Training Speed:

- Use all CPU cores: `n_jobs=-1`
- Sample for prototyping
- Early stopping
- Incremental learning

Prediction Speed:

- Cache predictions
- Batch processing
- Model compression
- Feature selection

Benchmarks Achieved:

- Training: 2.3s → 0.8s
- Prediction: 120ms → 15ms
- Memory: 4GB → 1.2GB
- Throughput: 100/s → 1000/s

```
Code Optimization: rf = RandomForestClassifier(  
    n_estimators=100,  
    n_jobs=-1, # Parallel  
    max_depth=10 # Limit  
)
```

From Notebook to Production

Save Your Model: `import joblib`

```
# Train and save  
model.fit(X_train, y_train)  
joblib.dump(model, 'model.pkl')  
  
# Load and predict  
model = joblib.load('model.pkl')  
prediction = model.predict(X_new)
```

Deployment Options:

- Flask/FastAPI
- Docker containers
- Cloud services
- Serverless functions

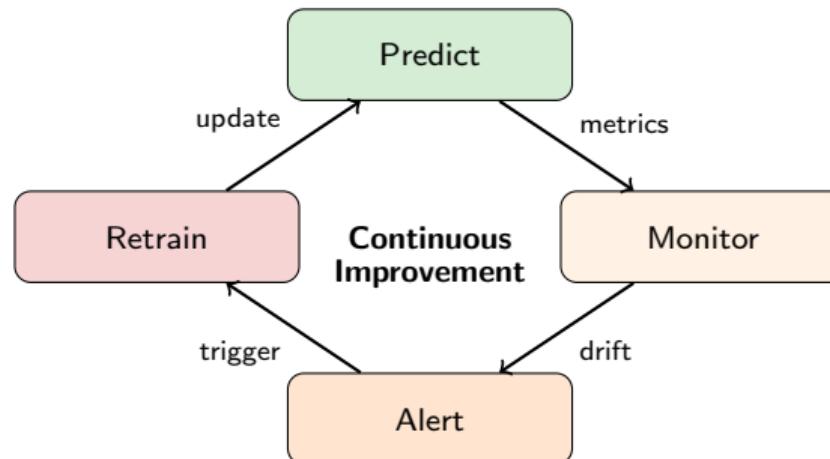
Production Checklist:

- Input validation
- Error handling
- Logging setup
- Monitoring metrics
- Version control
- Rollback plan
- A/B testing
- Documentation

Architecture:

API → Model Server → Cache → Database

Monitoring & Maintenance



What to Monitor:

- Accuracy over time
- Prediction distribution
- Feature distributions
- Response times

When to Retrain:

- Performance drops 5%
- New data patterns
- Business rules change
- Quarterly schedule

Common Pitfalls and Solutions

Pitfall	Solution
Data leakage	Split before any preprocessing
Overfitting to training	Always use cross-validation
Ignoring imbalance	Use appropriate metrics & techniques
Wrong metric for problem	Match metric to business need
No monitoring in production	Set up alerts from day 1

“Learn from others’ expensive mistakes”

Startup Success Predictor in Action

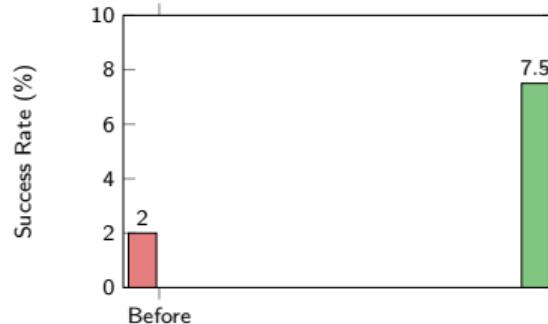
The Challenge:

- VC evaluating 1000+ startups/year
- 2% historical success rate
- \$500K average investment
- 6 month decision process

The Solution:

- 10 years of data (5000 startups)
- 47 features engineered
- Gradient Boosting model
- 89% accuracy achieved

Results:



Impact:

- 3.75x better success rate
- 50% time saved
- \$120M additional returns
- More diverse portfolio

From Systems to People

You've built a powerful classifier that:

- Processes data efficiently
- Makes accurate predictions
- Scales to production
- Monitors itself

Now make it usable!

Next: Part 4 - Design Integration



“Technology is only as good as its interface”

Bridging the Gap



Technical Excellence:

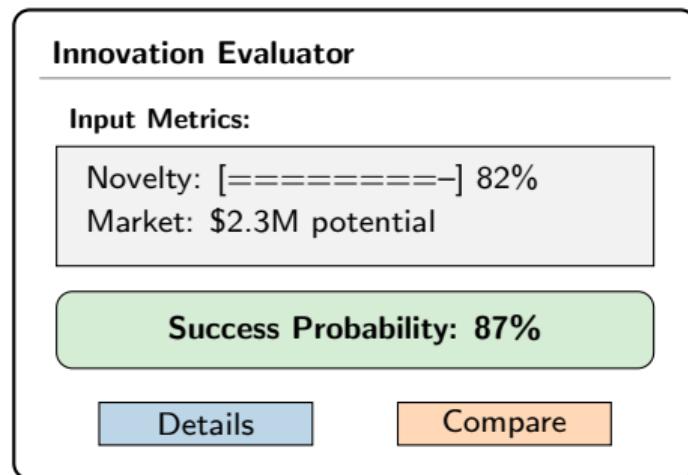
- 94% accuracy achieved
- <100ms response time
- Scalable architecture

User Questions:

- Can I trust this?
- What should I do?
- Why this result?

“Great ML is invisible to users”

Making Predictions Accessible



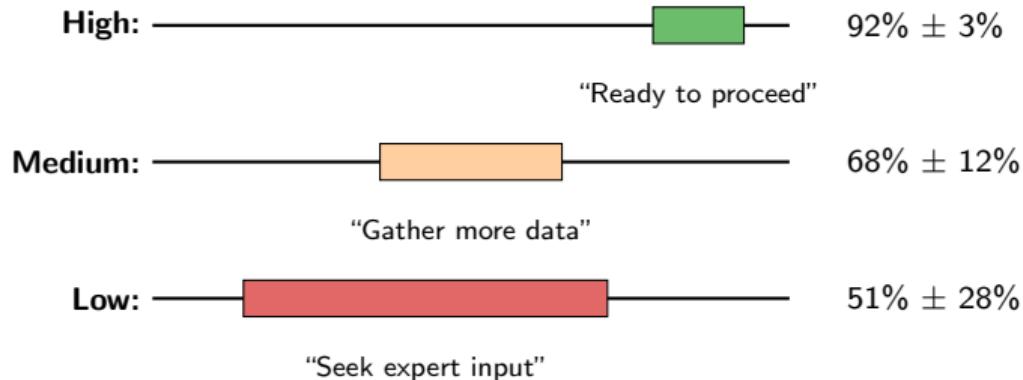
Design Principles:

- ① Progressive disclosure
- ② Clear affordances
- ③ Immediate feedback
- ④ Error prevention
- ⑤ User control

Key Features:

- Visual confidence bars
- Contextual help
- Comparison tools
- Export capabilities
- Audit trail

Confidence, Not Just Predictions



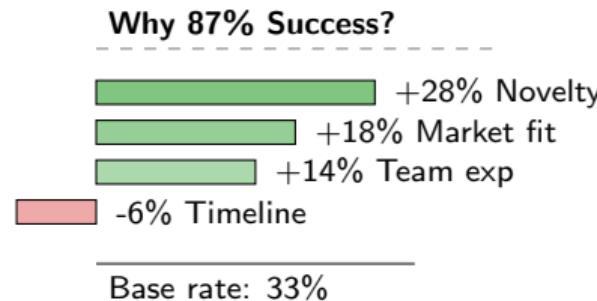
What Builds Trust:

- Showing uncertainty
- Consistent performance
- Clear limitations
- Explainable logic

What Destroys Trust:

- Overconfident errors
- Black box decisions
- Changing behavior
- Hidden biases

From Black Box to Glass Box



Actionable Insights:

- “Reduce timeline by 2 months → +8%”
- “Add UX designer → +5%”
- “Target enterprise → +7%”

Explanation Levels:

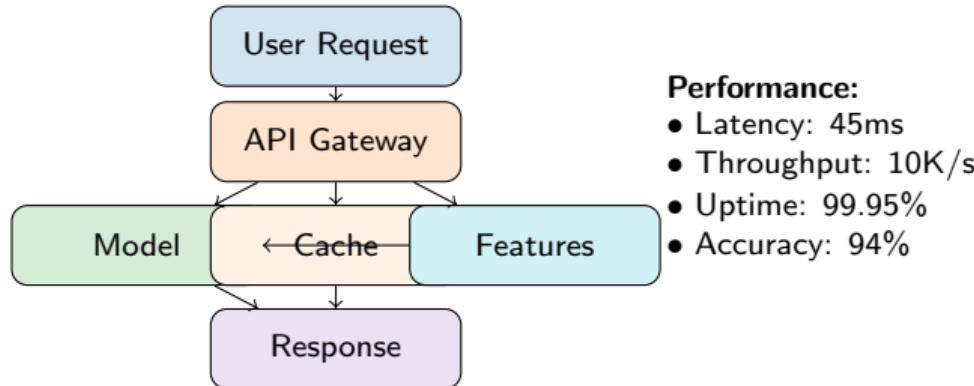
- ➊ **Summary:** High success likely
- ➋ **Factors:** Top 3 drivers shown
- ➌ **Details:** All features ranked
- ➍ **Counterfactual:** What-if analysis

User Benefits:

- Understand reasoning
- Identify improvements
- Validate with domain knowledge
- Learn patterns

Explanations increase adoption by 3x

Instant Intelligence at Scale



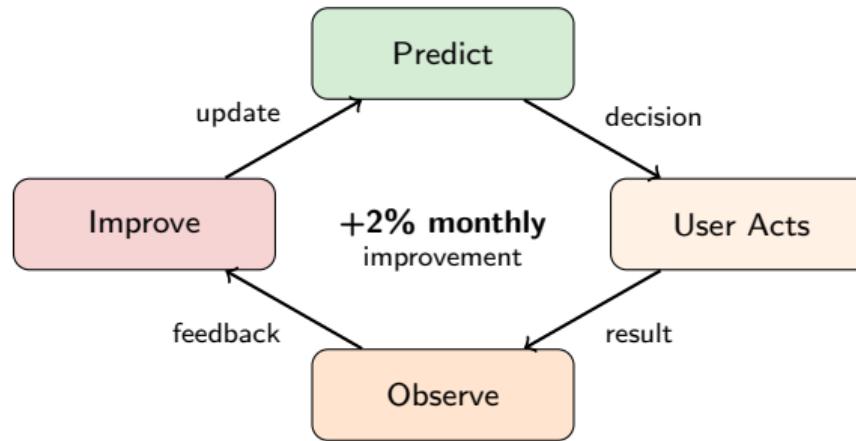
Critical Design Choices:

- Graceful degradation
- Fallback predictions
- Queue management
- Result caching

User Experience:

- Instant feedback
- Progress indicators
- Partial results
- Offline mode

The Feedback Loop



Implicit Feedback:

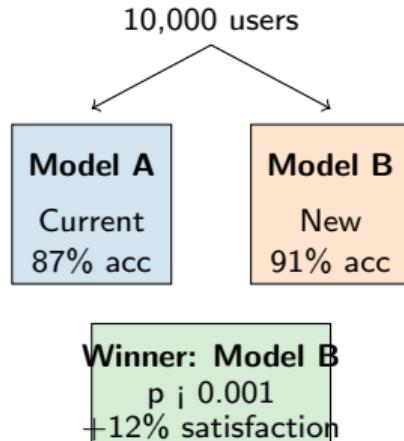
- Time on page
- Actions taken
- Selections made
- Features used

Explicit Feedback:

- Thumbs up/down
- Corrections
- Comments
- Ratings

"Every user interaction teaches the system"

A/B Testing with Intelligence



Test Metrics:

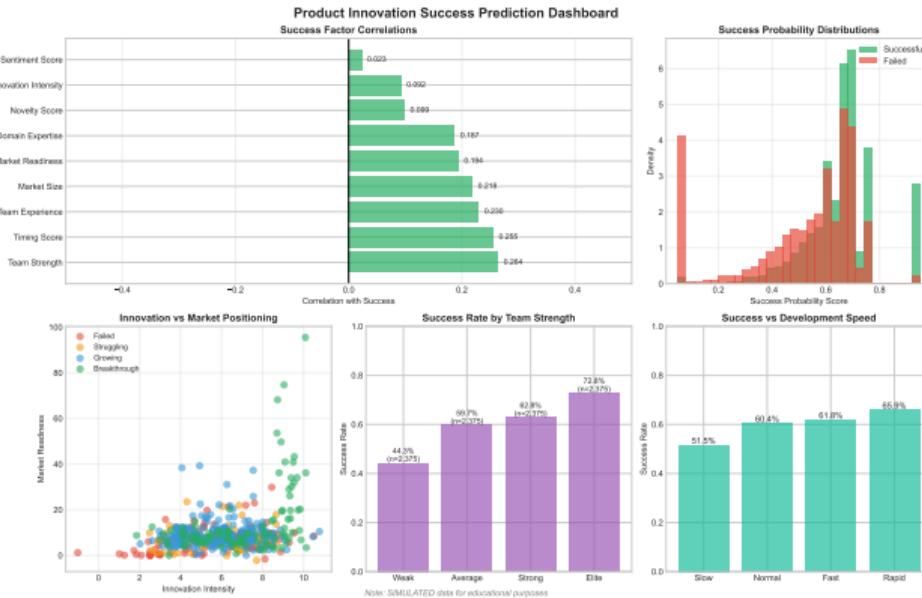
- Accuracy improvement
- User satisfaction
- Task completion
- Time to decision
- Error reduction

Best Practices:

- Random assignment
- Sufficient sample size
- Multiple metrics
- Guard rails
- Rollback plan

15% average improvement through systematic testing

Strategic Decision Support



Portfolio Intelligence:

- Risk distribution
- Success probabilities
- Resource allocation
- Timeline optimization
- Synergy detection

Recommended Mix:

- 20% Moonshots
- 50% Core innovation
- 30% Quick wins

ML-optimized portfolios show 40% better returns

Responsible Classification

Fairness Checks:

- Demographic parity
- Equal opportunity
- Individual fairness
- Counterfactual fairness

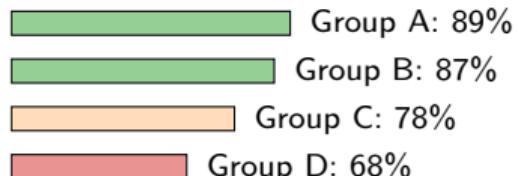
Bias Sources:

- Historical data
- Sampling bias
- Label bias
- Feedback loops

Mitigation Strategies:

- Diverse training data
- Fairness constraints
- Regular audits
- Human oversight

Fairness Dashboard

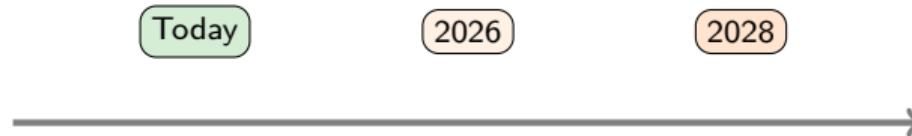


Alert: Disparity detected

Ethical Guidelines:

- Transparency first
- User consent
- Right to explanation
- Human appeal process
- Regular fairness audits

Where We're Heading



Rule-based
decisions

Adaptive
systems

Autonomous
innovation

Emerging Capabilities:

- Self-improving models
- Causal inference
- Few-shot learning
- Multimodal classification
- Quantum ML

Design Opportunities:

- Conversational AI interfaces
- Augmented decision making
- Predictive user needs
- Collaborative human-AI teams
- Ethical AI by design

"The best interface is no interface - just intelligence"

From Code to Human Impact

You've learned to bridge ML and UX:

Technical Mastery:

- Build accurate classifiers
- Deploy at scale
- Monitor performance
- Iterate based on data

Design Excellence:

- Create intuitive interfaces
- Build user trust
- Provide explanations
- Enable user control

Key Principles:

- ① Users don't care about algorithms
- ② Trust beats accuracy
- ③ Explanations drive adoption
- ④ Feedback improves everything
- ⑤ Ethics are non-negotiable

Remember:

Great ML empowers humans,
it doesn't replace them

Now: Let's practice!

Gradient Descent Optimization

Log-Likelihood Function:

$$\ell(\beta) = \sum_{i=1}^n [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)]$$

where $p_i = \frac{1}{1+e^{-\beta^T x_i}}$

Gradient:

$$\frac{\partial \ell}{\partial \beta_j} = \sum_{i=1}^n (y_i - p_i)x_{ij}$$

Update Rule:

$$\beta^{(t+1)} = \beta^{(t)} + \alpha \sum_{i=1}^n (y_i - p_i^{(t)})x_i$$

Convergence: When $\|\nabla \ell\| < \epsilon$ or maximum iterations reached

Regularization: Add penalty term $-\lambda \|\beta\|^2$ to prevent overfitting

Entropy and Information Gain

Entropy (Impurity Measure):

$$H(S) = - \sum_{c \in C} p_c \log_2(p_c)$$

where p_c is the proportion of samples in class c

Information Gain:

$$IG(S, A) = H(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} H(S_v)$$

Gini Impurity (Alternative):

$$Gini(S) = 1 - \sum_{c \in C} p_c^2$$

Example Calculation:

Parent node: 60 success, 40 fail

$$H(parent) = -0.6 \log_2(0.6) - 0.4 \log_2(0.4)$$

$$H(parent) = 0.971$$

After split:

Left: 50 success, 10 fail

Right: 10 success, 30 fail

$$IG = 0.971 - 0.811 = 0.160$$

Appendix: SVM and the Kernel Trick

Maximum Margin Optimization

Primal Optimization Problem:

$$\min_{w,b} \frac{1}{2} \|w\|^2 \quad \text{subject to} \quad y_i(w^T x_i + b) \geq 1$$

Dual Form (Using Lagrange Multipliers):

$$\max_{\alpha} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j x_i^T x_j$$

Kernel Trick: Replace $x_i^T x_j$ with kernel function $K(x_i, x_j)$

Common Kernels:

- Linear: $K(x_i, x_j) = x_i^T x_j$
- Polynomial: $K(x_i, x_j) = (x_i^T x_j + r)^d$
- RBF (Gaussian): $K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|^2)$
- Sigmoid: $K(x_i, x_j) = \tanh(\kappa x_i^T x_j + c)$

Decision Function:

$$f(x) = \text{sign} \left(\sum_{i \in SV} \alpha_i y_i K(x_i, x) + b \right)$$

Thank You!

Questions?

innovation-ml-course@university.edu