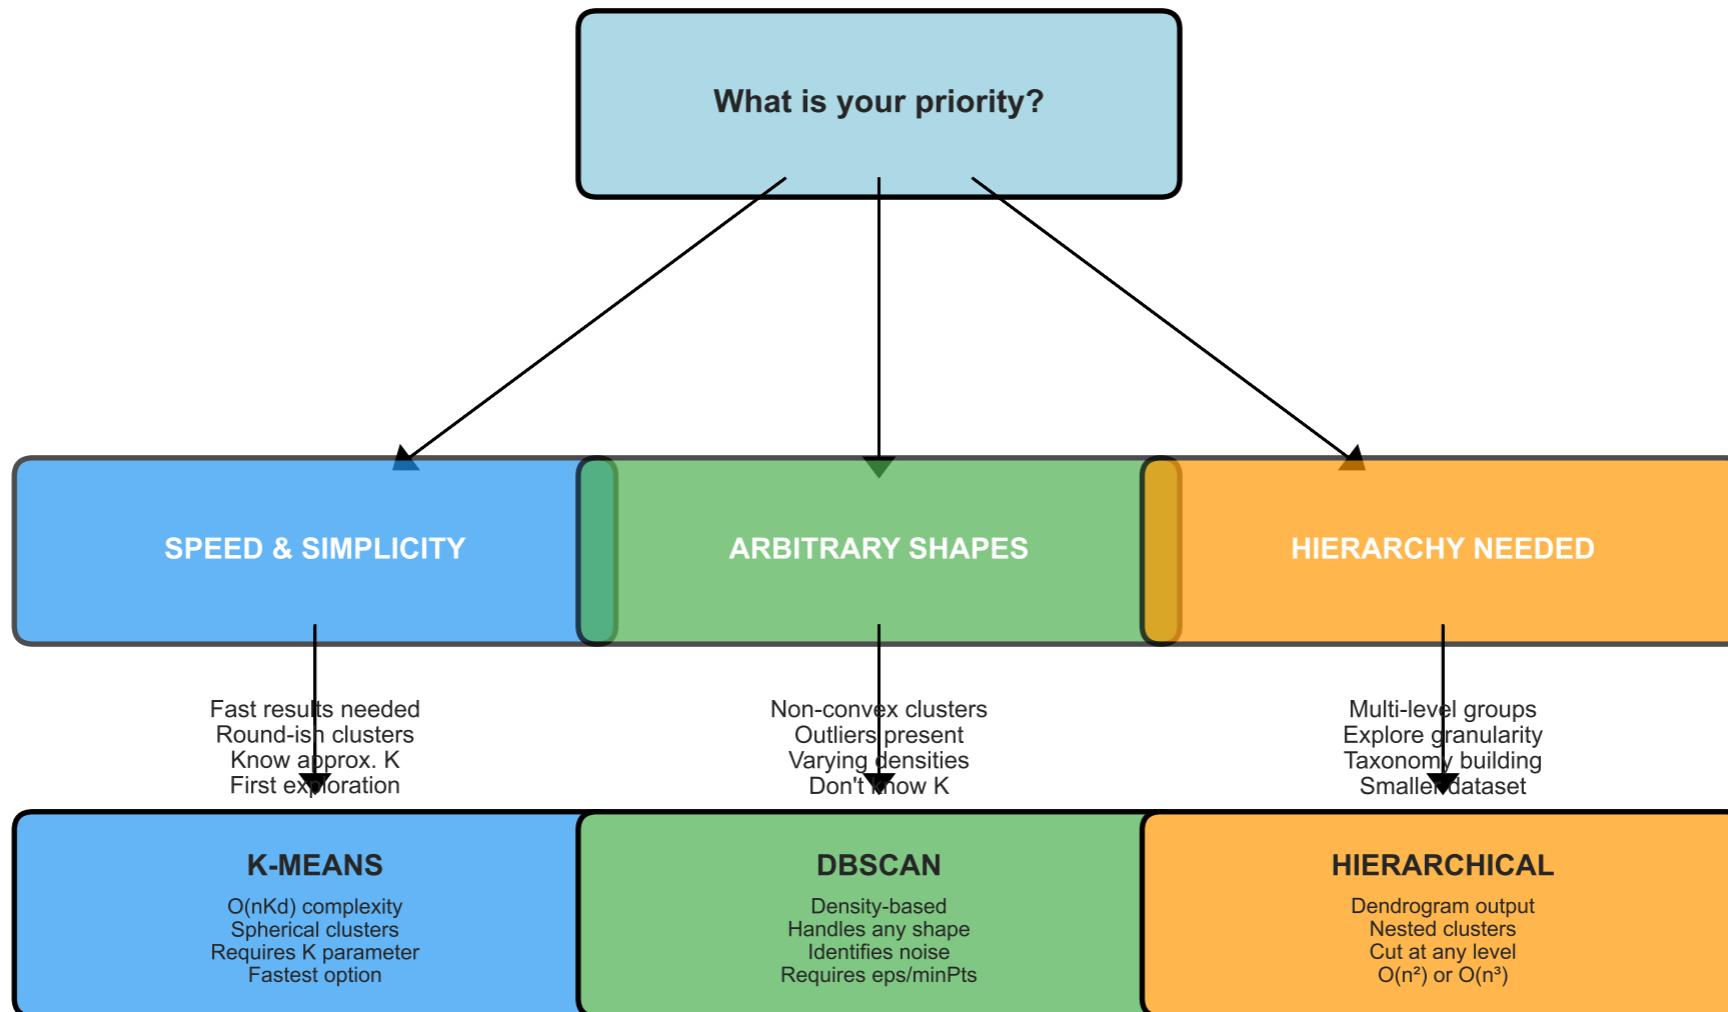


When to Use Which Clustering Algorithm: Decision Framework



Additional Considerations

Dataset Size: Very large ($>100K$ points) → MiniBatch K-means; Small ($<10K$) → Hierarchical feasible
Outliers Critical: Fraud detection, anomaly detection → DBSCAN preferred
Soft Assignments Needed: Mixed populations, uncertainty quantification → GMM (Gaussian Mixture)
High Dimensions: $d>20$ → Curse of dimensionality affects distance; Consider dimensionality reduction first
Reproducibility: Random init sensitivity → Use K-means++ or fixed seed; DBSCAN/Hierarchical deterministic
Production Deployment: Streaming data → BIRCH; Real-time → K-means; Batch → Any algorithm suitable