


Adversarial Debiasing Architecture



Predictor Loss:

$$L_P = -\text{Acc}(Y, \hat{Y})$$

Adversary Loss:

$$L_A = -\text{Acc}(A, \hat{A})$$

Combined:

$$\min_{\theta} \max_{\phi} L_P - \lambda L_A$$