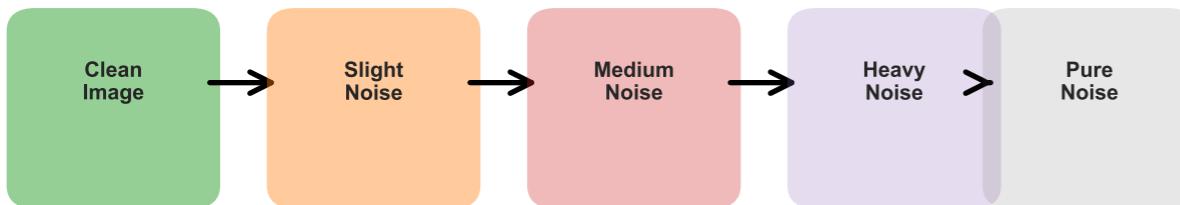


Diffusion: The Reverse Corruption Process

Forward Process: Add Noise

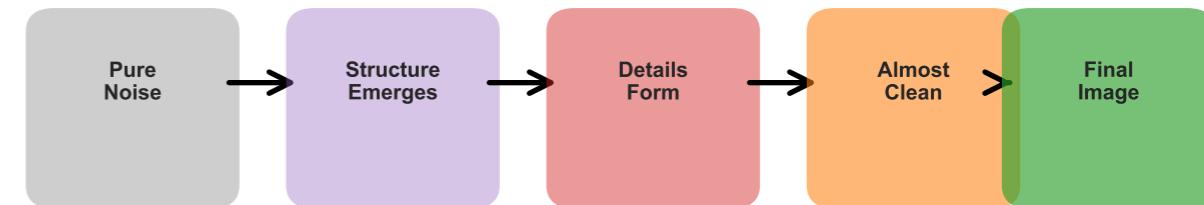


$q(x_t | x_{t-1}) \sim N(\sqrt{1-\beta_t} x_{t-1}, \beta_t I)$

Gradually corrupt clean data

1000 tiny steps

Reverse Process: Remove Noise



$p_\theta(x_{t-1} | x_t)$ - Neural network predicts noise

Learn to reverse corruption

1000 denoising steps