

Machine Learning for Smarter Innovation

Week 1: Clustering for Innovation Discovery

BSc Data Science & AI Program

Innovation & Design Thinking Lab

September 13, 2025

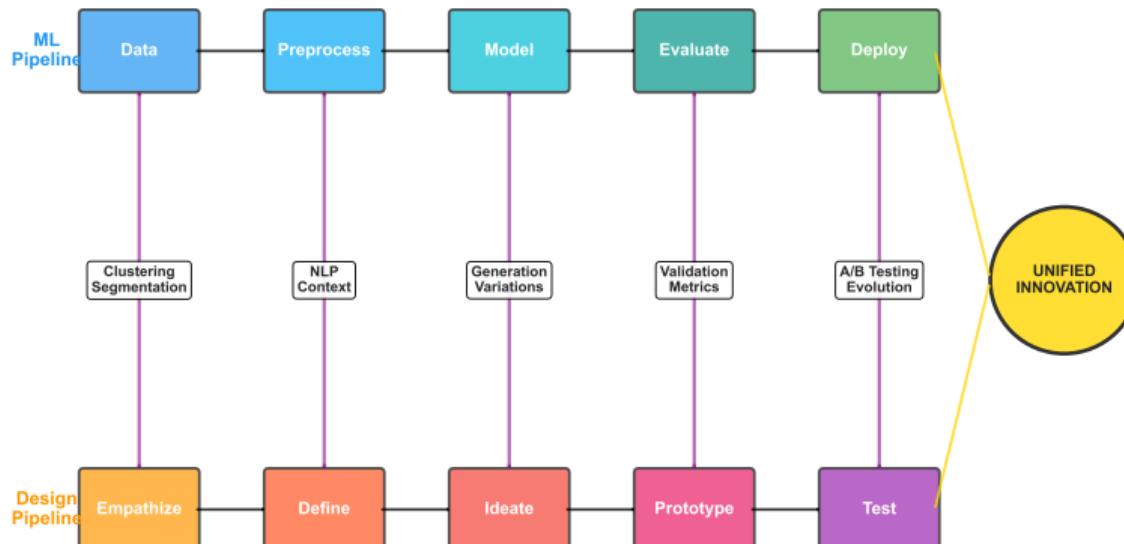
Week 1 Overview: Your Innovation Discovery Journey

From Chaos to Clarity Through Clustering

The Unified Innovation Pipeline

Where Technology Meets Human Creativity

Technical Mastery



Part 1

Foundation & Context

Understanding the Innovation Discovery Challenge

Part 1: Learning Objectives

What You'll Master in This Section

By the end of Part 1, you will:

- ① **Understand** why clustering is essential for innovation discovery
- ② **Identify** the challenges of pattern recognition in innovation data
- ③ **Recognize** how unsupervised learning differs from supervised approaches
- ④ **Connect** clustering to the empathize stage of design thinking

Key Concepts

- Innovation categories vs random ideas
- Pattern discovery in unstructured data
- The curse of dimensionality
- From chaos to actionable insights

Time: 10 minutes

The Innovation Discovery Challenge

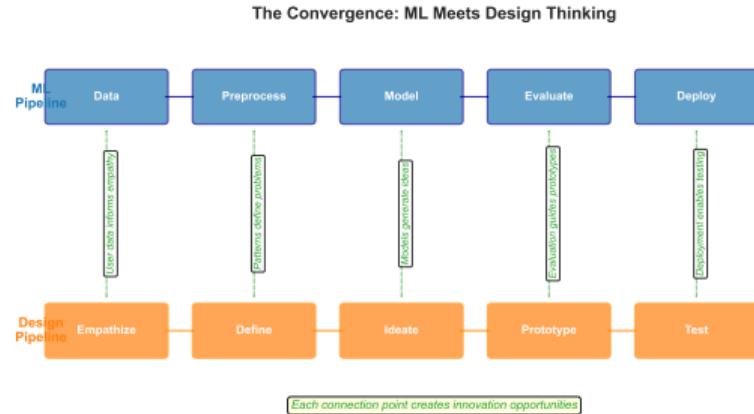
Why We Need Machine Learning

The Problem:

- 1000s of innovation ideas scattered
- No clear categories or patterns
- Hidden connections invisible
- Manual analysis takes months
- Biases cloud human judgment

The Opportunity:

- Discover natural groupings
- Find innovation white spaces
- Identify emerging themes
- Accelerate decision making



What is Clustering?

Finding Order in Innovation Chaos

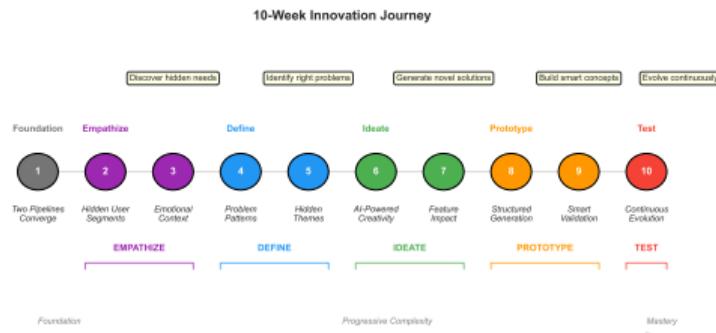
Definition: Unsupervised learning that groups similar items without predefined labels

Innovation Context:

- **Input:** Raw innovation data (ideas, features, feedback)
- **Process:** Algorithm finds natural groupings
- **Output:** Innovation categories and patterns

Key Difference:

- **Supervised:** You know the categories
- **Unsupervised:** You discover the categories



Knowledge Check: Part 1

Test Your Understanding of Foundation Concepts

Interactive Knowledge Checkpoints

Knowledge Check: Part 1

Innovation Discovery Foundation

1 of 3 Parts Complete

Q1: What is the main goal of clustering in innovation?

- A) To reduce data size
- B) To discover hidden patterns
- C) To predict outcomes
- D) To clean data

Q2: Which metric measures cluster cohesion?

- A) Accuracy
- B) Precision
- C) Silhouette Score
- D) F1 Score

Q3: Empathy mapping helps identify:

- A) Technical requirements
- B) User pain points
- C) System architecture
- D) Database schema

Knowledge Check: Part 2

Clustering Algorithms Deep Dive

2 of 3 Parts Complete

Q1: K-means time complexity is:

- A) $O(n)$
- B) $O(n \log n)$
- C) $O(n^k l^d)$
- D) $O(n^2)$

Q2: DBSCAN is best for:

- A) Spherical clusters
- B) Arbitrary shapes
- C) Fixed K clusters
- D) Linear data

Q3: GMM provides:

- A) Hard clustering
- B) Soft clustering
- C) No clustering
- D) Random clustering

Algorithm Quick Reference:

Key Concepts Covered:

• Unsupervised learning

Algorithm	Best For	Weakness
K-Means	Large datasets, spherical clusters	Sensitive to initial centroids, non-converges
DBSCAN	Arbitrary shapes, noisy data	Computational expensive
Gaussian Mixture Model	Soft clustering, complex distributions	Parameter sensitive

Knowledge Check: Part 3

Human-Centered Application

3 of 3 Parts Complete!

Q1: User archetypes are created from:

- A) Random assignment
- B) Cluster analysis
- C) Manual labeling
- D) Predictions

Q2: Innovation opportunities emerge from:

- A) Cluster gaps
- B) Dense regions
- C) Outliers
- D) All of above

Q3: Validation should include:

- A) Only metrics
- B) Domain experts
- C) Random checks
- D) Code review

Ready for Practice!

Part 2

Technical Deep Dive

Mastering Clustering Algorithms

Part 2: Learning Objectives

Technical Skills You'll Develop

By the end of Part 2, you will:

- ① **Master** four core clustering algorithms
- ② **Understand** algorithm complexity and scalability
- ③ **Apply** evaluation metrics effectively
- ④ **Select** the right algorithm for your data
- ⑤ **Optimize** parameters for best results

Algorithms Covered

- **K-Means:** Fast and simple
- **DBSCAN:** Density-based
- **Hierarchical:** Tree structure
- **GMM:** Soft clustering

Time: 20 minutes

Algorithm Complexity & Performance

Understanding Computational Requirements

Clustering Algorithm Complexity & Performance Guide

Algorithm Complexity Analysis

Big O Notation Comparison

Algorithm	Time Complexity	Space Complexity	Scalability
-----------	-----------------	------------------	-------------

K-means $O(n \cdot k \cdot d)$ $O(n \cdot d + k \cdot d)$ Excellent

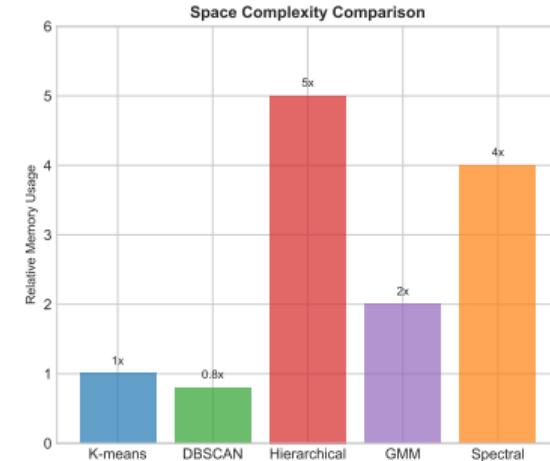
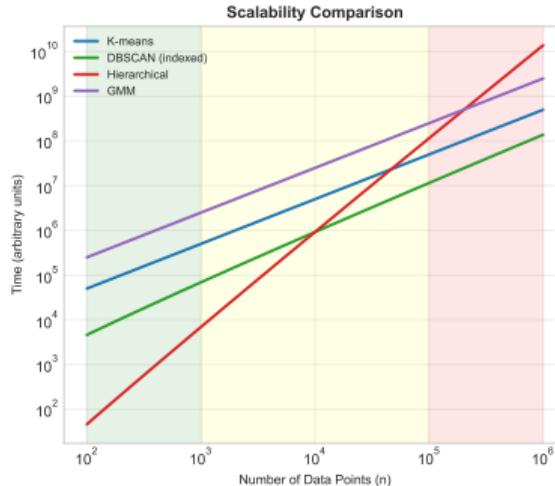
DBSCAN $O(n^2) / O(n \log n)^*$ $O(n)$ Good

Hierarchical $O(n^3) / O(n^2 \log n)^*$ $O(n^2)$ Poor

GMM $O(n \cdot k^2 \cdot i \cdot d)$ $O(k \cdot d^2)$ Moderate

Notation Guide:
n = number of data points
k = number of clusters
i = number of iterations
d = number of dimensions
* = with spatial index

$O(n^3)$ $O(n^2)$ Poor



Practical Recommendations

Small Data (<10K points)

→ Any algorithm works

Medium Data (10K-100K)

→ K-means or DBSCAN

Large Data (>100K)

→ MiniBatch K-means

High Dimensions (>50)

→ Consider PCA first

Real-time Requirements

→ Pre-computed K-means

Optimization Techniques

MiniBatch K-means:

- Samples subset of data
- 10-100x faster on large data

Spatial Indexing (DBSCAN):

- KD-tree or Ball-tree
- $O(n^2) \rightarrow O(n \log n)$

Dimensionality Reduction:

- PCA before clustering
- Reduces d in $O(n \cdot k \cdot d)$

Implementation Complexity

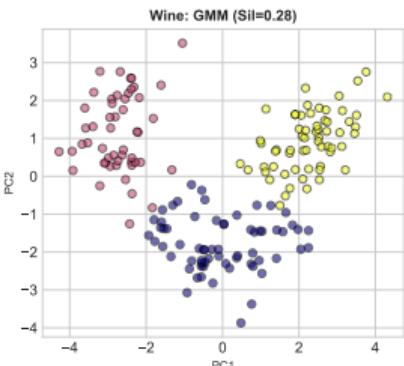
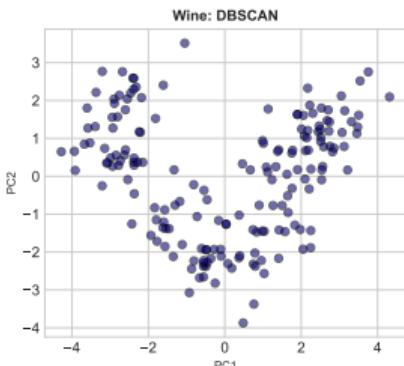
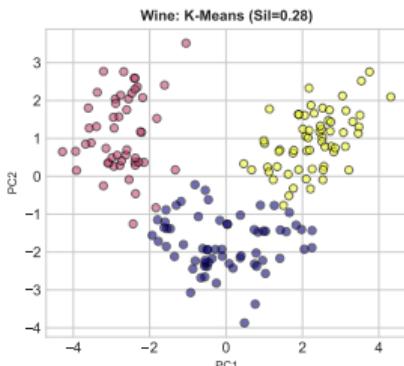
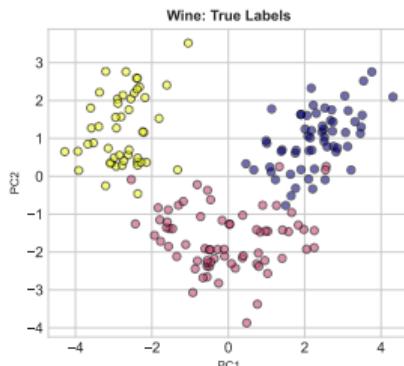
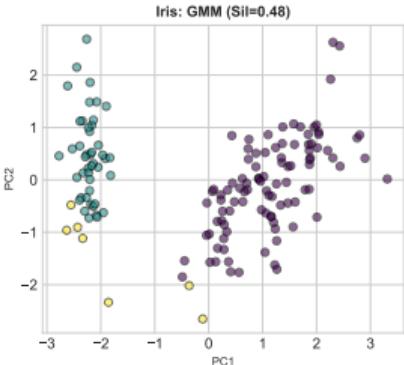
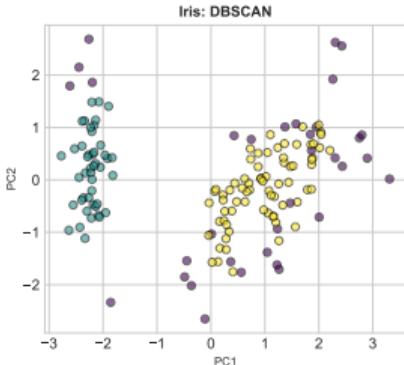
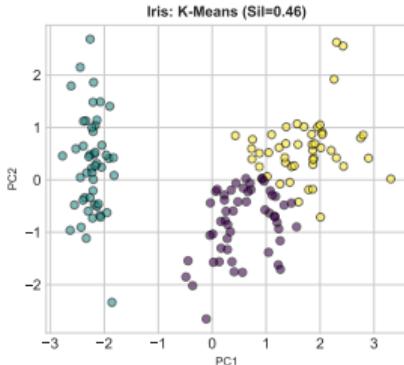
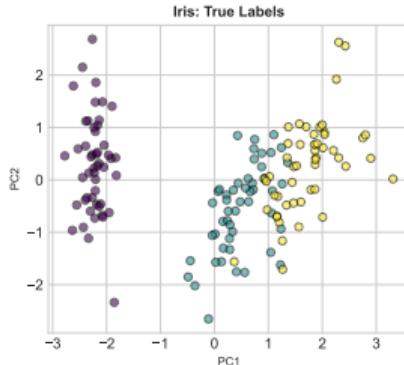
Algorithm	Ease	Lines of Code*	Tuning
K-means	Easy	~50	Simple
DBSCAN	Moderate	~100	Tricky
Hierarchical	Easy	~30	Simple
GMM	Hard	~200	Complex
Spectral	Hard	~150	Complex

Clustering on Real Datasets

Iris, Wine, and Customer Segmentation Examples

Real Dataset Clustering Comparison

Iris (150 samples, 4 features) | Wine (178 samples, 13 features) | Customers (200 samples, 3 features)



Customers: True Segments



Customers: K-Means (Sil=0.56)



Customers: DBSCAN



Customers: GMM (Sil=0.54)



Case Study: Spotify's Music Discovery

How Clustering Powers Personalized Playlists

The Challenge:

- 100+ million songs in catalog
- 500+ million users globally
- Diverse musical tastes
- Need personalized discovery

The Solution:

- **Audio Features:** Extract 13 dimensions
- **Clustering:** Group similar songs
- **User Profiles:** Map listening to clusters
- **Recommendations:** Adjacent clusters

Results:

- 40% increase in discovery
- 2.7B Discover Weekly streams
- 30% longer listening sessions
- 75% user retention

Key Insight:

"Clustering revealed micro-genres users didn't know they loved"

Knowledge Check: Part 2

Test Your Technical Understanding

Algorithm Selection Quiz:

- ① Large dataset (1M points)?
→ K-Means or MiniBatch K-Means
- ② Non-spherical clusters?
→ DBSCAN
- ③ Need probability scores?
→ GMM
- ④ Want dendrogram?
→ Hierarchical

Complexity Check:

- ✓ K-means: $O(n \cdot k \cdot i \cdot d)$
- ✓ DBSCAN: $O(n \log n)$ with index
- ✓ Hierarchical: $O(n^2)$ memory
- ✓ GMM: Soft clustering

Ready for Design Integration?

Let's apply this to innovation!

Part 3

Design Integration

Applying Clustering to Innovation Discovery

Ethical Considerations in Clustering

Responsible AI for Innovation Discovery

Potential Biases:

- **Selection Bias:** Who's included in data?
- **Feature Bias:** What dimensions matter?
- **Algorithmic Bias:** Distance metrics assumptions
- **Interpretation Bias:** Label assignment

Mitigation Strategies:

- Diverse data collection
- Multiple algorithm comparison
- Expert validation
- Transparent documentation

Fair Clustering Checklist:

- Representative sampling?
- Protected attributes removed?
- Cluster balance checked?
- Minority groups visible?
- Results interpretable?
- Decisions reversible?

Remember:

"Clusters are hypotheses, not truth"

Scaling Your Clustering: Cloud & Distributed Options

From Prototype to Production

Local Development:

- Scikit-learn (\downarrow 100K points)
- Jupyter notebooks
- Single machine
- Rapid prototyping

Cloud Platforms:

- **AWS SageMaker:** Built-in algorithms
- **Google Cloud AI:** AutoML clustering
- **Azure ML:** Drag-and-drop interface
- **Databricks:** Spark MLLib integration

Distributed Computing:

Apache Spark MLLib:

- Handles billions of points
- Distributed K-means
- Bisecting K-means
- Gaussian Mixture

Cost-Performance Trade-offs:

- Local: Free but limited
- Cloud: Pay-per-use, scalable
- On-premise: High initial, unlimited use

Knowledge Check: Part 3

Design Integration Mastery

Application Quiz:

- ① Clusters reveal what?
→ Innovation patterns
- ② Validation requires?
→ Domain experts
- ③ Ethical concerns include?
→ Bias & fairness
- ④ Scale with?
→ Cloud platforms

You Can Now:

- ✓ Choose algorithms wisely
- ✓ Apply to real data
- ✓ Consider ethics
- ✓ Scale solutions
- ✓ Extract insights

Next: Practice Time!

Part 4

Summary & Practice

Putting It All Together

Practice Exercise: Innovation Clustering

Your Turn to Discover Patterns

The Challenge: Analyze 500 innovation proposals from a hackathon

Starter Code Template:

```
import pandas as pd
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

# Load data
data = pd.read_csv('innovations.csv')

# Preprocess
scaler = StandardScaler()
X = scaler.fit_transform(data)

# Cluster
kmeans = KMeans(n_clusters=?)
labels = kmeans.fit_predict(X)
```

Your Tasks:

- ① Choose optimal K
- ② Apply clustering
- ③ Evaluate results
- ④ Interpret patterns
- ⑤ Present findings

Resources Provided:

- Jupyter notebook template
- Sample dataset
- Evaluation rubric
- Solution walkthrough

Submit by: Next Week

Week 1: Key Takeaways

Your Innovation Discovery Toolkit

Foundation

- Clustering finds hidden patterns
- Unsupervised learning
- No labels needed
- Connects to empathy stage

Technical

- 4 algorithms mastered
- Complexity understood
- Metrics applied
- Real data processed

Application

- Innovation patterns found
- Ethical considerations
- Scalability options
- Practice exercise ready

You're Ready to Discover Innovation Patterns!

Resources & Next Week

Continue Your Learning Journey

Resources:

Documentation:

- Scikit-learn clustering guide
- Course GitHub repository
- Jupyter notebook templates

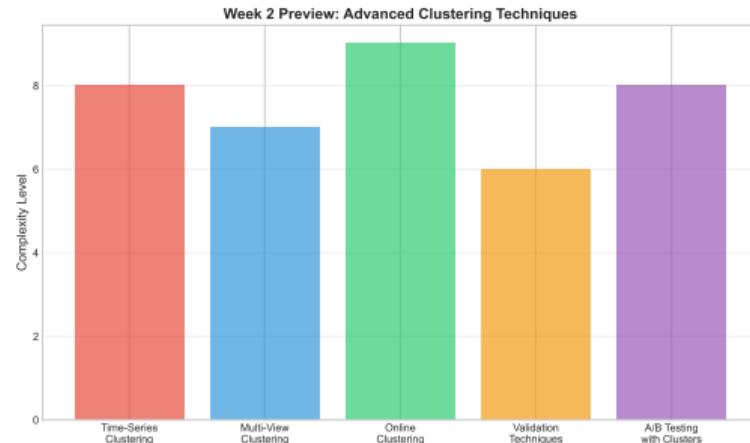
Datasets:

- UCI ML Repository
- Kaggle competitions
- Innovation dataset collection

Community:

- Course Slack channel
- Office hours: Wed 3-5pm
- Peer study groups

Next Week Preview:



Week 2: Advanced Clustering

- Spectral clustering
- Mean shift algorithm
- Affinity propagation
- Ensemble methods

Glossary of Technical Terms

Key Concepts Quick Reference

Clustering Algorithms:

- **K-Means:** Partitions data into K predefined clusters
- **DBSCAN:** Density-based spatial clustering
- **Hierarchical:** Builds cluster tree (dendrogram)
- **GMM:** Gaussian Mixture Models, soft clustering

Key Parameters:

- **K:** Number of clusters
- **eps:** Neighborhood radius (DBSCAN)
- **min_samples:** Minimum points for density
- **n_init:** Number of random initializations

Evaluation Metrics:

- **Silhouette:** Cluster cohesion vs separation [-1,1]
- **Inertia:** Sum of squared distances to centroids
- **Davies-Bouldin:** Ratio of within to between distances
- **Calinski-Harabasz:** Ratio of dispersions

Innovation Terms:

- **Empathy Mapping:** Understanding user perspectives
- **Pain Points:** User problems/frustrations
- **User Archetypes:** Representative user groups
- **Innovation Ecosystem:** Connected stakeholders

Implementation Checklist

Your Step-by-Step Guide to Success

Data Preparation:

- Collect innovation feedback data
- Clean and remove duplicates
- Handle missing values
- Normalize/standardize features
- Create feature vectors

Algorithm Selection:

- Analyze data distribution
- Choose appropriate algorithm
- Set initial parameters
- Prepare validation strategy

Implementation:

- Run clustering algorithm
- Calculate evaluation metrics
- Visualize results (PCA/t-SNE)
- Validate with domain experts
- Iterate and refine

Innovation Application:

- Map clusters to user personas
- Identify innovation opportunities
- Create targeted solutions
- Design prototype features
- Test with user groups

Ready? Start with data preparation and work your way down!