Machine Learning for Innovation

Week 1: Finding Patterns with Clustering

BSc Data Science & AI Program

No prerequisites required

September 13, 2025

By the end of today:

- Understand what clustering does
- Know when to use it
- Apply 3 different methods
- Interpret the results
- Practice with real examples

No math required!

We'll use pictures and examples instead.

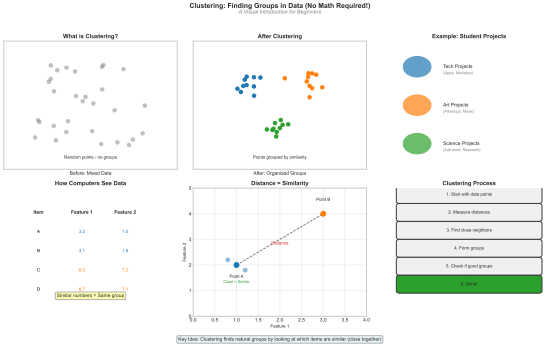

Part 1

Understanding the Problem

Why we need to find patterns

Imagine you have:

- 1000 customer reviews
- 500 product ideas
- 10,000 survey responses

The problem:

- Too much to read manually
- Hidden patterns we can't see
- Takes too long to analyze

The solution:

Let computers find the groups for us!

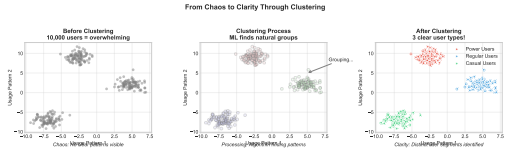

From chaos to organized groups

What is Clustering?

A simple explanation

Clustering: Finding Groups in Data (No Math Required!)

A Visual Introduction for Beginners

What is Clustering?

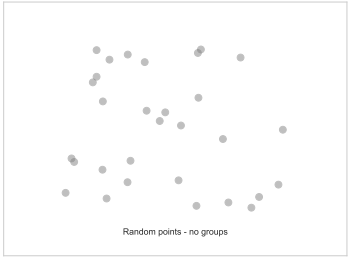

Before: Mixed Data

After Clustering

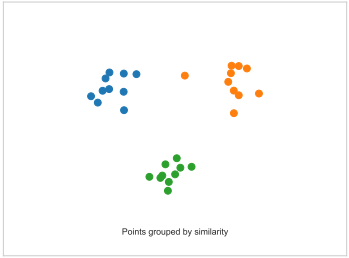

After: Organized Groups

Example: Student Projects

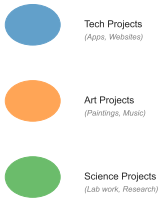

How Computers See Data

Item	Feature 1	Feature 2
A	3.2	1.5
B	3.1	1.6
C	8.5	7.2
D	8.7	7.1

Distance = Similarity

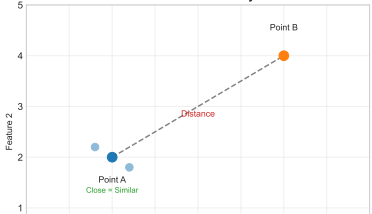

Clustering Process

1. Start with data points
2. Measure distances
3. Find close neighbors
4. Form groups
5. Check if good groups
6. Done!

Where We Use Clustering

Everyday examples

Business:

- Group similar customers
- Organize products in store
- Find spending patterns

Science:

- Group similar genes
- Classify star types
- Identify weather patterns

Daily Life:

- Music playlists (similar songs)
- Friend suggestions (similar interests)
- News categories (similar topics)

Innovation:

- Group similar ideas
- Find user needs
- Identify opportunities

All these use clustering to find patterns!

Part 2

How Clustering Works

Three different approaches

Three Ways to Find Groups

Each has strengths and weaknesses

Clustering Algorithms: Simple Comparison

Three Different Ways to Find Groups

K-Means: How it Works

1. Pick 3 center points
2. Assign each point to nearest center
3. Move centers to middle of groups
4. Repeat until stable

Like organizing by neighborhoods

K-Means Result

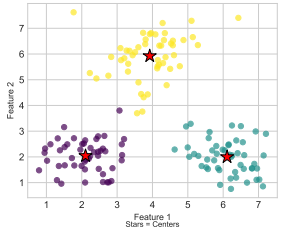

When to Use K-Means

Good for:

- Round groups
- Similar sizes
- Fast results

Not good for:

- Weird shapes
- Different sizes

Speed & Difficulty

Speed: FAST

Difficulty: EASY

You need to know:

- How many groups (K)

DBSCAN: How it Works

1. Look at each point
2. Count neighbors nearby
3. If enough neighbors → core point
4. Connect core points → groups

Like finding crowds at a party

DBSCAN Result

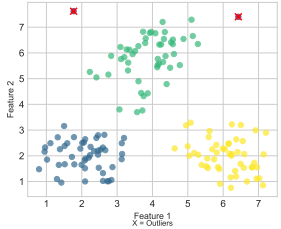

When to Use DBSCAN

Good for:

- Any shape groups
- Finding outliers
- Unknown group count

Not good for:

- Different densities
- Need exact K groups

Speed & Difficulty

Speed: MEDIUM

Difficulty: MEDIUM

You need to know:

- Distance (eps)
- Min neighbors

Hierarchical: How it Works

Hierarchical Result

When to Use Hierarchical

Speed & Difficulty

K-Means: The Most Common Method

Like organizing by neighborhoods

How it works:

1. Decide how many groups (K)
2. Place K center points randomly
3. Assign each point to nearest center
4. Move centers to middle of their group
5. Repeat until stable

Good for:

- Round, similar-sized groups
- When you know how many groups
- Need fast results

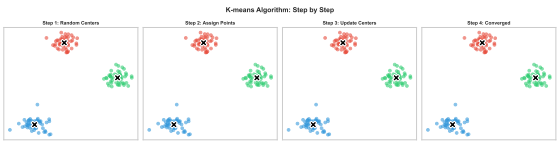

Watch the centers move to find groups

DBSCAN: Finding Any Shape

Like finding crowds at a party

How it works:

1. Look at each point
2. Count neighbors within distance
3. If enough neighbors → core point
4. Connect core points → groups
5. Points with few neighbors → outliers

Good for:

- Weird-shaped groups
- Finding outliers
- Don't know how many groups

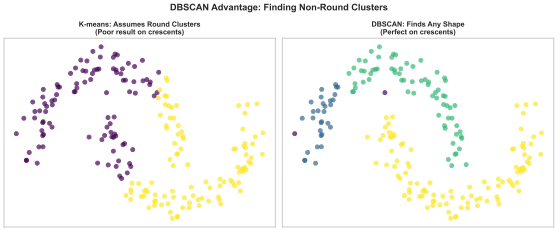

Can find any shape, even crescents!

How Good Are Our Groups?

Checking if clustering worked

What makes groups good?

- Points in same group are similar
- Different groups are different
- Groups make sense for your problem

Simple checks:

1. **Visual**: Do groups look separated?
2. **Size**: Are groups reasonable sizes?
3. **Meaning**: Can you explain each group?
4. **Stability**: Same result if run again?

Silhouette Score:
-1 = Bad (overlapping)
0 = OK (touching)
+1 = Good (separated)

Quick Decision Tree

Use K-Means if:

- Need speed
- Know number of groups
- Groups are round
- Similar sizes expected

Use DBSCAN if:

- Groups have weird shapes
- Want to find outliers
- Don't know group count
- Groups have different densities

Use Hierarchical if:

- Small dataset (¡500 points)
- Want to see all groupings
- Need tree structure
- Can wait for results

Not sure? Try K-Means first - it's simplest!

Part 3

Let's Practice

A simple example you can follow

Practice: Student Study Patterns

Follow along with this example

Practice Example: Grouping Students by Study Habits

A Simple Clustering Exercise

The Problem

A teacher wants to understand different study patterns in class.

Data collected:

- Hours studied per week
- Number of questions asked

50 students total

Goal: Find study groups

Student Data

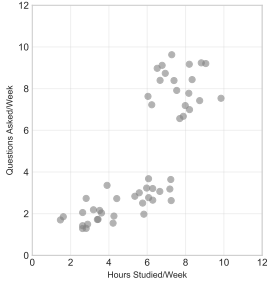

Each dot = 1 student

After Clustering

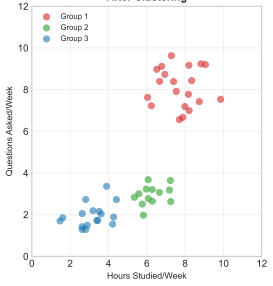

What We Found

Group 1:

Need extra help

- *Low hours, few questions*

Group 2:

Independent learners

- *Good hours, few questions*

Group 3:

Highly engaged

- *Many hours, many questions*

Your Turn: Step 1

Load the data:

```
import pandas as pd
data = pd.read_csv("students.csv")
```

Look at it:

```
print(data.head())
```

You should see:

```
  hours  questions
```

Your Turn: Step 2

Prepare the data:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(data)
```

Why scale?

Makes features comparable

Apply clustering:

Your Turn: Step 3

See the results:

```
import matplotlib.pyplot as plt
plt.scatter(data["hours"],
            data["questions"],
            c=labels)
plt.xlabel("Hours")
plt.ylabel("Questions")
plt.show()
```

Check Your Work

- ☐ Did you find 3 groups?
- ☐ Are groups visually separated?
- ☐ Do groups make sense?

What to look for:

- Clear differences between groups
- Similar students in same group
- Groups tell a story

Your Turn: Exercise

Try it yourself

Dataset: Store Products

You have data about 200 products:

- Price (0-100)
- Customer rating (1-5 stars)
- Sales per month

Your tasks:

1. Load the data
2. Apply K-Means with K=3
3. Plot the results
4. Describe each group
5. Try K=4, which is better?

Starter code:

```
import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

# Load data
data = pd.read_csv('products.csv')

# Apply clustering
kmeans = KMeans(n_clusters=3)
labels = kmeans.fit_predict(data)

# Plot results
plt.scatter(data['price'],
data['rating'],
c=labels)
plt.show()
```

Common Mistakes to Avoid

Learn from these errors

Mistake 1: Forgetting to scale

- Problem: Price (0-100) vs Rating (1-5)
- Price dominates because bigger numbers
- Solution: Always scale your data first!

Mistake 2: Wrong K

- Too few groups: Miss important patterns
- Too many groups: Overly complex
- Solution: Try different K values

Mistake 3: Ignoring outliers

- One weird point can ruin groups
- K-Means pulls centers toward outliers
- Solution: Check for outliers first

Mistake 4: Not checking results

- Algorithm always gives an answer
- Doesn't mean it's meaningful!
- Solution: Always visualize and interpret

Remember: The computer finds patterns, YOU decide if they make sense!

Week 1 Summary

What you learned today

Concepts

- Clustering finds groups
- No labels needed
- Distance = similarity
- Multiple methods exist

Methods

- K-Means (fast, simple)
- DBSCAN (any shape)
- Hierarchical (tree view)
- Each has trade-offs

Skills

- Choose algorithm
- Apply clustering
- Check quality
- Interpret results

You can now find patterns in data!

Next week: More advanced clustering techniques

Resources for Practice

Where to learn more

Practice datasets:

- Iris flowers (150 samples, 4 features)
- Wine quality (178 samples, 13 features)
- Mall customers (200 samples, 5 features)

Python libraries:

- scikit-learn (all algorithms)
- pandas (data handling)
- matplotlib (visualization)

Online resources:

- scikit-learn.org/stable/modules/clustering
- Google Colab (free Python online)
- Kaggle Learn (free courses)

Help available:

- Office hours: Wed 3-5pm
- Course forum
- Study groups

Questions? Just ask - no question is too simple!