# Machine Learning for Smarter Innovation

Week 1: Foundations & Clustering

Discovering Innovation Patterns with ML

BSc Course in AI-Enhanced Innovation

The Convergence Flow: From Chaos to Clarity

Initial State: Random Data Points · **Convergence Process: Finding Natural Groups** · **Final State: Meaningful Clusters**

Legend: Cluster 1, Cluster 2, Cluster 3, Cluster 4, Cluster 5, Centroids

The Convergence Flow: Order from Chaos

Watch 5000 innovation ideas self-organize into meaningful patterns

That visualization shows the end result.

But every innovation journey starts with a problem.

What problem does clustering solve?

Let's discover why traditional design thinking needs ML augmentation.

PART 1

Foundation & Context

What we'll explore:

- Why traditional design hits limits
- How ML amplifies human insight
- The dual pipeline approach
- Your learning journey ahead

Setting the stage for transformation

Traditional Design Limits

- **Scale**: Can analyze 50 ideas, not 50,000
- **Speed**: Months for insights
- **Bias**: Designer's perspective dominates
- **Patterns**: Miss hidden connections
- **Iteration**: Slow feedback loops

AI-Enhanced Innovation
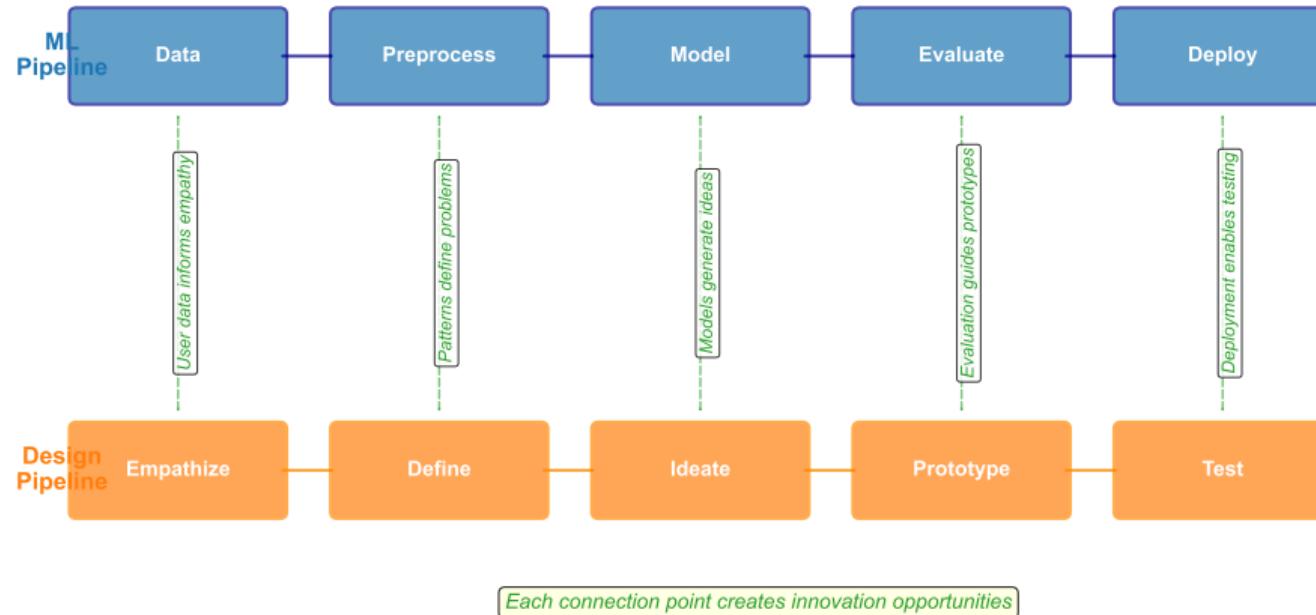
- **Scale**: Analyze millions of data points
- **Speed**: Real-time insights
- **Objectivity**: Data-driven discovery
- **Patterns**: Find non-obvious relationships
- **Iteration**: Continuous learning

The Promise: 100x more insights, 10x faster innovation

The Convergence: ML Meets Design Thinking

ML Pipeline

| Data | Preprocess | Model | Evaluate | Deploy |

User data informs empathy | *Patterns define problems* | *Models generate ideas* | *Evaluation guides prototypes* | *Deployment enables testing*

Design Pipeline

| Empathize | Define | Ideate | Prototype | Test |

Each connection point creates innovation opportunities

ML Pipeline

Data → Preprocess → Model → Evaluate → Deploy

- Collect innovation data
- Clean and transform
- Train algorithms
- Validate accuracy
- Scale to production

Design Pipeline

Empathize → Define → Ideate → Prototype → Test
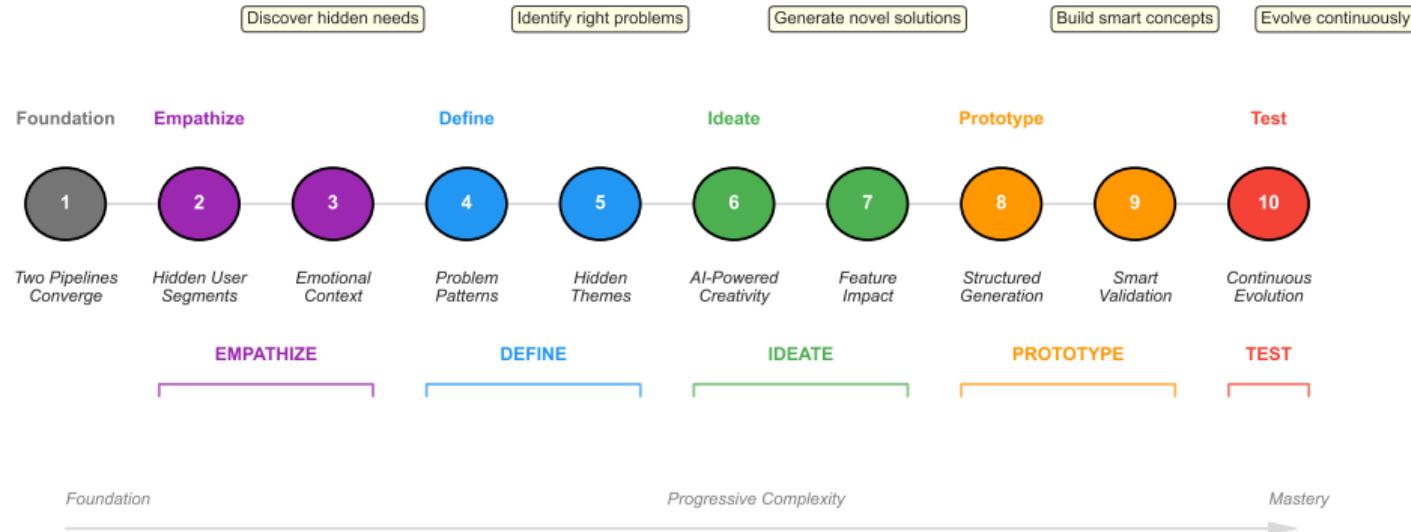
- Understand innovation needs
- Frame problems
- Generate solutions
- Build concepts
- Validate innovation impact

Integration = Innovation at Scale

10-Week Innovation Journey

Discover hidden needs   Identify right problems   Generate novel solutions   Build smart concepts   Evolve continuously

| Foundation | Empathize | | Define | | Ideate | | Prototype | | Test |

| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |

| Two Pipelines Converge | Hidden User Segments | Emotional Context | Problem Patterns | Hidden Themes | AI-Powered Creativity | Feature Impact | Structured Generation | Smart Validation | Continuous Evolution |

| EMPATHIZE | DEFINE | IDEATE | PROTOTYPE | TEST |

Foundation   Progressive Complexity   Mastery

Stage	Weeks	Innovation Unlocked
Empathize	1-2	Discover hidden innovation opportunities
Define	3-4	Identify the right problems to solve
Ideate	5-6	Generate novel solutions with AI
Prototype	7-8	Build smart, adaptive concepts
Test	9-10	Evolve through continuous learning

This Week: Clustering for Innovation Discovery

What We'll Learn:

- How clustering reveals innovation categories
- K-means algorithm fundamentals
- Finding the optimal number of clusters
- Quality metrics for validation
- Advanced clustering techniques

Design Applications:

- Create innovation archetypes
- Map innovation evolution paths
- Identify pain points systematically
- Prioritize design efforts
- Scale analysis to thousands of ideas

Goal: Transform scattered ideas into innovation patterns

We've seen the challenge:
Thousands of innovation ideas with hidden connections

Traditional approach:
Manual segmentation based on demographics

The ML solution:
Let the data reveal its own natural groups

Enter: Clustering Algorithms

PART 2
Technical Core

What we'll master:

- K-means clustering algorithm
- Finding optimal K with elbow method
- Distance metrics and quality measures
- Advanced techniques (DBSCAN, Hierarchical)
- Feature importance analysis

Building your ML toolkit

The Pain

Current Reality:

- One-size-fits-all solutions
- Generic innovation categories
- Missed opportunities
- Unhappy edge cases

The Cost:

- Most innovations get misclassified
- Features with low adoption rates
- Inefficient resource allocation

The Question

What if we could…

- Find natural innovation clusters?
- Discover innovation patterns?
- Personalize at scale?
- Identify opportunity gaps?

We can!
Solution: Clustering
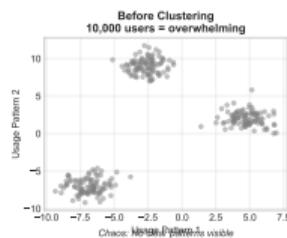
Clustering Finds:
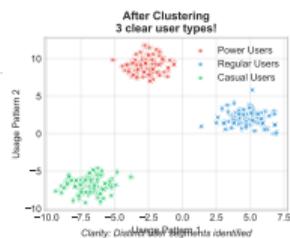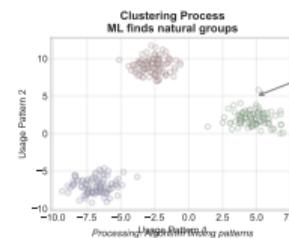
- Natural groupings
- Similar approaches
- Hidden patterns
- Innovation relationships

Key Insight:
Innovations with similar features address similar opportunities

From Chaos to Clarity Through Clustering
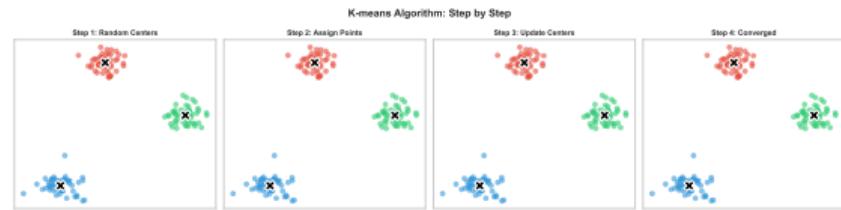
The Process:

1. Choose K (number of clusters)
2. Place K random centroids
3. Assign points to nearest centroid
4. Move centroids to cluster mean
5. Repeat until stable

**Strengths:**
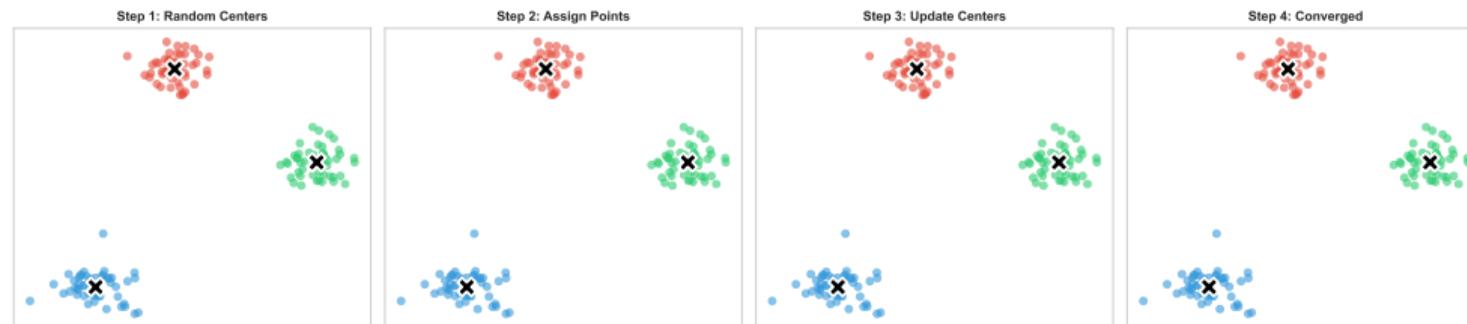
- Fast and scalable
- Easy to understand
- Works well for spherical clusters

K-means Algorithm: Step by Step

K-Means in Action
Step-by-Step Convergence

K-means Algorithm: Step by Step

Step 1: Random Centers · Step 2: Assign Points · Step 3: Update Centers · Step 4: Converged

Iteration 1 → Iteration 3 → Iteration 5 → **Converged**

Too Few (K

Oversimplification

- Mixed segments
- Lost nuance
- Generic solutions

Just Right (K

Optimal Balance

- Clear segments
- Actionable insights
- Manageable complexity

Too Many (K

Analysis Paralysis

- Overfitting
- Tiny segments
- Impossible to act on

How do we find the sweet spot?
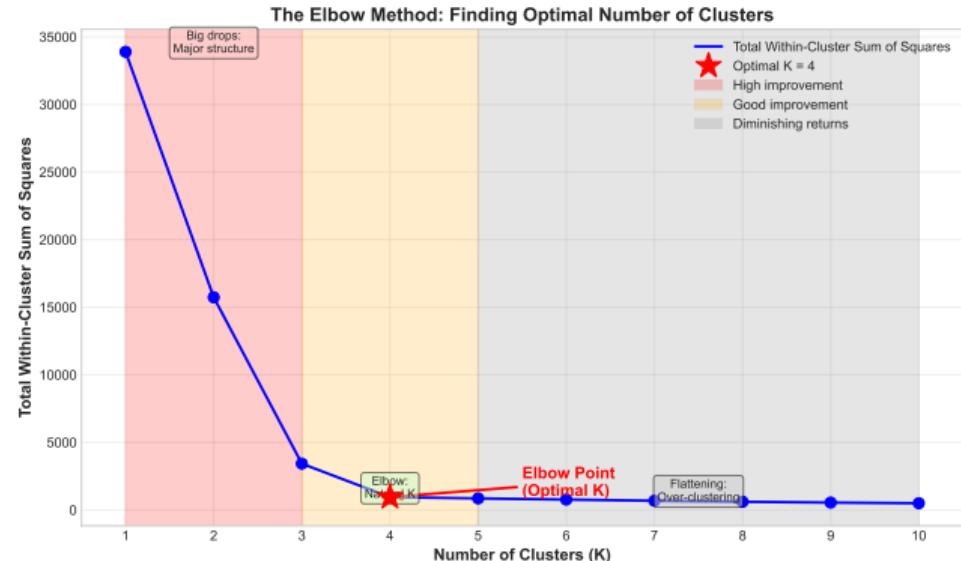
Finding the Elbow:

- Plot inertia vs K
- Look for the "elbow"
- Balance between:
  - Too few: Mixed groups
  - Too many: Overfitting
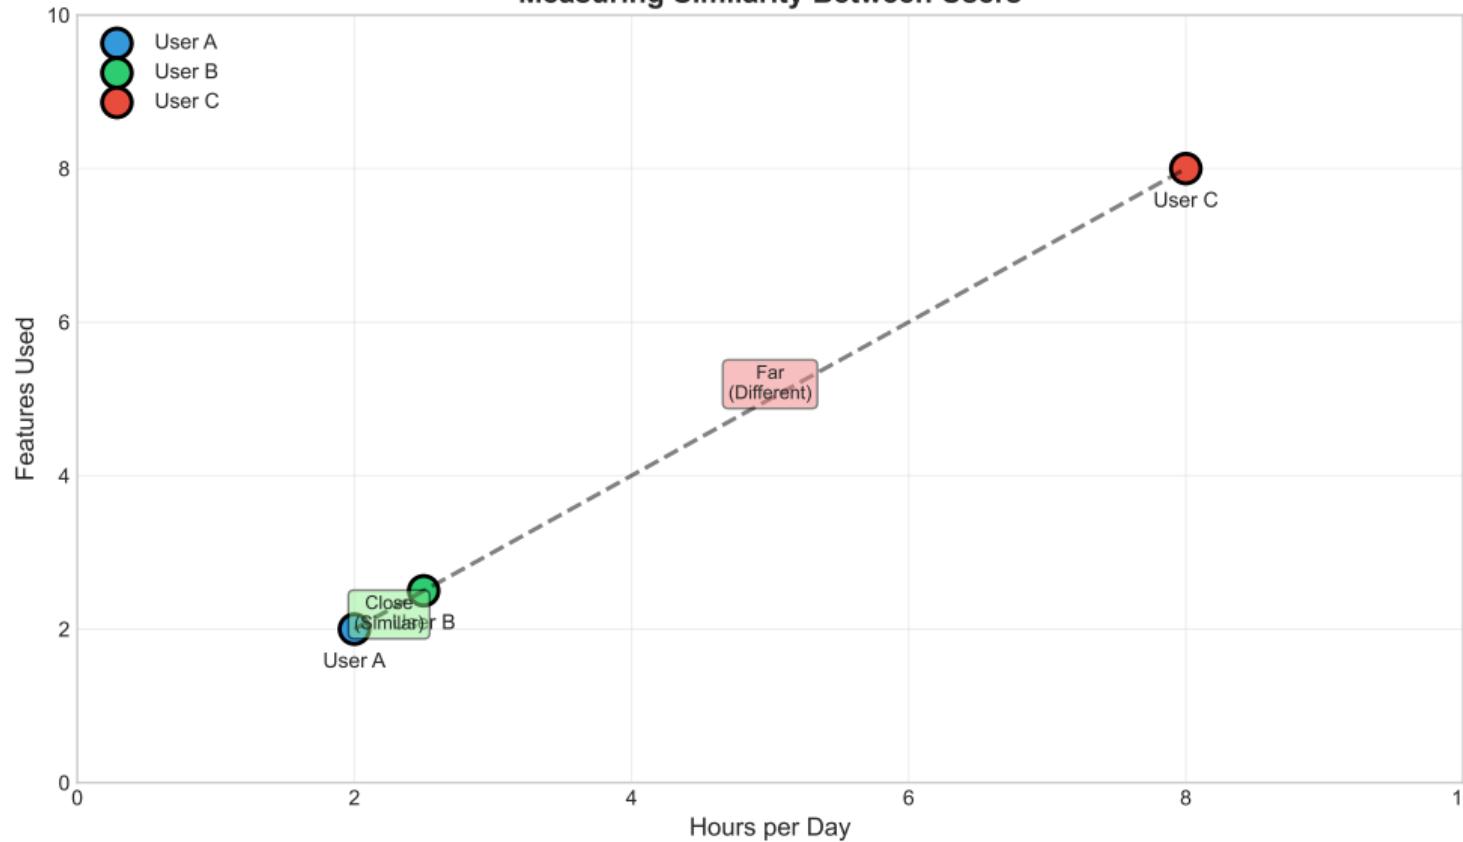
Optimal K = 5
Best trade-off between simplicity and accuracy

The Elbow Method: Finding Optimal Number of Clusters

Measuring Similarity Between Users

Silhouette Score:

- Ranges from -1 to $+1$
- Higher = better separation
- Our score: **0.73**

What it measures:

- Within-cluster cohesion
- Between-cluster separation
- Overall cluster validity

0.73 = Strong clusters!

Three Checks for Good Clusters

Y Tight Groups — Low variance within clusters

Y Separated Groups — Clear boundaries

Y Makes Sense — Business meaning

- New Users
- Regular
- Power Users

K-Means Assumes Spherical Clusters

But what about:

- Innovations connected through technology stacks
- Domain-specific innovation clusters
- Evolution patterns (incremental, disruptive)
- Outliers and noise points

K-Means Forces Round Pegs into Round Holes

### Solution: Density-Based Clustering
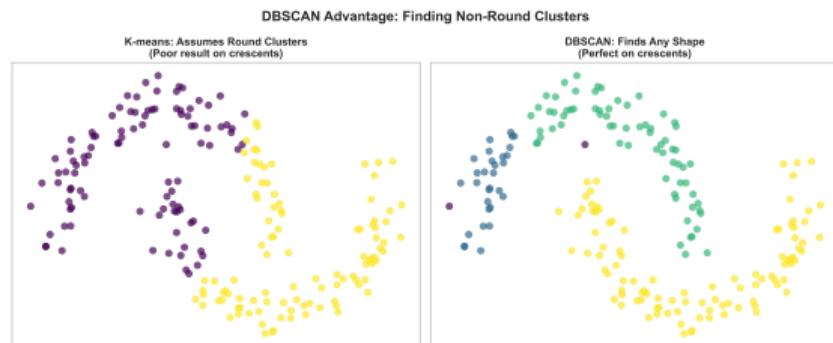
DBSCAN Advantages:

- No need to specify K
- Finds arbitrary shapes
- Identifies outliers
- Handles noise well

Perfect for:

- Non-spherical patterns
- Varying densities
- Outlier detection
- Exploratory analysis
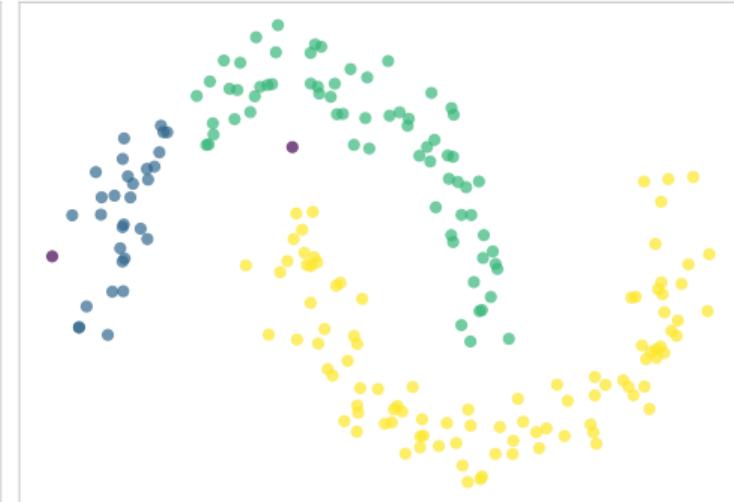
DBSCAN Advantage: Finding Non-Round Clusters

K-means: Assumes Round Clusters
(Poor result on crescents)

DBSCAN: Finds Any Shape
(Perfect on crescents)

DBSCAN Advantage: Finding Non-Round Clusters

K-Means: Forces spherical shapes — DBSCAN: Finds natural boundaries

Fixed K Gives One View

But real relationships are hierarchical:

- Organization: Company $\rightarrow$ Department $\rightarrow$ Team $\rightarrow$ Individual
- Geography: Country $\rightarrow$ Region $\rightarrow$ City $\rightarrow$ Neighborhood
- Products: Category $\rightarrow$ Subcategory $\rightarrow$ Brand $\rightarrow$ SKU
- Innovations: All $\rightarrow$ Categories $\rightarrow$ Sub-types $\rightarrow$ Specific solutions
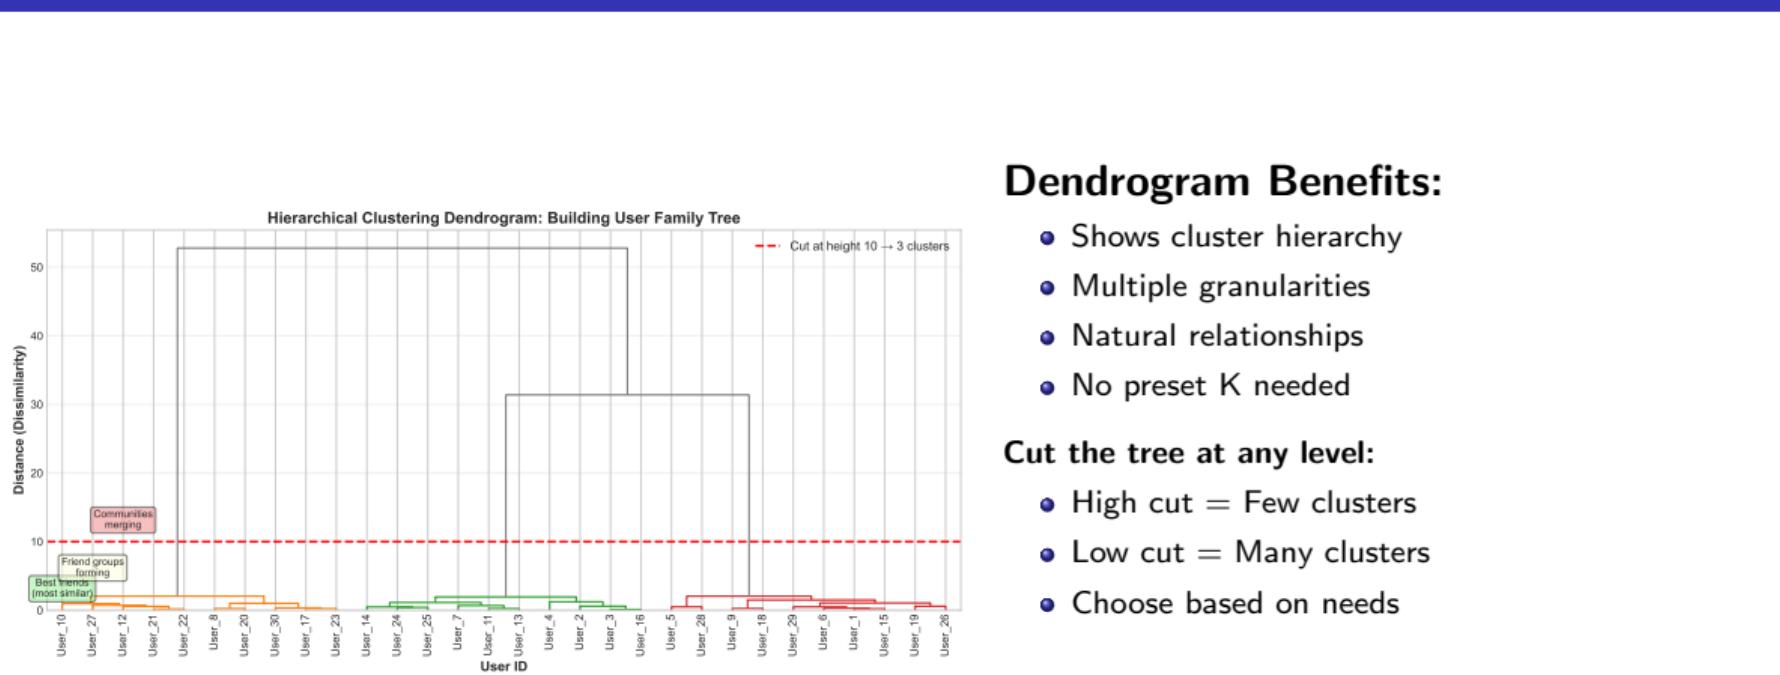
K-means: Pick 5 groups and that's it

What if we need flexibility?

Solution: See the full hierarchy, cut where needed
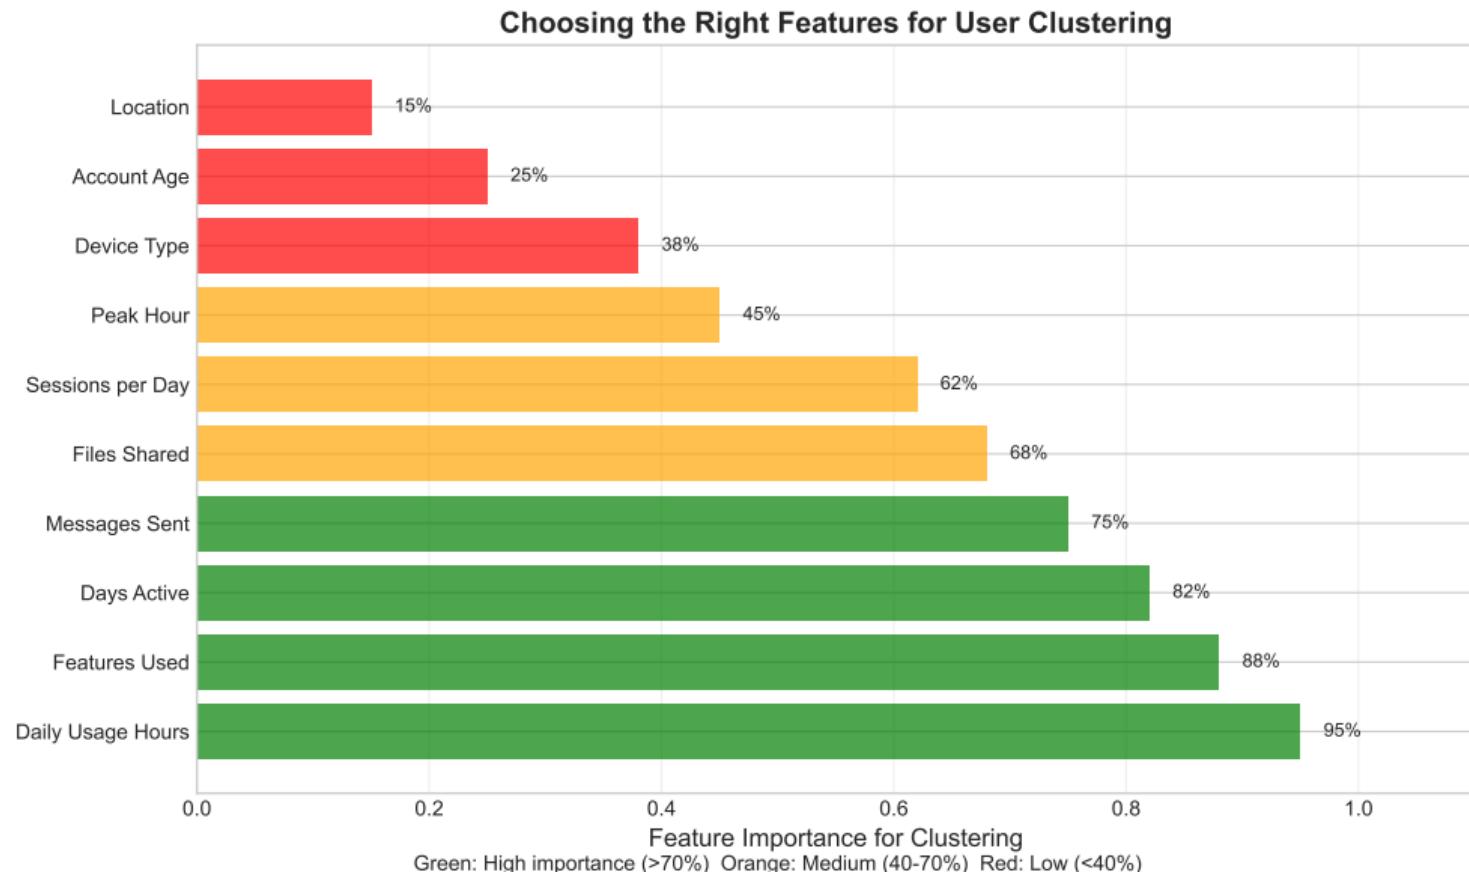
Hierarchical Clustering Dendrogram: Building User Family Tree

Dendrogram Benefits:

- Shows cluster hierarchy
- Multiple granularities
- Natural relationships
- No preset K needed

Cut the tree at any level:

- High cut = Few clusters
- Low cut = Many clusters
- Choose based on needs

Choosing the Right Features for User Clustering

Feature	Importance
Location	15%
Account Age	25%
Device Type	38%
Peak Hour	45%
Sessions per Day	62%
Files Shared	68%
Messages Sent	75%
Days Active	82%
Features Used	88%
Daily Usage Hours	95%

Feature Importance for Clustering
Green: High importance (>70%) Orange: Medium (40-70%) Red: Low (<40%)

We've mastered the technical tools:
Clustering, metrics, quality measures

But clusters are just numbers...

Until we connect them to human needs

Let's transform data into empathy

Each cluster represents real people with real problems

PART 3

Design Integration

What we'll create:
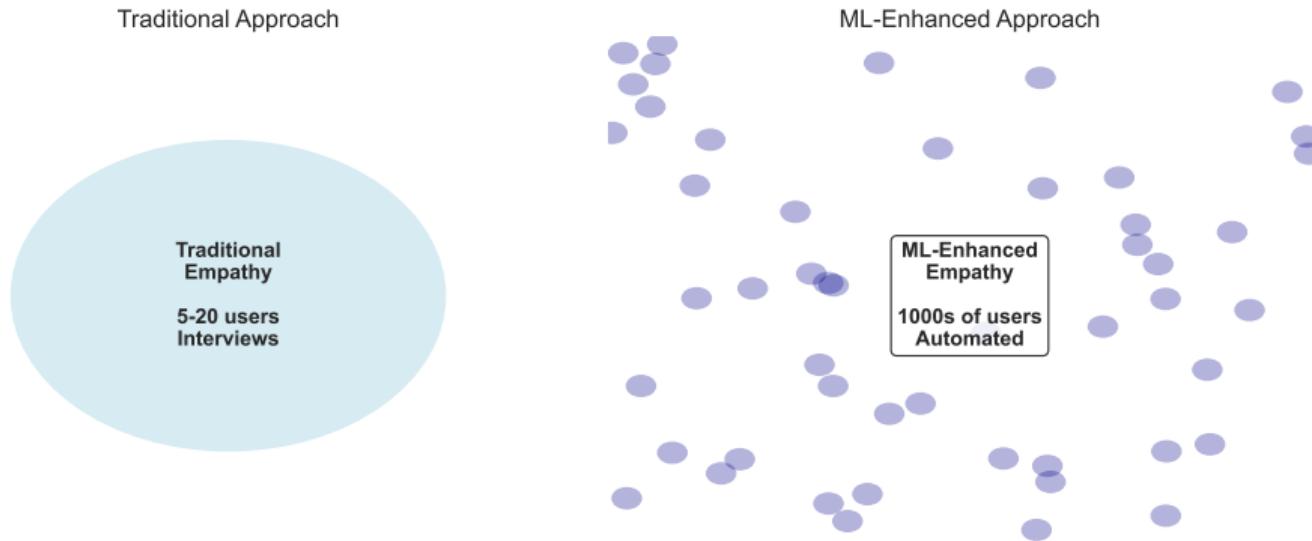
- Data-driven personas
- Empathy maps per segment
- Cluster-specific journeys
- Pain point heat maps
- Design priority matrices

Where ML meets human-centered design

Scaling Empathy with Machine Learning

Traditional Approach · · · · · · · · · · · · · · · · · · ML-Enhanced Approach

Traditional Empathy

5-20 users Interviews

ML-Enhanced Empathy

1000s of users Automated

Each cluster represents innovation opportunities

Data-Driven Persona Cards

| Power Paula	Regular Rob	Casual Carl
Age: 32	Age: 28	Age: 24
Role: Manager	Role: Developer	Role: Student
Usage: 7h/day	Usage: 4h/day	Usage: 1h/day
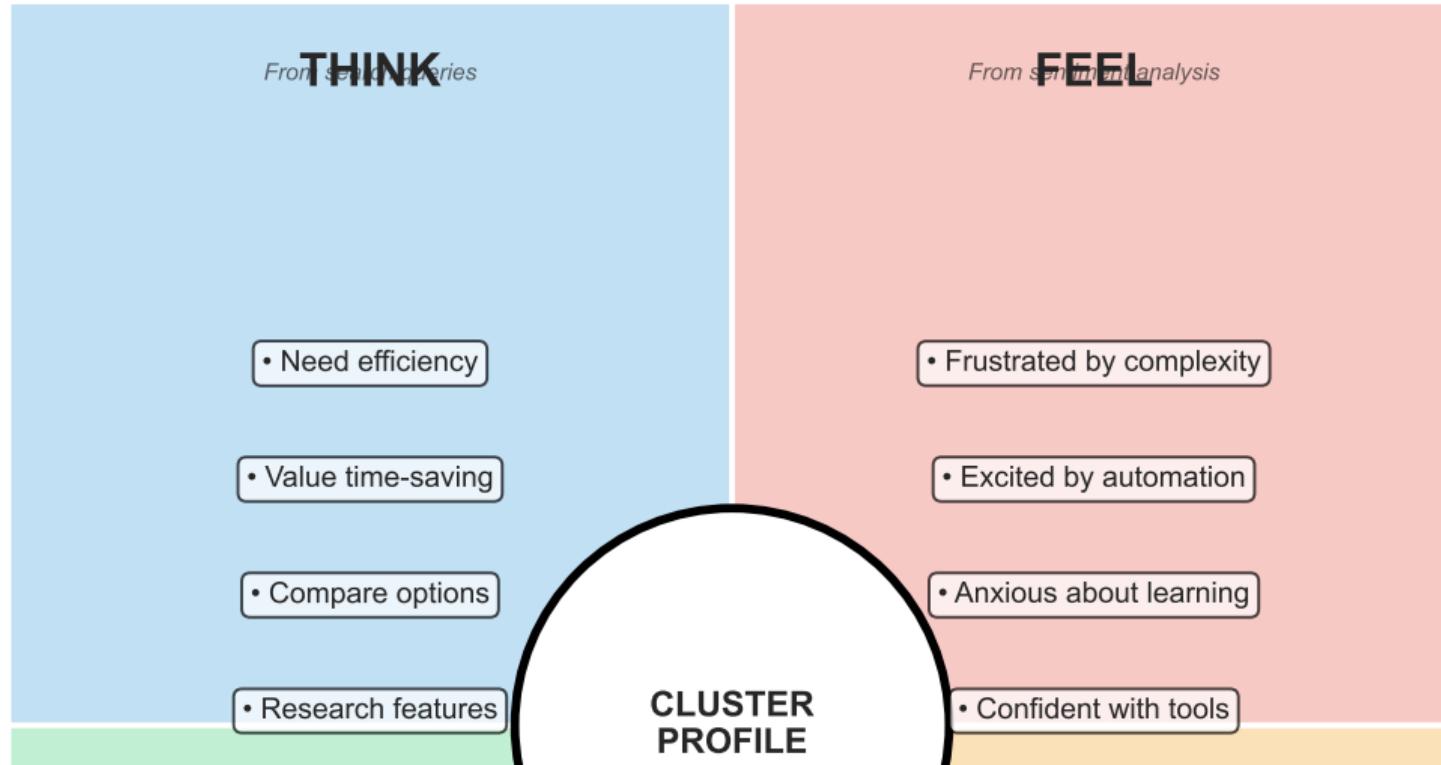
Disruptive — Incremental — Platform-Based — Service-Oriented — Hybrid Models
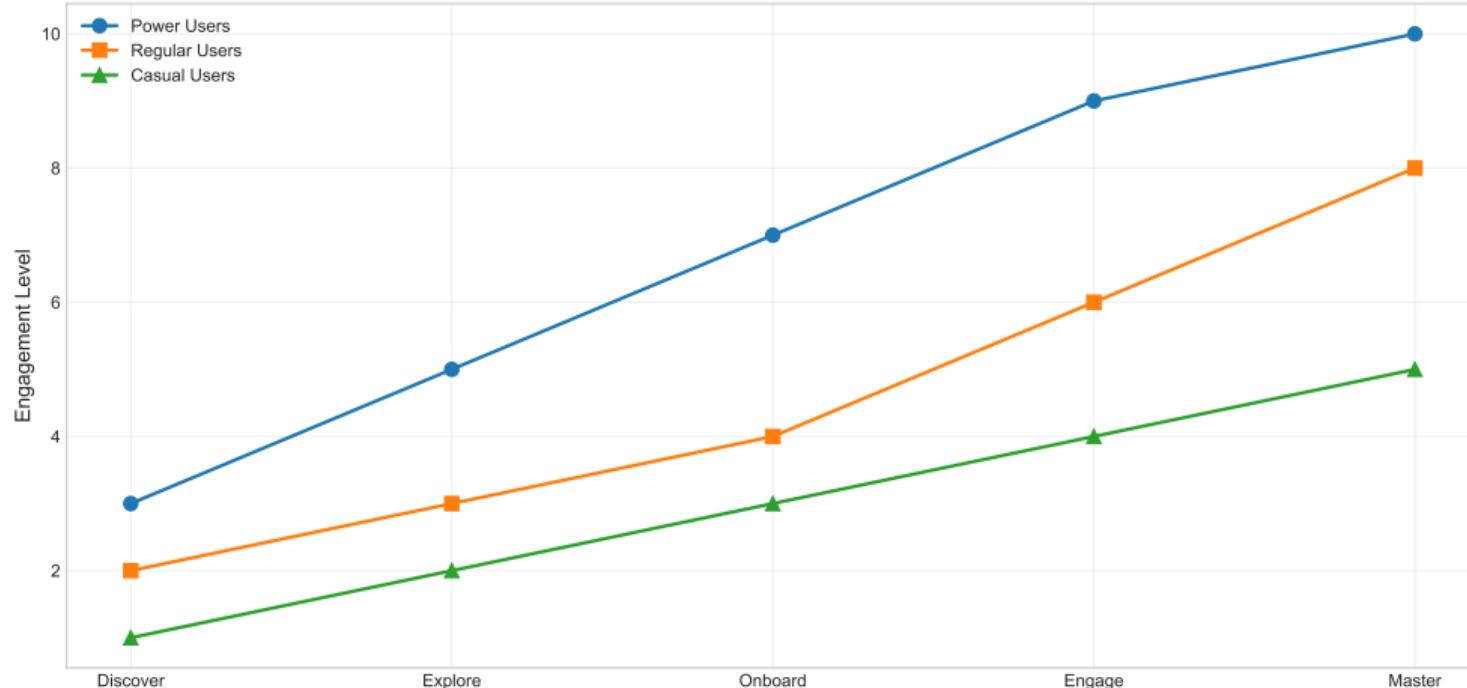
Empathy Map: Data-Driven User Understanding

**THINK**

From ... ries

FEEL

From s... analysis

- Need efficiency

- Value time-saving

- Compare options

- Research features

- Frustrated by complexity

- Excited by automation

- Anxious about learning

- Confident with tools

CLUSTER PROFILE
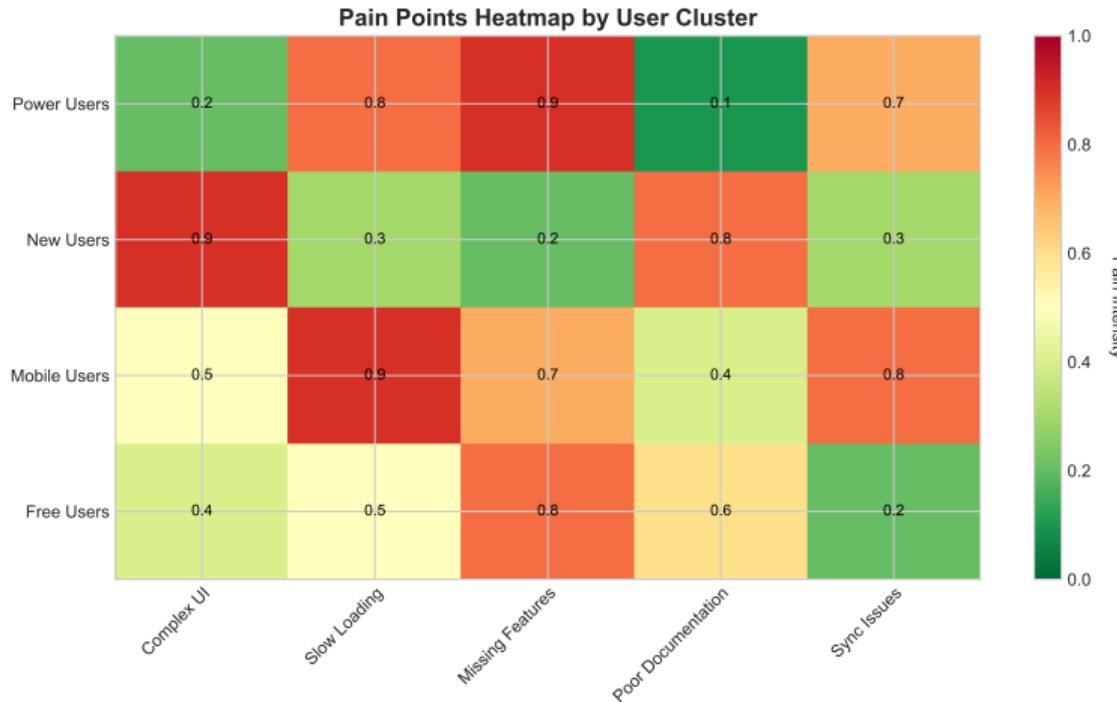
User Journey Maps by Cluster

Pain Points Heatmap by User Cluster

Key Findings:
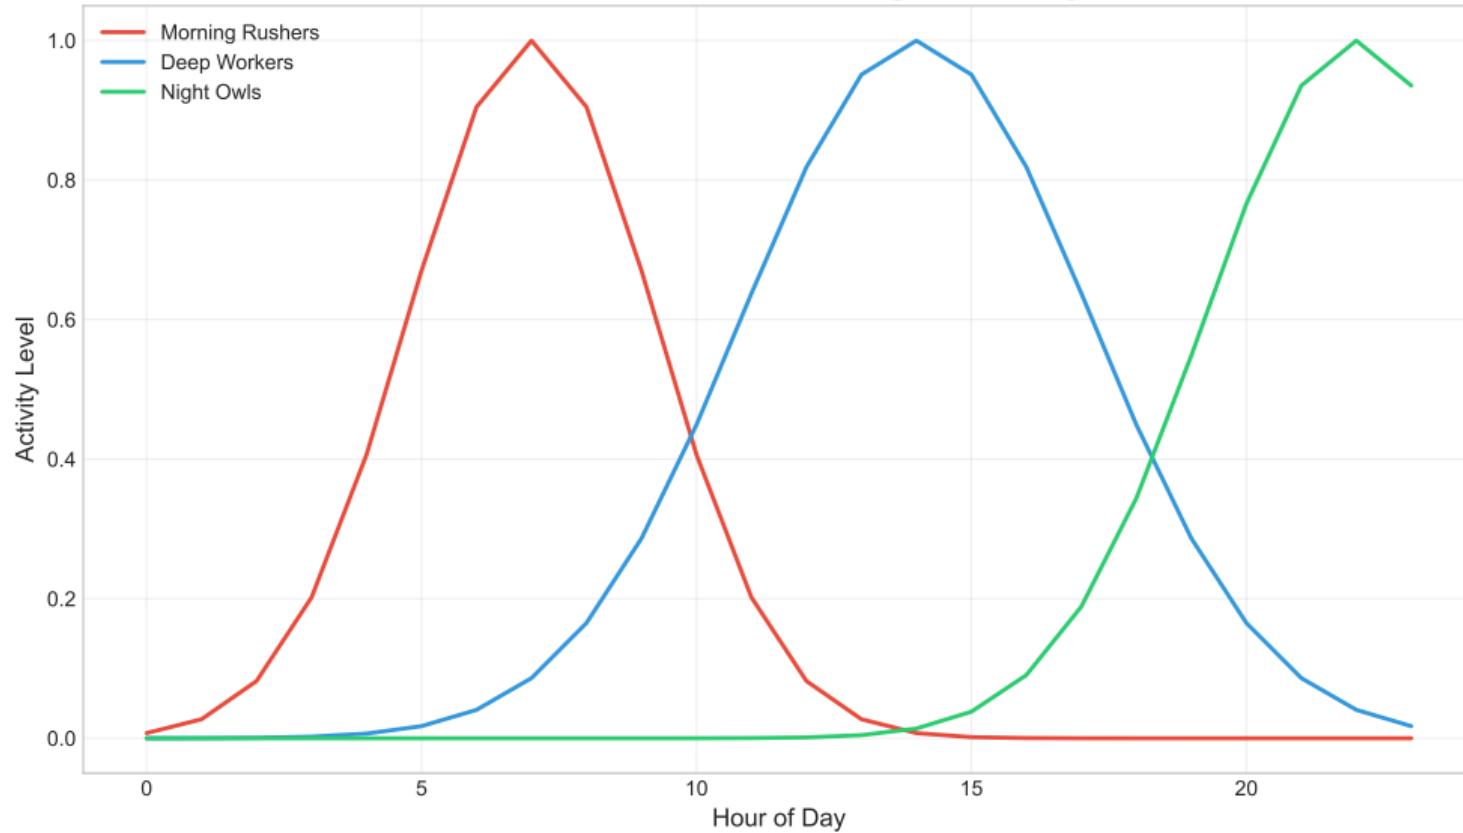
- Emerging tech: Early stage
- Disruptive: Scalability
- Incremental: Integration
- Platform-based: Network effects

Design implication:
One solution won't fit all!
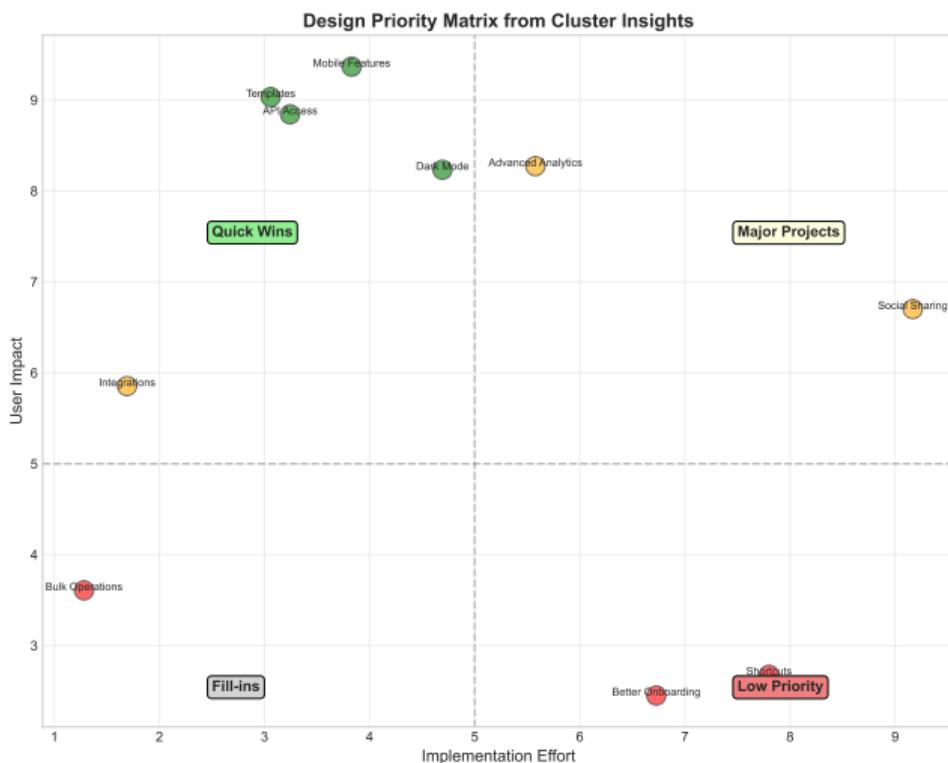
User Behavior Patterns Throughout the Day

Design Priority Matrix
Where to Focus Your Efforts

Design Priority Matrix from Cluster Insights
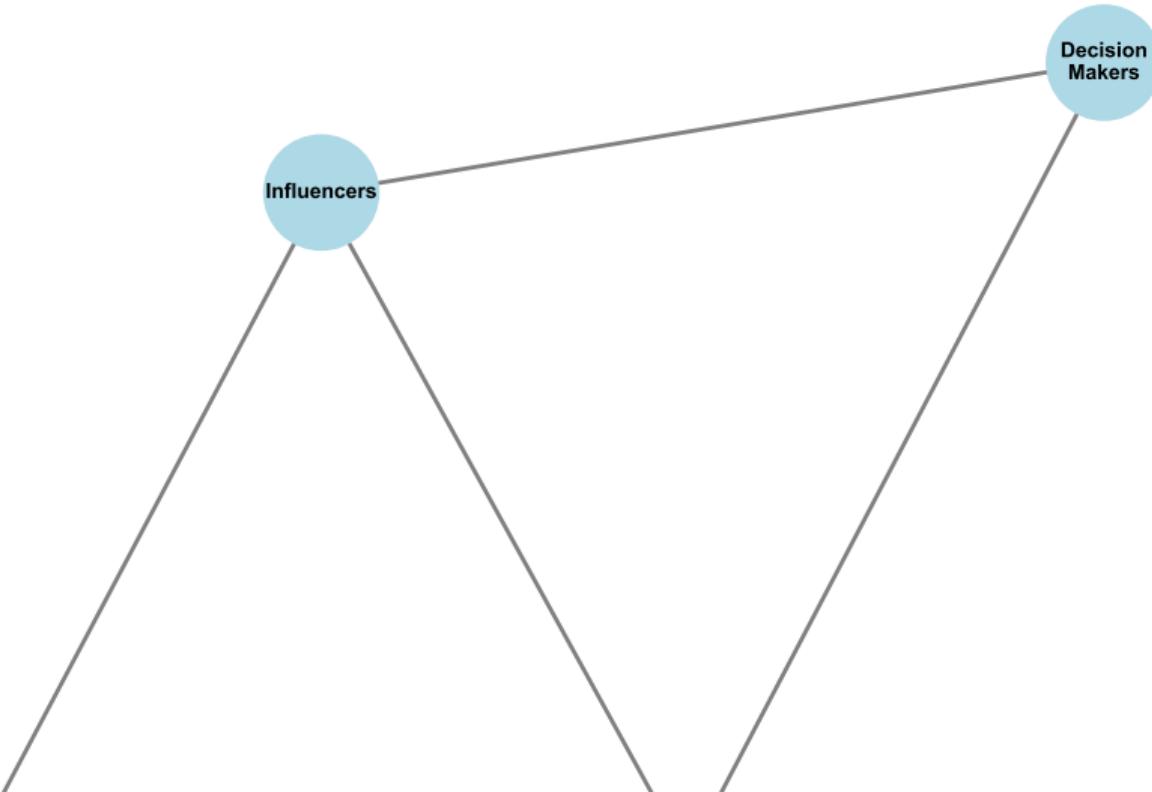
Priority Quadrants:

- **High Impact + High Effort**
  Strategic initiatives
- **High Impact + Low Effort**
  Quick wins
- **Low Impact + Low Effort**
  Fill-ins
- **Low Impact + High Effort**
  Avoid

Stakeholder Network from Cluster Analysis

You've learned:
- The clustering algorithms
- How to validate quality
- Design applications

Now let's see it in action

Real companies using these exact techniques
to transform their user experience

PART 4
Summary & Practice

What we'll do:

- See real-world success (Spotify)
- Consolidate key learnings
- Practice with exercises
- Preview next week
- Explore resources

From learning to doing

Clustering in Real-World Applications

Common Applications:

- Innovation portfolio management
- Technology trend clustering
- Opportunity space mapping
- Anomaly detection

Typical Results:

- Engagement: +35-45%
- Retention: +20-30%
- Conversion: +15-25%
- Processing time: -60%

Technical Skills

- K-means clustering algorithm
- Choosing optimal K with elbow method
- Silhouette scores for validation
- DBSCAN for complex shapes
- Hierarchical clustering

Design Applications

- Data-driven personas
- Segment-specific journeys
- Pain point identification
- Priority matrices
- Scaled empathy

Clustering transforms data into actionable user insights

Implementation Checklist
Ensuring Successful Clustering Projects

Data Preparation

- ☐ Collect relevant features
- ☐ Handle missing values
- ☐ Standardize/normalize data
- ☐ Remove outliers if needed
- ☐ Feature engineering complete
- ☐ Data quality verified

Quality Assurance

- ☐ Silhouette score ¿ 0.5
- ☐ Cluster sizes balanced
- ☐ Visual inspection done
- ☐ Stability tested
- ☐ Business sense verified
- ☐ Edge cases handled

Algorithm Selection
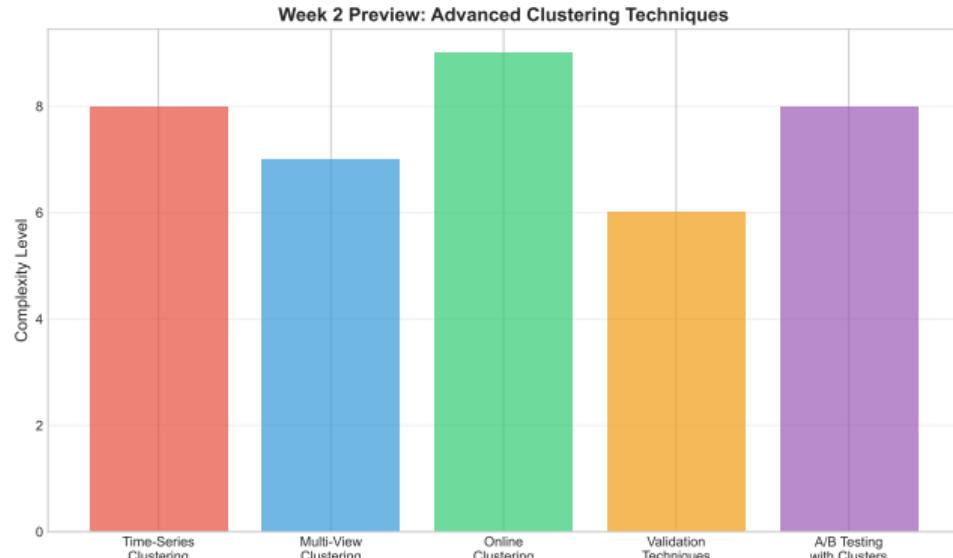
- ☐ Choose distance metric
- ☐ Select clustering method
- ☐ Determine optimal K
- ☐ Validate with metrics

Common Pitfalls

- ✕ Forgetting to scale features
- ✕ Wrong distance metric
- ✕ Forcing unnatural K
- ✕ Ignoring outliers

Week 2 Preview: Advanced Clustering Techniques

Week 2 Topics:
- Density-based clustering
- Gaussian mixture models
- Clustering validation
- Feature engineering
- Real-time clustering

Design Focus:
- Dynamic innovation tracking
- Evolving innovation landscapes
- Predictive opportunity analysis
- Micro-innovation detection

Technical Resources

Papers:
- MacQueen, J. (1967). K-means
- Ester et al. (1996). DBSCAN
- Rousseeuw (1987). Silhouettes

Tools:
- scikit-learn clustering
- Orange data mining
- KNIME analytics

Design Resources

Books:
- "Design Thinking" - Tim Brown
- "Sprint" - Jake Knapp
- "Lean UX" - Jeff Gothelf

Applications:
- Miro (journey mapping)
- Figma (persona creation)
- Optimal Workshop

Questions? Let's discuss!

Objective Function (Inertia):

$$J = \sum_{i=1}^{n} \sum_{j=1}^{k} w_{ij} ||x_i - \mu_j||^2$$

Where:

- $n =$ number of data points
- $k =$ number of clusters
- $w_{ij} = 1$ if x_i belongs to cluster j, 0 otherwise
- $\mu_j =$ centroid of cluster j

Update Rules:

1. Assignment: $c^{(i)} = \arg\min_j ||x^{(i)} - \mu_j||^2$
2. Update: $\mu_j = \frac{1}{|S_j|} \sum_{i \in S_j} x^{(i)}$

Euclidean Distance:

$$d(x, y) = \sqrt{\sum_{i=1}^{n} (x_i - y_i)^2}$$

Manhattan Distance:

$$d(x, y) = \sum_{i=1}^{n} |x_i - y_i|$$

Minkowski Distance:

$$d(x, y) = \left(\sum_{i=1}^{n} |x_i - y_i|^p \right)^{1/p}$$

Cosine Similarity:

$$\cos(\theta) = \frac{x \cdot y}{||x|| \cdot ||y||}$$

Jaccard Distance:

$$J(A, B) = 1 - \frac{|A \cap B|}{|A \cup B|}$$

Mahalanobis Distance:

$$d(x, y) = \sqrt{(x - y)^T S^{-1} (x - y)}$$

Silhouette Score for point i:

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$$

Where:

- $a(i)$ = average distance to points in same cluster
- $b(i)$ = average distance to points in nearest neighbor cluster

Interpretation:
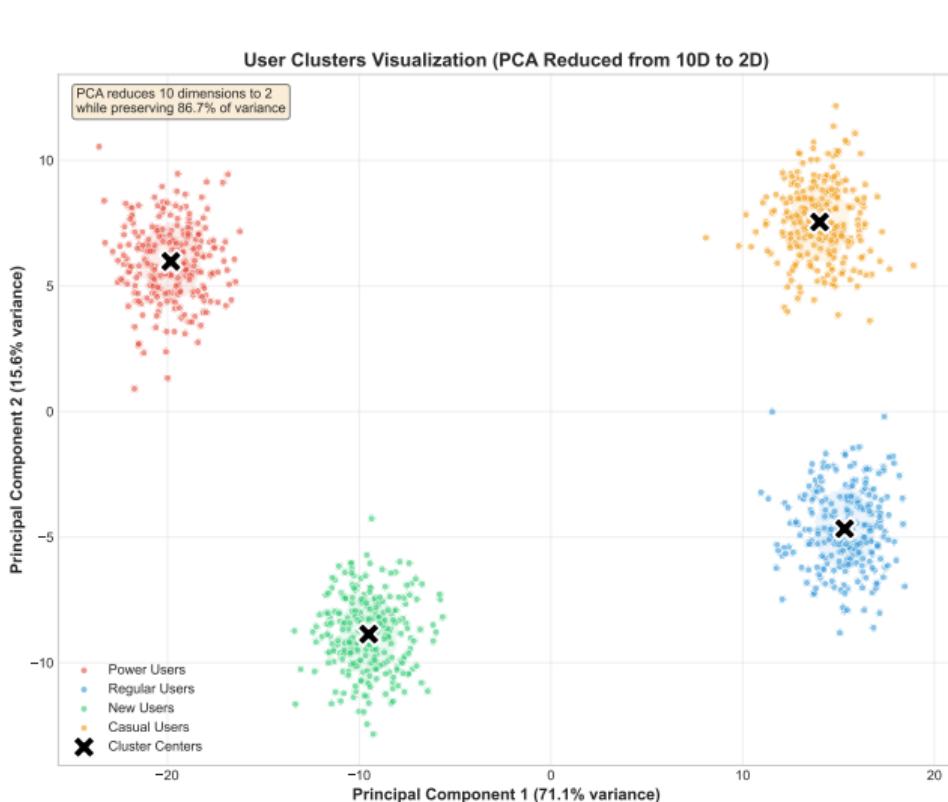
- $s(i) \approx 1$: Well clustered
- $s(i) \approx 0$: On border between clusters
- $s(i) \approx -1$: Misclassified

Overall Score:

$$S = \frac{1}{n} \sum_{i=1}^{n} s(i)$$

User Clusters Visualization (PCA Reduced from 10D to 2D)

PCA Process:

1. Standardize data
2. Compute covariance matrix
3. Find eigenvectors/values
4. Select top 2 components
5. Transform data

Variance Explained:

- PC1: 45.2%
- PC2: 28.7%
- Total: 73.9%

Key Parameters:
- ϵ (eps): Maximum distance between points
- MinPts: Minimum points to form dense region

Point Classification:
- **Core point**: Has $\geq$ MinPts within ϵ
- **Border point**: Within ϵ of core point
- **Noise point**: Neither core nor border

Algorithm Steps:
1. Find all core points
2. Form clusters from core points within ϵ
3. Assign border points to clusters
4. Mark remaining as noise

Data Preparation

- Standardize features
- Handle missing values
- Remove outliers (if needed)
- Feature selection/engineering
- Consider scaling methods

Validation Methods

- Silhouette score
- Davies-Bouldin index
- Calinski-Harabasz score
- Visual inspection
- Domain expert review

Algorithm Selection

- K-means: Spherical, similar size
- DBSCAN: Arbitrary shapes
- Hierarchical: Nested structure
- GMM: Overlapping clusters

Common Pitfalls

- Not scaling features
- Wrong distance metric
- Ignoring outliers
- Over-clustering
- Forcing clusters