

# Week 1: Foundation + Clustering

## ML/AI Design Thinking: Empathize Phase




BSc Data Science & Design

2025

# The Convergence Flow

How Random Data Points Become Meaningful Groups

The Convergence Flow: From Chaos to Clarity



*Watch as 1000 users organize themselves into natural groups*

## Traditional Design Thinking

- Empathize with users
- Define problems
- Ideate solutions
- Prototype ideas
- Test and iterate

## + Machine Learning Power

- Analyze thousands of users
- Find hidden patterns
- Generate insights automatically
- Validate with data
- Scale your understanding

This Week: Using clustering to truly understand your users

# The Empathize Phase: Understanding Your Users

## What is Empathizing?


- Walking in your users' shoes
- Understanding their needs, wants, fears
- Discovering what they don't tell you
- Finding patterns in behavior

## Traditional Methods:

- Interviews (5-20 people)
- Observations (days/weeks)
- Surveys (100s of responses)

## ML-Enhanced Methods:

- Analyze millions of interactions
- Find patterns humans miss
- Work 24/7 automatically
- Unbiased grouping



# Why Clustering Helps Us Understand People

## From Chaos to Clarity

### Before Clustering




10,000 users = overwhelming

### Clustering Process



ML finds natural groups

### After Clustering



5 clear user types!

- **Power Users:** Heavy usage, all features
- **Casuals:** Weekend usage, basic features
- **Professionals:** Business hours, productivity focus
- **Students:** Evening usage, collaboration features
- **Explorers:** Try everything once

# What is Clustering?

**Clustering = Finding Natural Groups in Data**



## Real-World Examples:

- Spotify: Grouping similar listeners
- Netflix: Finding viewer types
- Amazon: Customer segments
- Instagram: Content categories
- Gmail: Organizing emails

**Key Idea:** Items in the same group are more similar to each other than to items in other groups

No labels needed - the algorithm finds groups automatically!

# How Do We Measure “Similar”?

**Distance = Difference Between Things**


## Simple Example: App Usage

- User A: 2 hours/day, 10 features used
- User B: 2.5 hours/day, 12 features used
- User C: 8 hours/day, 50 features used

Who is more similar?

- A and B are close (similar usage)
- C is far from both (power user)

Think of it like: “How different are these users?”



## Common Measures:

- Straight line (Euclidean)
- City blocks (Manhattan)
- Correlation-based

# K-means: Finding Group Centers

Like Finding the Best Meeting Points for Groups

## How K-means Works:

- ① Pick K center points randomly
- ② Assign each user to nearest center
- ③ Move centers to group middle
- ④ Repeat until stable

## Real Example:

Finding 3 types of coffee drinkers:

- Morning rushers
- Afternoon socializers
- All-day workers



Pros: Fast, simple, scalable

Cons: Need to know K, assumes round clusters

# Hierarchical: Building a Family Tree


## Bottom-Up Approach:

- ① Start: Everyone is separate
- ② Find two most similar users
- ③ Group them together
- ④ Repeat with groups
- ⑤ Stop when all connected

## Like Making Friends:

- Best friends first
- Then friend groups
- Then communities
- Finally, everyone connected

## Dendrogram: The Family Tree



# How Do We Know Our Groups Are Good?

## Three Simple Checks

### 1. Tight Groups



Users in same group should be close together

### 2. Separated Groups



Different groups should be far apart

### 3. Makes Sense



Groups should mean something real

## Simple Metrics:

- **Elbow Method:** Plot error vs. number of clusters, look for “elbow”
- **Silhouette Score:** -1 (bad) to +1 (perfect), aim for  $\geq 0.5$
- **Business Sense:** Can you name and use each group?

## Getting Data Ready for Clustering

```
1 import pandas as pd
2 from sklearn.preprocessing import StandardScaler
3
4 # Load your user data
5 users = pd.read_csv('user_behavior.csv')
6
7 # Select features for clustering
8 features = ['daily_usage_hours', 'features_used',
9             'days_active', 'messages_sent']
10
11 # Handle missing values
12 users[features] = users[features].fillna(users[features].mean())
13
14 # Normalize: Make all features same scale (0-1)
15 scaler = StandardScaler()
16 users_normalized = scaler.fit_transform(users[features])
17
18 print("Before:", users[features].iloc[0].values)
19 # [8.5, 45, 28, 156]
20 print("After:", users_normalized[0])
21 # [1.2, 0.8, 1.1, 0.9] - all similar scale!
```

Why normalize? So “hours used” doesn’t dominate “features used”

# Your First K-means in Python

## Just 5 Lines to Find User Groups!

```
1 from sklearn.cluster import KMeans
2 import matplotlib.pyplot as plt
3
4 # Create and fit K-means (let's find 4 groups)
5 kmeans = KMeans(n_clusters=4, random_state=42)
6 users['cluster'] = kmeans.fit_predict(users_normalized)
7
8 # See the groups
9 print(users.groupby('cluster')[features].mean())
```

### Cluster 0: Power Users

- 8.2 hours/day
- 52 features used

### Cluster 1: Casual Users

- 1.5 hours/day
- 8 features used

### Cluster 2: Regular Users

- 4.1 hours/day
- 25 features used

### Cluster 3: New Users

- 0.8 hours/day
- 3 features used

That's it! You've segmented thousands of users in seconds

# How Many Groups? The Elbow Method


## Finding the “Just Right” Number of Clusters

```
# Try different numbers of clusters
inertias = []
K_range = range(2, 11)

for k in K_range:
    kmeans = KMeans(n_clusters=k)
    kmeans.fit(users_normalized)
    inertias.append(kmeans.inertia_)

# Plot the elbow curve
plt.plot(K_range, inertias, 'bo-')
plt.xlabel('Number of Clusters')
plt.ylabel('Total Distance')
plt.title('The Elbow Method')
plt.show()
```

Look for the “elbow” - where adding more clusters doesn’t help much



In this example:

- 2-3 clusters: Big improvement
- 4-5 clusters: Good improvement
- 6+ clusters: Diminishing returns
- Choose: 4 or 5 clusters

# Creating Dendrograms (Tree Diagrams)


## Visualizing How Users Group Together

```
from scipy.cluster.hierarchy import dendrogram, linkage
import matplotlib.pyplot as plt

# Create the linkage matrix
linkage_matrix = linkage(users_normalized,
                         method='ward')

# Plot dendrogram
plt.figure(figsize=(10, 6))
dendrogram(linkage_matrix,
            labels=users['user_id'].values,
            leaf_rotation=90)
plt.title('User Clustering Dendrogram')
plt.xlabel('User ID')
plt.ylabel('Distance')
plt.show()

# Cut tree to get 4 clusters
from scipy.cluster.hierarchy import fcluster
users['h_cluster'] = fcluster(linkage_matrix,
                               t=4,
                               criterion='maxclust')
```



### Reading the Tree:

- Bottom: Individual users
- Height: How different groups are
- Branches: Groups forming
- Cut line: Your chosen clusters

# Finding Dense Areas with DBSCAN

## When Your Groups Aren't Round


```
from sklearn.cluster import DBSCAN

# DBSCAN: Finds dense regions
dbscan = DBSCAN(eps=0.5,           # neighborhood size
                 min_samples=5) # min points
users['db_cluster'] = dbscan.fit_predict(
    users_normalized)

# Check results
print(f"Found {len(set(users['db_cluster']))-1} clusters")
print(f"Outliers: {sum(users['db_cluster']==-1)}")

# Visualize
colors = ['red', 'blue', 'green', 'yellow', 'purple']
for i in range(max(users['db_cluster'])+1):
    if i == -1: # Outliers
        plt.scatter(X[users['db_cluster']==i, 0],
                     X[users['db_cluster']==i, 1],
                     c='gray', marker='x', alpha=0.3)
    else:
        plt.scatter(X[users['db_cluster']==i, 0],
                     X[users['db_cluster']==i, 1],
                     c=colors[i], alpha=0.6)
```

## DBSCAN Advantages:



- Finds any shape clusters
- Identifies outliers (noise)
- No need to specify K
- Great for unusual patterns

# Choosing the Right Features

## What to Measure for Good Clustering

### Good Features for User Clustering:


- Usage frequency
- Feature adoption
- Time patterns
- Interaction types
- Content preferences

### Avoid These:

- User ID (unique)
- Registration date (if not relevant)
- Random identifiers
- Highly correlated features

### Feature Engineering Example:

Raw Data	Engineered Feature
Login times	Morning/Evening user
Click events	Clicks per session
Page views	Depth of exploration
Purchase history	Spending tier
Support tickets	Frustration level



# Making Clusters Visible (2D Plots)


## Reducing Dimensions to See Patterns

```
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

# Reduce to 2D for visualization
pca = PCA(n_components=2)
users_2d = pca.fit_transform(users_normalized)

# Plot with cluster colors
plt.figure(figsize=(10, 8))
colors = ['#e74c3c', '#3498db', '#2ecc71', '#f39c12']
for i in range(4):
    cluster_data = users_2d[users['cluster'] == i]
    plt.scatter(cluster_data[:, 0],
                cluster_data[:, 1],
                c=colors[i],
                label=f'Cluster {i}',
                alpha=0.6, s=50)

plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('User Clusters Visualization')
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()
```



### What PCA Does:

- Combines all features into 2D
- Preserves most important differences
- Makes patterns visible
- Great for presentations!

# Working with Categories

## Clustering with Non-Numeric Data

```
# Convert categories to numbers
from sklearn.preprocessing import LabelEncoder

# Example: Device type
le = LabelEncoder()
users[‘device_encoded’] = le.fit_transform(
    users[‘device_type’])
# ‘mobile’ -> 0, ‘desktop’ -> 1, ‘tablet’ -> 2

# Better: One-hot encoding for clustering
device_dummies = pd.get_dummies(
    users[‘device_type’],
    prefix=‘device’)
# Creates: device_mobile, device_desktop, device_tablet

# Combine with numeric features
features_all = pd.concat([
    users[numeric_features],
    device_dummies,
    pd.get_dummies(users[‘subscription_type’])])
], axis=1)

# Now cluster as normal
kmeans = KMeans(n_clusters=4)
users[‘cluster’] = kmeans.fit_predict(features_all)
```

### Category Examples:

- Device type
- Subscription level
- Country/Region
- Product categories
- User role

**Rule:** Use one-hot encoding for clustering,  
not label encoding


# Real Example: E-commerce Customer Segmentation

## Finding Customer Types in Online Shopping Data

```
1 # Real e-commerce features
2 customer_features = [
3     'total_spent', 'order_frequency', 'avg_cart_size',
4     'categories_browsed', 'return_rate', 'review_count'
5 ]
6
7 # Cluster and analyze
8 kmeans = KMeans(n_clusters=5)
9 customers['segment'] = kmeans.fit_predict(customers_scaled)
10
11 # Results interpretation
12 for i in range(5):
13     segment = customers[customers['segment'] == i]
14     print(f"\nSegment {i}: {len(segment)} customers")
15     print(segment[customer_features].mean())
```

### Discovered Segments:

- **VIP Shoppers:** High spend, low returns
- **Bargain Hunters:** Sale-focused, high cart
- **Window Shoppers:** Browse, rarely buy
- **Loyal Regulars:** Consistent, medium spend
- **One-timers:** Single purchase, dormant



# Interactive Demo: Try It Yourself!

## Complete Clustering Pipeline

```
1 # Complete working example you can run
2 import pandas as pd
3 import numpy as np
4 from sklearn.cluster import KMeans
5 from sklearn.preprocessing import StandardScaler
6 import matplotlib.pyplot as plt
7
8 # Generate sample data (replace with your data)
9 np.random.seed(42)
10 n_users = 1000
11 data = {
12     'usage_hours': np.random.exponential(3, n_users),
13     'features_used': np.random.poisson(15, n_users),
14     'days_active': np.random.randint(1, 31, n_users)
15 }
16 df = pd.DataFrame(data)
17
18 # Standardize
19 scaler = StandardScaler()
20 X_scaled = scaler.fit_transform(df)
21
22 # Cluster
23 kmeans = KMeans(n_clusters=3, random_state=42)
24 df['cluster'] = kmeans.fit_predict(X_scaled)
25
26 # Visualize and interpret
27 print(df.groupby('cluster').mean())
28 # Try changing n_clusters and see what happens!
```

# From Data Clusters to User Personas

## Turning Numbers into People

### Cluster Statistics:

Metric	Cluster 0
Avg. Usage	7.2 hrs/day
Features Used	45/50
Peak Time	9am-5pm
Device	Desktop (85%)
Retention	95%

### Persona Created:

**“Professional Paula”**

*Power user, 32, Marketing Manager*

**Goals:** Maximize productivity

**Needs:** Advanced features, shortcuts

**Pain:** Slow load times

**Quote:** “This tool is my office”


## Transformation Process:

- ① Analyze cluster statistics
- ② Identify defining characteristics
- ③ Create realistic profile
- ④ Add human elements (name, photo, quote)
- ⑤ Validate with real user interviews

# Creating Empathy Maps from Data

## What Your Clusters Think, Feel, Say, and Do

### Empathy Map: Data-Driven User Understanding



# Finding User Pain Points

## Where Each Cluster Struggles


### Pain Point Detection Methods:

- High exit rates at specific features
- Support ticket clustering
- Feature abandonment patterns
- Error message frequency
- Negative sentiment spikes

### Cluster-Specific Pain Points:


- **Power Users:** Need bulk operations
- **New Users:** Overwhelming interface
- **Mobile Users:** Desktop-only features
- **Free Users:** Paywall friction

Different clusters = Different problems = Different solutions needed



# Discovering Behavior Patterns


## How Different Clusters Use Your Product



# Building User Journey Maps

## Each Cluster's Path Through Your Product

User Journey Maps by Cluster




# Identifying Stakeholder Groups

## Who Really Uses Your Product?

### Discovered Stakeholders:

- **Decision Makers (5%)**
  - Admin features
  - Billing pages
  - Team management
- **Daily Users (60%)**
  - Core features
  - Regular patterns
  - Productivity focus
- **Influencers (15%)**
  - Share features
  - Invite others
  - Write reviews
- **Evaluators (20%)**
  - Trial users
  - Comparison shoppers
  - Feature testers

Stakeholder Network from Cluster Analysis




Network shows how different groups interact and influence each other

# Visual Persona Cards

## Making Clusters Memorable and Actionable

### Data-Driven Persona Cards



Each Card Includes:

BSc Data Science & Design

Week 1: Foundation + Clustering

2025

27 / 1

# Making Design Decisions from Clusters

## From Insights to Action


### Cluster-Driven Decisions:

Finding	Design Decision
3 skill levels found	Progressive disclosure UI
Mobile vs Desktop split	Responsive-first design
Power users frustrated	Advanced mode option
New users confused	Better onboarding
Social cluster exists	Add sharing features

### Prioritization Framework:

- ① Size of cluster (impact)
- ② Pain intensity (urgency)
- ③ Business value (ROI)
- ④ Implementation cost (feasibility)

Design for your biggest, most valuable, or most struggling clusters first



# Other Clustering Methods Overview

## Beyond K-means: More Tools in Your Toolkit

### Gaussian Mixture Models

- Soft clustering (probability-based)
- Handles overlapping groups
- Good for uncertain boundaries

### Mean Shift

- Finds density peaks automatically
- No need to specify K
- Great for image segmentation


### Spectral Clustering

- Handles complex shapes
- Uses graph theory
- Good for social networks

### OPTICS

- Like DBSCAN but better
- Handles varying densities
- Creates reachability plots

Clustering Methods Comparison



# Industry Case Study: Spotify's Discover Weekly

## How Clustering Powers Personalized Playlists

### The Challenge:


- Hundreds of millions of users
- Tens of millions of songs
- Create unique playlists weekly
- Feel personally curated

### The Solution:

- ① Cluster users by listening history
- ② Cluster songs by audio features
- ③ Find songs your cluster likes that you haven't heard
- ④ Mix in variety from adjacent clusters
- ⑤ Result: 30 new songs every Monday

Clustering at scale: From understanding users to delighting them

Spotify's Discover Weekly: Clustering in Action



### Impact:

- Significant user engagement
- Major impact on listening behavior
- Transformed music discovery

## What You've Learned

### Technical Skills:

- K-means clustering implementation
- Hierarchical clustering with dendrograms
- DBSCAN for density-based groups
- Data preparation and scaling
- Cluster evaluation methods
- Visualization techniques

### Design Skills:

- Creating data-driven personas
- Building empathy maps
- Identifying pain points
- Journey mapping
- Stakeholder identification

### Key Insights:

- Clustering reveals hidden user groups
- Different algorithms for different data
- Always validate with business sense
- Clusters drive design decisions
- Scale empathy with data

**Remember:** Clustering is about understanding, not just grouping

# Preview: Week 2


## Advanced Clustering + Deep Empathy

### Next Week You'll Learn:

- Time-series clustering for behavior evolution
- Multi-view clustering (combining data sources)
- Online clustering for real-time segmentation
- Clustering validation techniques
- A/B testing with clusters
- Emotional journey mapping
- Micro-moment identification
- Cluster-based personalization

### Practical Project:

Build a complete user segmentation system for a real app



### Homework:

- Practice K-means on your data
- Create one persona from a cluster
- Read: Chapter 2 materials

See you next week for deeper dives into clustering!

## Core Algorithms:

- MacQueen, J. (1967). Some methods for classification and analysis of multivariate observations
- Ester et al. (1996). A density-based algorithm for discovering clusters (DBSCAN)
- Lloyd, S. (1982). Least squares quantization in PCM (K-means)

## Design Thinking Integration:

- Brown, T. (2009). Change by Design: How Design Thinking Transforms Organizations
- IDEO Design Thinking Toolkit

## Tools & Libraries:

- scikit-learn: Machine Learning in Python
- matplotlib & seaborn: Visualization libraries
- Course repository: [github.com/ml-design-thinking](https://github.com/ml-design-thinking)

## Distance Metrics Formulas

### Euclidean Distance:

$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

### Manhattan Distance:

$$d(x, y) = \sum_{i=1}^n |x_i - y_i|$$

### Cosine Similarity:

$$\text{similarity}(x, y) = \frac{x \cdot y}{\|x\| \cdot \|y\|} = \frac{\sum_{i=1}^n x_i y_i}{\sqrt{\sum_{i=1}^n x_i^2} \cdot \sqrt{\sum_{i=1}^n y_i^2}}$$

### Minkowski Distance:

$$d(x, y) = \left( \sum_{i=1}^n |x_i - y_i|^p \right)^{1/p}$$

# Appendix: K-means Algorithm Details

## Complete Algorithm Specification

### Objective Function (Minimize):

$$J = \sum_{i=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2$$

where  $\mu_i$  is the mean of cluster  $C_i$

### Algorithm Steps:

- ① Initialize: Choose  $k$  points as initial centroids
- ② Assignment:  $C_i = \{x_p : \|x_p - \mu_i\|^2 \leq \|x_p - \mu_j\|^2 \forall j, 1 \leq j \leq k\}$
- ③ Update:  $\mu_i = \frac{1}{|C_i|} \sum_{x \in C_i} x$
- ④ Repeat until convergence:  $\|\mu_i^{(t+1)} - \mu_i^{(t)}\| < \epsilon$

**Complexity:**  $O(n \cdot k \cdot d \cdot i)$  where  $n$  = points,  $k$  = clusters,  $d$  = dimensions,  $i$  = iterations

# Appendix: Silhouette Coefficient Calculation

## Cluster Quality Metric

For each point  $i$ :

- $a(i)$  = average distance to points in same cluster
- $b(i)$  = minimum average distance to points in different cluster

## Silhouette Coefficient:

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$$

## Interpretation:

- $s(i) \approx 1$ : Well clustered
- $s(i) \approx 0$ : On border between clusters
- $s(i) \approx -1$ : Misclassified

## Overall Score:

$$S = \frac{1}{n} \sum_{i=1}^n s(i)$$

## Probabilistic Clustering Model:

$$p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$$

where:

- $\pi_k$  = mixing coefficient (prior probability)
- $\mu_k$  = mean of component  $k$
- $\Sigma_k$  = covariance matrix of component  $k$

## EM Algorithm:

- **E-step:** Compute responsibilities

$$\gamma_{nk} = \frac{\pi_k \mathcal{N}(x_n|\mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n|\mu_j, \Sigma_j)}$$

- **M-step:** Update parameters

$$\mu_k = \frac{\sum_n \gamma_{nk} x_n}{\sum_n \gamma_{nk}}$$

## t-Distributed Stochastic Neighbor Embedding

High-dimensional similarity:

$$p_{j|i} = \frac{\exp(-\|x_i - x_j\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|x_i - x_k\|^2 / 2\sigma_i^2)}$$

Low-dimensional similarity (Student t-distribution):

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}}$$

Objective (KL divergence):

$$C = \sum_i KL(P_i || Q_i) = \sum_i \sum_j p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

Gradient:

$$\frac{\partial C}{\partial y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j)(1 + \|y_i - y_j\|^2)^{-1}$$

# Appendix: Computational Complexity Analysis

## Time and Space Complexity of Clustering Algorithms

Algorithm	Time Complexity	Space Complexity
K-means	$O(n \cdot k \cdot d \cdot i)$	$O(n \cdot d + k \cdot d)$
Hierarchical	$O(n^2 \log n)$ to $O(n^3)$	$O(n^2)$
DBSCAN	$O(n \log n)$ average	$O(n)$
GMM	$O(n \cdot k \cdot d^2 \cdot i)$	$O(n \cdot d + k \cdot d^2)$
Spectral	$O(n^3)$	$O(n^2)$
Mean Shift	$O(n^2 \cdot i)$	$O(n \cdot d)$

### Legend:

- $n$  = number of data points
- $k$  = number of clusters
- $d$  = dimensionality
- $i$  = number of iterations

### Scalability Tips:

- Use Mini-batch K-means for  $n > 10,000$
- Consider sampling for hierarchical clustering
- Use approximate nearest neighbors for DBSCAN

## Key Papers and Resources

### Foundational Papers:

- MacQueen, J. (1967). "Some methods for classification and analysis of multivariate observations"
- Ester et al. (1996). "A density-based algorithm for discovering clusters" (DBSCAN)
- Ng et al. (2002). "On spectral clustering: Analysis and an algorithm"

### Modern Applications:

- Sculley, D. (2010). "Web-scale k-means clustering" (Google)
- McInnes et al. (2017). "hdbscan: Hierarchical density based clustering"
- Spotify Research. "Understanding Music through Machine Learning"

### Online Resources:

- scikit-learn clustering documentation
- Google's Machine Learning Crash Course
- Fast.ai Practical Deep Learning course