

# **NLP for Emotional Context**

From Words to Understanding What Users Really Feel

Week 3: Machine Learning for Smarter Innovation

Transform 50,000 Reviews into Actionable Design Insights

## Four Stages of Understanding

1. **The Challenge** - Why understanding emotion in text is impossibly hard
2. **First Solution & Its Limits** - Traditional NLP works... until it doesn't
3. **The Breakthrough** - How Transformers changed everything
4. **Design Synthesis** - From ML insights to user empathy

**Core Question:** How can we understand the emotions of 50,000 users without reading 5 million words?

---

Large-scale emotional analysis enables mass personalization - automated sentiment understanding processes millions of user expressions beyond manual capacity

## Which Reviews Reveal Why Users Really Quit?

### The Scenario You Face:

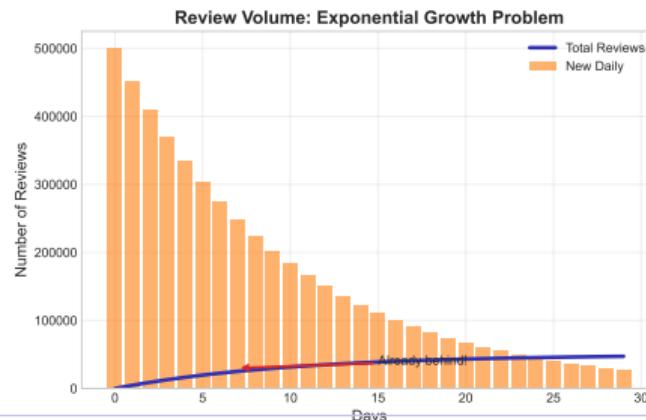
- Your product has 50,000 user reviews
- Average review: 100 words
- Total: 5 million words to process
- Hidden inside: The 10 reviews that explain 80% of churn

### The Human Limit:

- Reading speed: 200 words/minute
- Time to read all: 417 hours (10 weeks!)
- Reviews arrive: 500 new ones daily
- You're already behind before you start

### Real Review Examples:

- "Love it but..." (quit in 2 weeks)
- "Not bad for the price" (renewed 3 years)
- "Just what I expected!" (1-star rating)
- "Finally someone gets it" (5-star champion)



Data velocity exceeds human processing capacity - automated analysis becomes necessity rather than optimization

# Building Block 1: Words Don't Have Fixed Meanings

## Context Changes Everything

Same Word, Different Meanings:

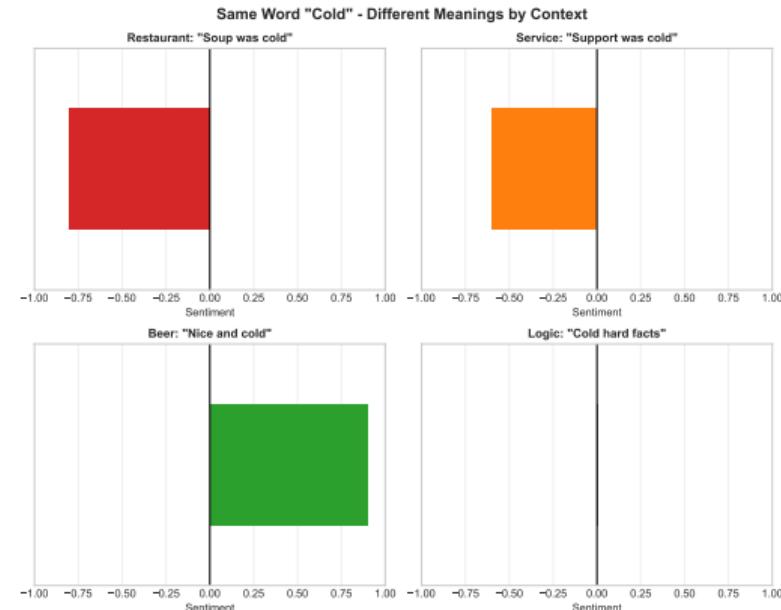
“Cold”

- Restaurant: “The soup was cold” → Negative
- Service: “Support was cold” → Negative
- Beer: “Nice and cold” → Positive
- Logic: “Cold hard facts” → Neutral

“Fast”

- Delivery: “Super fast!” → Positive
- Battery: “Dies too fast” → Negative
- Customer service: “Too fast, felt rushed” → Negative

The Computer’s Problem:



**Key Insight:** A word's emotional meaning depends on ALL surrounding words, not just the word itself.

### What People Say What They Feel

#### Layer 1: Literal

- "Great product"
- Clear positive
- Easy to detect
- 15% of reviews

#### Example:

"This is excellent. I love every feature. Will buy again."

→ Genuinely positive

#### Layer 2: Sarcastic

- "Just wonderful"
- Opposite meaning
- Context reveals truth
- 25% of reviews

#### Example:

"Oh great, another update that breaks everything"

→ Actually negative

#### Layer 3: Cultural

- "It's okay"
- Varies by culture
- UK: Negative
- US: Neutral

#### Example:

"It does what it says"

UK: Disappointed

US: Satisfied


**The Complexity:** 60% of emotional meaning is NOT in the words themselves but in how they're used together

---

Interpretive complexity limits consensus - emotional content carries inherent ambiguity even among expert human evaluators

## Why This Problem Explodes in Complexity

Complexity Explosion: 600,000 Combinations



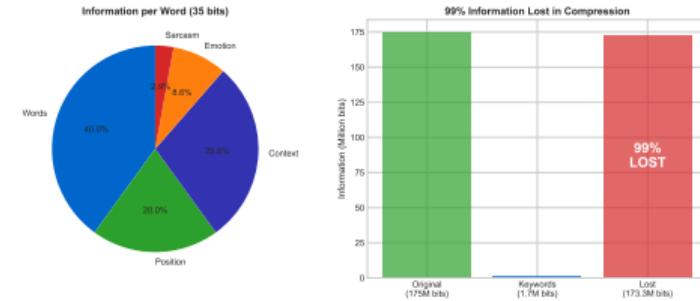
## The Mathematical Impossibility

### Information Content Analysis:

Component	Bits	Information
One word	14 bits	Which of 10,000 words
+ Position	7 bits	Where in sentence
+ Emotion	3 bits	8 emotion types
+ Context	10 bits	Surrounding words
+ Sarcasm	1 bit	Yes/No
<b>Total per word</b>	<b>35 bits</b>	
<b>Per review (100 words)</b>	<b>3,500 bits</b>	
<b>50,000 reviews</b>	<b>175M bits</b>	<b>22 MB</b>

### Traditional Compression:

- Keyword counts: 10,000 → 100 words
- Compression ratio: 100:1
- Information lost: 99%



### The Dilemma:

Compress too much → Lose emotional nuance

Keep everything → Impossible to process

**Need:** Selective preservation of emotional signal

## Can we preserve emotion while compressing 100:1?

Information-theoretic limits constrain compression - high-entropy signals require proportionally complex representations to preserve fidelity

## How Would YOU Quickly Scan 1000 Reviews?

### Human Strategy:

1. Look for emotion words: "love", "hate", "terrible"
2. Count positive vs negative
3. More positive → Happy customer
4. More negative → Unhappy customer

### Computer Implementation:

- Build word lists
- Positive: [great, excellent, love, amazing...]
- Negative: [bad, terrible, hate, awful...]
- Count occurrences
- Calculate: Positive% - Negative%


### Example Analysis:

"I **love** the design but **hate** waiting. The quality is **excellent** though shipping was **terrible**."

**Count:** 2 positive, 2 negative

**Score:**  $50\% - 50\% = \text{Neutral (0)}$

Seems reasonable!



Simplicity enables initial progress - crude approximations provide baseline performance that reveals limitation patterns

# Bag of Words: Converting Text to Numbers

## The Classic Approach in Action

### Step-by-Step Process:

#### 1. Original Review:

"The product quality is excellent excellent excellent but customer service terrible terrible"


#### 2. Tokenize (split into words):

The, product, quality, is, excellent, excellent, excellent, but, customer, service, terrible, terrible

	Word	Count		Word	Count
3. Count Each Word:	excellent	3		terrible	2
	product	1		service	1
	quality	1		customer	1

#### 4. Convert to Vector:

1, 1, 1, 3, 0, 0, 0, 1, 1, 2, 0, ...  
(10,000 dimensions, mostly zeros)



### What We Keep:

- Word frequencies
- Vocabulary presence

### What We Lose:

- Word order
- Grammar

## Not All Words Are Equal

### The Problem with Raw Counts:

- "The" appears 1000 times → Important?
- "Revolutionary" appears once → Not important?
- Common words dominate
- Rare but meaningful words ignored

### TF-IDF Solution:

$$\text{TF-IDF} = \underbrace{\frac{\text{count in doc}}{\text{total words}}}_{\text{Term Frequency}} \times \log \underbrace{\frac{\text{total docs}}{\text{docs with word}}}_{\text{Inverse Document Freq}}$$

- TF: How often in THIS review
- IDF: How rare across ALL reviews
- Product: Important AND distinctive

### Example Calculation:

Word	TF	IDF	TF-IDF
"the"	0.15	0.01	0.0015
"product"	0.05	0.69	0.0345
"excellent"	0.10	1.38	0.138
"revolutionary"	0.02	3.40	0.068

**Result:** Meaningful words now weighted higher than common words!

Rarity signals importance - terms appearing selectively carry more discriminative power than ubiquitous vocabulary

## When Simple Reviews Get Perfect Scores

Review Text	Human	BoW+TFIDF	Match
"This product is absolutely fantastic! Best purchase ever!"	Positive	Positive (95%)	
"Terrible quality. Waste of money. Very disappointed."	Negative	Negative (92%)	
"Good product, fair price, happy with purchase"	Positive	Positive (88%)	
"Broken on arrival. Customer service un- helpful. Never again."	Negative	Negative (94%)	
"Excellent! Exceeded all expectations! Highly recommend!"	Positive	Positive (97%)	

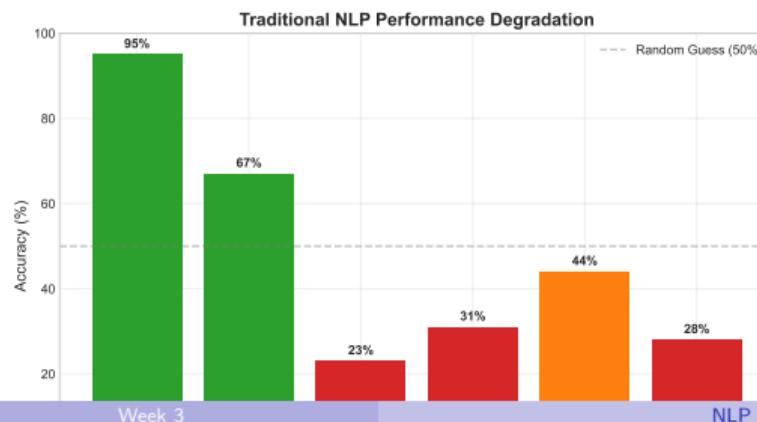
**Average Accuracy: 93.2%**

This is why Bag of Words dominated for 40 years!

Explicit signals enable simple methods - literal expression reduces analytical complexity compared to implicit communication

## Performance Degradation with Complexity

Review Type	Example	Human	BoW	Accuracy
Simple & Direct	"Great product!"	Pos	Pos	95%
Mixed Sentiment	"Good but overpriced"	Neg	Pos	67%
Sarcasm	"Oh great, it broke. Just perfect!"	Neg	Pos	23%
Context Dependent	"Not bad for the price"	Pos	Neg	31%
Subtle Emotion	"It's fine, I guess"	Neg	Neu	44%
Negation	"Not the worst I've seen"	Neu	Neg	28%



### The Pattern:

- Simple: 95% → Works great!
- Real-world: 44% → Worse than coin flip
- Sarcasm: 23% → Actively wrong

Average: 51% (Random guessing: 50%)

## Tracing the Failure Through Real Examples

Example: “Not bad for the price”

What BoW Sees:

- Words: [not, bad, for, the, price]
- “bad” → Negative word (-1)
- “not” → Negation word (ignored)
- Score: Negative

What Got Lost:

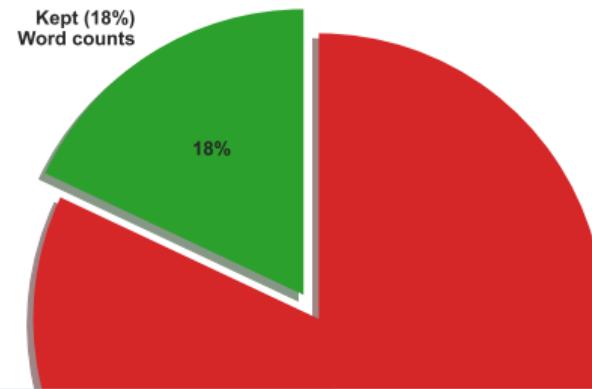
- “not bad” = Actually positive
- “for the price” = Qualified satisfaction
- Relationship between words
- Overall: Mild positive sentiment

**Root Cause:** Word order contains meaning!

### Information Loss Calculation:

Information Type	Bits Lost
Word positions	420 bits
Word relationships	1,200 bits
Grammar structure	350 bits
Contextual meaning	890 bits
<b>Total Lost</b>	<b>2,860 bits</b>
<b>Total Kept</b>	<b>640 bits</b>
<b>Kept Percentage</b>	<b>18%</b>

Information Preserved in Bag of Words



## Stop: How Do YOU Understand “The Bank is Nice”?

### Let's Be Honest About Our Mental Process

#### Trace Your Thinking:

Sentence: “I went to the bank. The water was nice.”

1. Read “I went to the bank”
2. Your brain: Could be financial or river...
3. Read “The water was nice”
4. Your brain: Ah! River bank, not financial
5. Re-interpret: “bank” = riverbank
6. Understand: Person enjoyed the riverside

#### The Key Observation:

You DON'T compress the sentence into a summary.

You KEEP all words available and SELECTIVELY ATTEND to relevant ones when needed.

This is the breakthrough insight!

#### What You Actually Did:

- Held multiple meanings open
- Used later words to resolve earlier ones
- Selectively focused on “water” to understand “bank”

Human Attention: “water” helps understand “bank”

I      went      to      the      bank      .      The      water      was      nice      .



## A Fundamentally Different Approach

### Old Way (Bag of Words):

- Take all words
- Compress to counts
- Lose relationships
- Try to reconstruct meaning

### Analogy:

Like taking a book, counting word frequencies, throwing away the book, then trying to understand the story from counts.

**Result:** Lost 82% of information  
Can't recover context

### New Way (Attention):

- Keep all words
- Keep all positions
- For each word, look at all others
- Decide how much to focus on each

### Analogy:

Like keeping the entire book and using a smart highlighter that knows which sentences help understand other sentences.

**Result:** Keep 100% of information  
Use what's relevant when needed

**Instead of asking “What to keep?” we ask “What to focus on?”**

---

Paradigm shifts replace constraints with capabilities - architectural innovations enable previously impossible approaches

# How Attention Works: Focus Percentages

## No Math Yet - Just Percentages

Example: Understanding "bank" in: "The bank by the river is peaceful"

### Step 1: For word "bank", look at all other words

Word	Relevance to "bank"
The	Low
by	Medium
the	Low
<b>river</b>	<b>Very High</b>
is	Low
peaceful	Medium

### Step 2: Convert to percentages (must sum to 100%)

Word	Focus %
The	5%
by	15%
the	5%
<b>river</b>	<b>55%</b>
is	5%
peaceful	15%

### Step 3: Use these percentages to understand "bank"

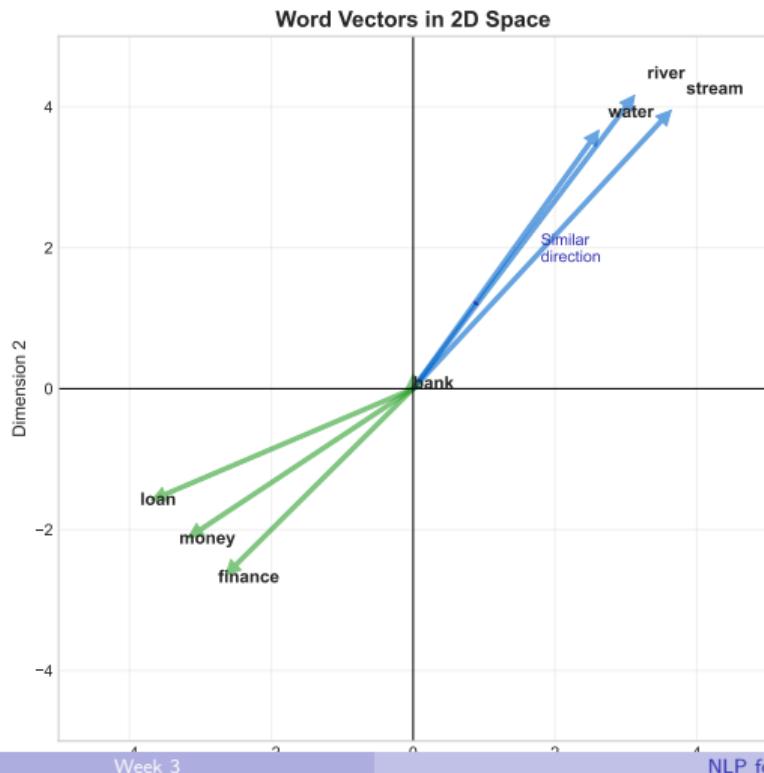
"bank" meaning =  
5% × meaning("The") +  
15% × meaning("by") +  
5% × meaning("the") +  
**55% × meaning("river")** +  
5% × meaning("is") +  
15% × meaning("peaceful")  
= Mostly "river" context  
= Riverbank, not financial bank!

These percentages ARE the attention weights!

Normalized weighting enables comparison - probability distributions force explicit trade-offs among competing information sources

## Words as Directions in Space

Start Simple: 2D Space



Scale to Real NLP: 768D Space

Same principle, more dimensions:

- 2D:  $[x, y]$  coordinates
- 768D:  $[x, x, \dots, x]$  coordinates

More dimensions = More nuanced meaning  
Can distinguish "bank (river)" from "bank (money)" from "bank (slope)"

The Attention Calculation:

1. Compute alignment (dot product) with all words
2. Higher alignment = Pay more attention
3. Convert to percentages
4. Use percentages to blend meanings

Geometry gives us semantic similarity!

# The Attention Algorithm: Three Essential Steps

## Each Step Has a Purpose

### Step 1: SCORE - Find Relevance

- Action: For each word, compute alignment with all others
- Why: Identify which words help understand this one
- Math:  $\text{Score}_{ij} = Q_i \cdot K_j$
- Plain: How relevant is word  $j$  to understanding word  $i$ ?

### Step 2: NORMALIZE - Create Percentages

- Action: Convert scores to probabilities (0-100%)
- Why: Need weights that sum to 1.0 for averaging
- Math:  $\alpha_{ij} = \frac{e^{\text{Score}_{ij}}}{\sum_k e^{\text{Score}_{ik}}}$
- Plain: What percentage of attention should word  $j$  get?

### Step 3: AGGREGATE - Blend Information

- Action: Weighted sum of all word meanings
- Why: Combine context based on attention weights
- Math:  $\text{Output}_i = \sum_j \alpha_{ij} \cdot V_j$
- Plain: New understanding using focused context  
Parallelization scales computational throughput - simultaneous processing across elements eliminates sequential bottlenecks

Attention Mechanism: Three Essential Steps



### The Complete Formula:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Where:

- $Q$  = Query (what am I looking for?)
- $K$  = Keys (what info is available?)
- $V$  = Values (what to extract?)
- $\sqrt{d_k}$  = Scaling factor for stability

## Real Numbers, Real Calculation

Task: Computing attention for “nice” in: “The service was nice”

Given: Word Vectors (Simplified to 2D)

The = [1, 0], service = [2, 3], was = [0, 1], nice = [3, 2]

Step 1: Calculate Scores

Query (nice) = [3, 2]

Word	Key	Q-K	Score
The	[1, 0]	$3 \times 1 + 2 \times 0$	3
service	[2, 3]	$3 \times 2 + 2 \times 3$	12
was	[0, 1]	$3 \times 0 + 2 \times 1$	2
nice	[3, 2]	$3 \times 3 + 2 \times 2$	13

Step 2: Convert to Percentages (Softmax)

Word	e^Score	Attention %
The	20	0.7%
service	162,755	63.1%
was	7	0.0%
nice	442,413	36.2%

Step 3: Aggregate

Values: The=[1,1], service=[4.5], was=[1,2], nice=[5,4]

Output for “nice” =

$$0.007 \times [1,1] +$$

$$\textcolor{green}{0.631} \times \textcolor{blue}{[4,5]} +$$

$$0.000 \times [1,2] +$$

$$0.362 \times [5,4]$$

$$= [0.01, 0.01] + \textcolor{green}{[2.52, 3.16]} + [0, 0] + [1.81, 1.45]$$

$$= [4.34, 4.62]$$

Result: 63% attention on “service”

The word “nice” is understood primarily in context of “service” → Customer service sentiment!

Low-dimensional principles generalize to high dimensions - mathematical operations remain identical regardless of vector size

# BERT: Reading in Both Directions Changes Everything

## The Power of Looking Forward AND Backward

### Traditional (Left-to-Right Only):

Sentence: "The bank charges are too high"

Reading "bank":

- Can see: "The"
- Cannot see: "charges are too high"
- Guess: Could be river or financial...
- **50% chance of error**

Bidirectional Understanding Changes Everything

Traditional (Left-to-Right): 50% chance of error



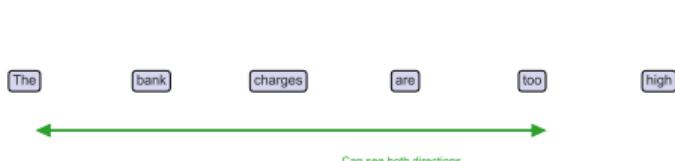
### BERT (Bidirectional):

Same sentence: "The bank charges are too high"

Reading "bank":

- Can see: "The" AND "charges are too high"
- "charges" → financial context
- Certain: Financial bank
- **99% accurate**

BERT (Bidirectional): 99% accurate



### BERT's Training Trick:

1. Mask random words: "The [MASK] is nice"
2. Predict masked word using ALL context
3. Must understand both directions to succeed

## Standing on the Shoulders of Human Knowledge

### BERT's Pre-training Data:

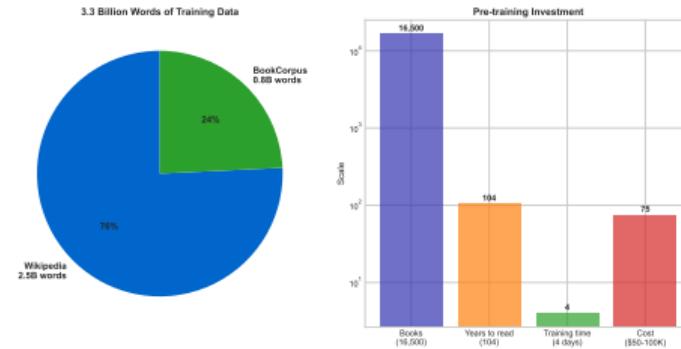
- Wikipedia: 2.5 billion words
- BookCorpus: 800 million words
- Total: 3.3 billion words
- Equivalent: 16,500 books (200 pages each)
- Reading time: 104 years non-stop

### What BERT Learned:

- Language patterns
- World knowledge
- Cultural contexts
- Emotional expressions
- Sarcasm patterns
- Domain vocabularies

### Training Cost:

- Time: 4 days on 64 TPUs
- Cost: \$50,000-100,000



### The Key Insight:

You DON'T need to train from scratch!

BERT already knows:

- Grammar
- Vocabulary
- Context patterns
- Emotional language

You just teach it YOUR specific task

## From General Knowledge to Your Reviews

### Your Fine-tuning Process:

#### 1. Start with Pre-trained BERT

- Has general language understanding
- Knows emotions, sarcasm, context
- But doesn't know YOUR products

#### 2. Prepare Your Data

- 1,000 labeled reviews
- Positive, Negative, Neutral labels
- Your product-specific language

#### 3. Fine-tune (Transfer Learning)

- Show BERT your reviews
- It adjusts its parameters slightly
- Learns your domain specifics
- Keeps general knowledge intact

#### 4. Time & Cost

- Time: 2-3 hours on GPU
- Cost: \$10-50 cloud compute
- Data needed: As few as 500 examples

### Example: Learning Company Slang

#### Before Fine-tuning:

"This app is sick!" → Negative (illness)

#### Your Training Data:

"Sick UI design!" → Positive

"Sick features!" → Positive

"Sick performance!" → Positive

#### After Fine-tuning:

"This app is sick!" → Positive (cool)

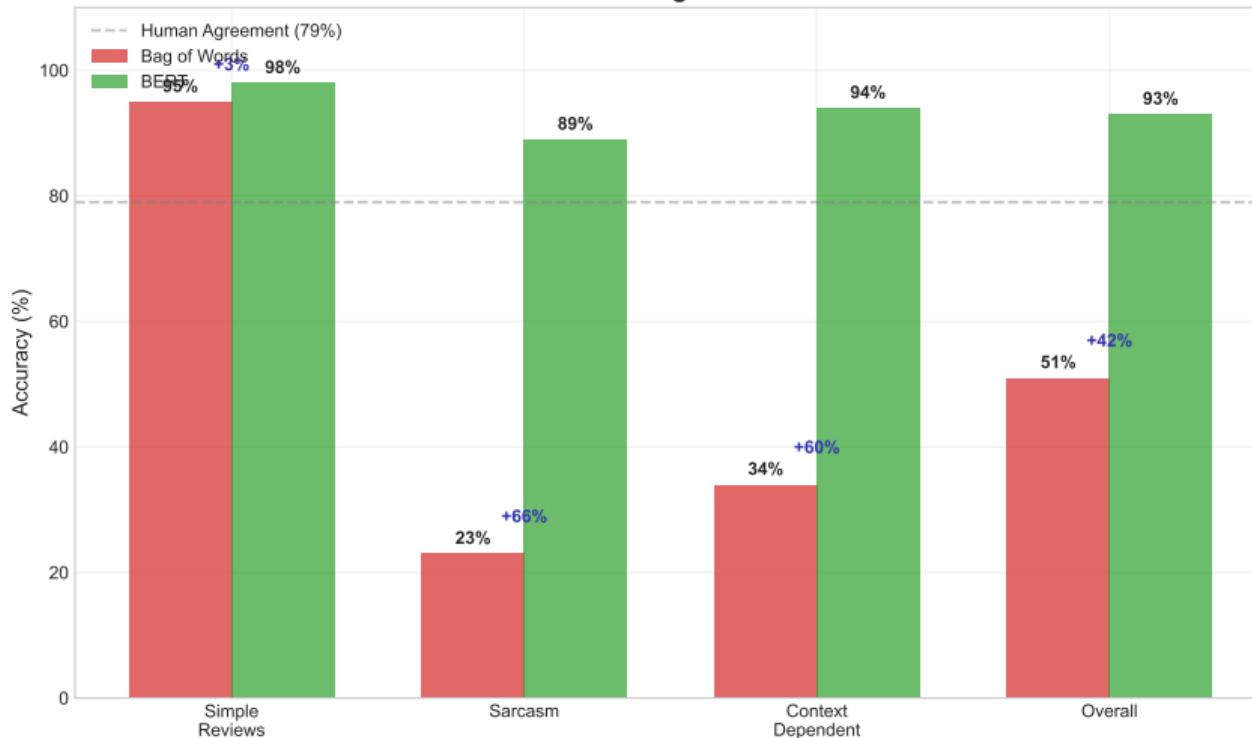
### Fine-tuning: Teaching BERT Your Specific Task




## The Results: From 34% to 94% Accuracy

### Transformers Solve the Context Problem

The Transformer Breakthrough: Context Problem Solved



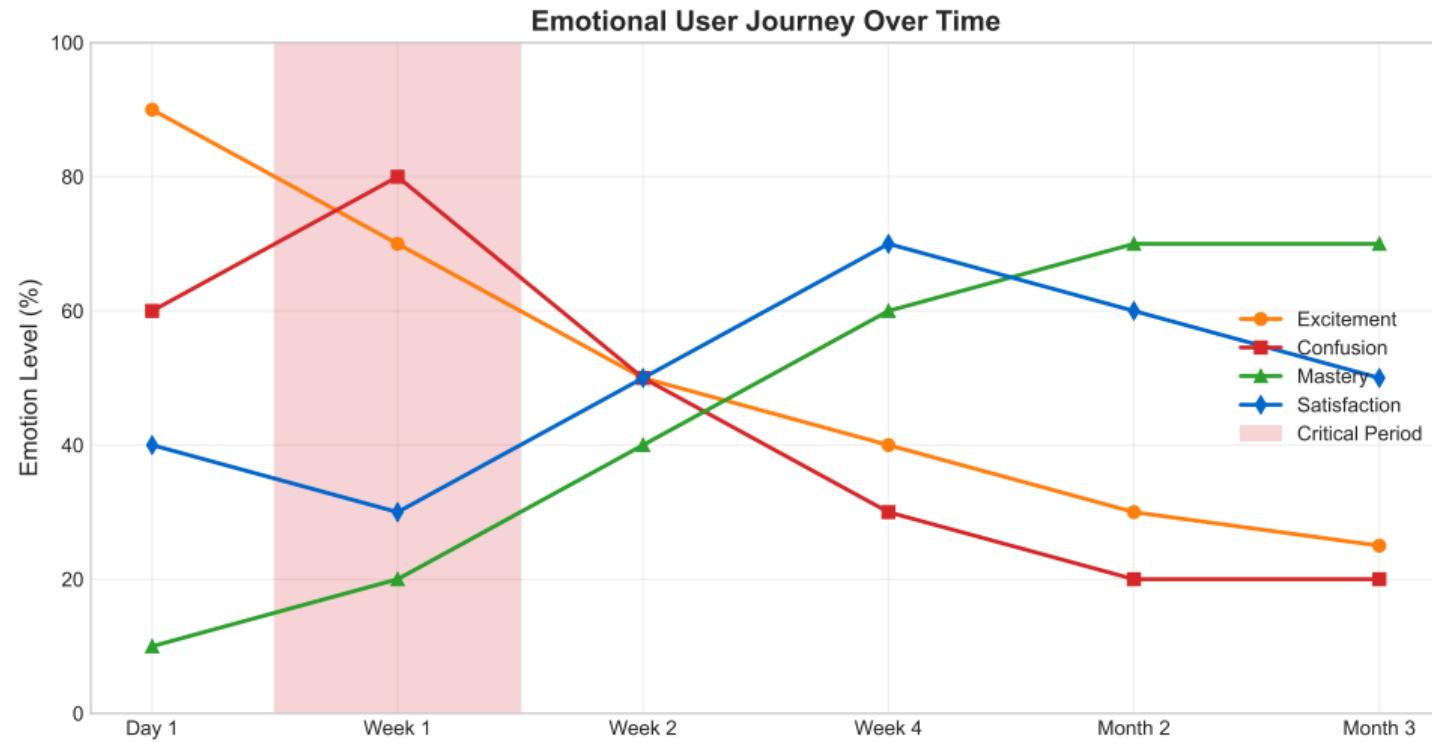
## What BERT Found in 50,000 Reviews



### 6 Core Emotion Clusters:

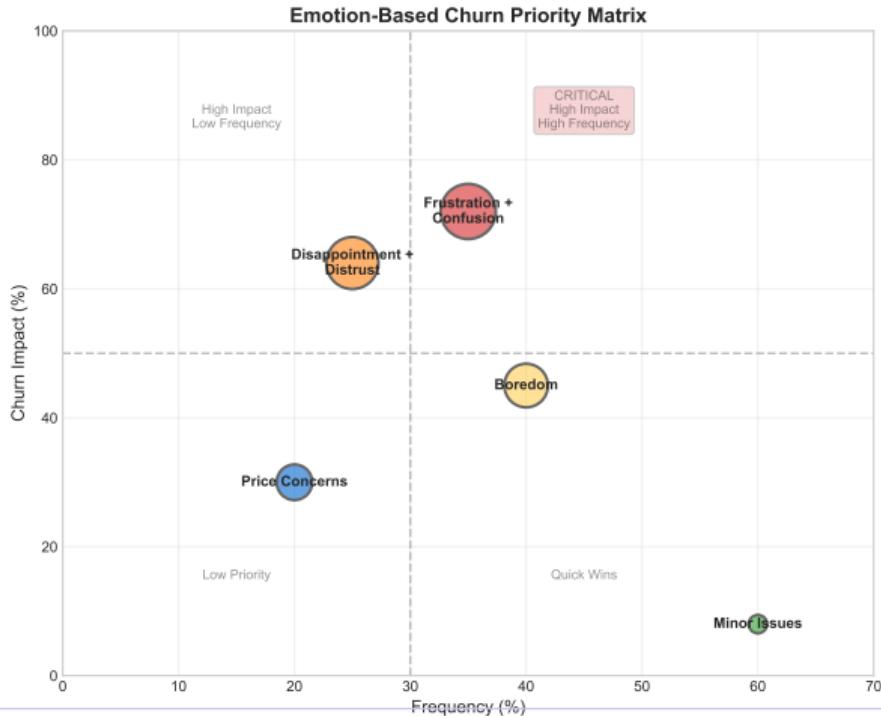
1. **Frustration (28%)**
  - Long wait times
  - Complex interfaces
  - Missing features
2. **Delight (22%)**
  - Unexpected features
  - Beautiful design
  - Fast performance
3. **Confusion (18%)**
  - Unclear instructions
  - Hidden functions
  - Inconsistent behavior
4. **Trust (15%)**
  - Data security
  - Reliable service
  - Transparent pricing
5. **Disappointment (12%)**
  - Unmet expectations
  - Quality issues
  - Broken promises
6. **Satisfaction (5%)**

## When and Why Emotions Shift



# The Churn Priority Matrix

## Which Emotions Predict User Loss?



Interaction effects dominate simple effects - isolated emotions mislead when combinations produce emergent meanings

### High Impact on Churn:

- 1. Frustration + Confusion (72% quit)**  
"Can't figure it out and support doesn't help"  
**Design Action:** Better onboarding + in-app help

- 2. Disappointment + Distrust (64% quit)**  
"Not what was promised, feels sketchy"  
**Design Action:** Align marketing with reality

### Low Impact on Churn:

- 3. Minor Frustrations (8% quit)**  
"Annoying but I deal with it"  
**Design Action:** Fix in regular updates

## Real Users, Not Imagined Ones

### The Enthusiast

15% of users

#### Language:

- "Love it!"
- "Game changer"
- "Can't wait for..."

#### Emotions:

- Delight: 78%
- Anticipation: 22%

#### Design Need:

Advanced features

### The Struggler

35% of users

#### Language:

- "Trying to..."
- "Can't find..."
- "How do I..."

#### Emotions:

- Confusion: 61%
- Frustration: 39%

#### Design Need:

Better guidance

### The Pragmatist

30% of users

#### Language:

- "It works"
- "Does the job"
- "Fair price"

#### Emotions:

- Satisfaction: 82%
- Neutral: 18%

#### Design Need:

Reliability

### The Critic

20% of users

#### Language:

- "Should have..."
- "Compared to X..."
- "Missing..."

#### Emotions:

- Disappointment: 54%
- Frustration: 46%

#### Design Need:

Feature parity

These personas emerged from BERT clustering - not designer assumptions

Language-based segmentation reveals latent groups - behavioral patterns emerge from expression style rather than demographics

## Personalized Understanding, Automated Delivery

### The Scale Challenge:

- 1 million users
- 100 reviews each
- 100 million opinions
- Impossible to read manually

### BERT's Solution:

- Process all in 24 hours
- Understand each individually
- Group by emotional need
- Generate targeted responses

### Personalization Examples:

- Frustrated user → Proactive support offer
- Confused user → Tutorial recommendation
- Delighted user → Feature beta invitation
- Disappointed user → Feedback survey

### Empathy at Scale: From Millions to One



Real Impact:

## Emotion-Driven Engagement at Scale

### The Challenge:


- 400 million users
- Make each feel special
- Drive social sharing
- Increase engagement

### NLP Analysis Revealed:

- Users want validation of taste
- Nostalgia drives sharing
- Uniqueness matters most
- Discovery excites users

### Design Response:

- Personal emotion words: "Your year was Rebellious"
- Unique statistics: "Top 0.5% of fans"
- Nostalgic moments: "You played X 47 times in March"
- Social proof: "Share your unique taste"



### Results:

- 120M users engaged
- 60M social shares
- 40% increase in app usage
- 21% increase in subscriptions

**ROI: 400% on NLP investment**

Emotion understanding → Personalization → Engagement → Business value

## Apply These Techniques to Real Data

Workshop Exercise (45 minutes):

Dataset:

- 5,000 app store reviews
- Your choice of app category
- Mix of ratings (1-5 stars)
- Real user language

Your Tasks:

1. Load pre-trained BERT model
2. Fine-tune on 500 labeled reviews
3. Analyze remaining 4,500 reviews
4. Discover emotion clusters
5. Identify top 3 pain points
6. Generate design recommendations

Tools Provided:

- Jupyter notebook template
- Pre-processed data

Expected Outputs:

1. **Emotion Distribution Chart**  
Show percentages of each emotion
2. **Pain Point Priority List**  
Ranked by impact on ratings
3. **User Segment Personas**  
Data-driven, not assumed
4. **Design Recommendations**  
Specific, actionable, prioritized

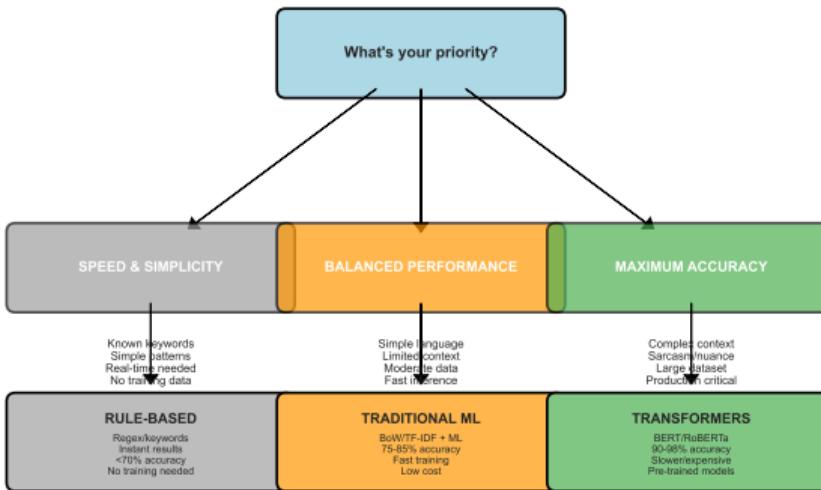
Learning Objectives:

- Use BERT for emotion analysis
- Interpret attention weights
- Convert ML insights to design actions
- Experience the power of scale

From 5,000 reviews to 5 key insights in 45 minutes!

# When to Use Which NLP Method: Judgment Criteria

## When to Use Which NLP Method: Decision Framework



### Additional Considerations

Data Volume: <3K samples - Rule-based or few-shot; 3K-10K - Traditional ML; >10K - Transformers viable  
Languages: Multi-lingual needs - Multilingual BERT (mBERT, XLM-RoBERTa); English only - simpler models  
Domain: Medical/legal - Fine-tune domain-specific transformer; General - Use pre-trained as-is  
Latency: Real-time (<100ms) - Rule-based or cached ML; Batch processing - Transformers acceptable  
Budget: Limited - Traditional ML (10-100x cheaper); Enterprise - Transformers for best results  
Explainability: High need - Rule-based (transparent) or LIME/SHAP on ML; Black box OK - Transformers

Principle: Start simple (rules/traditional ML), upgrade to transformers only when context/nuance critical

## Three Levels of Hands-on Learning

### Exercise 1: Basic Sentiment Analysis

Time: 20 minutes

Difficulty: Beginner

Task:

- Use pre-trained model
- Analyze 100 reviews
- Classify positive/negative
- Calculate accuracy

#### Learning Goal:

Understand basic NLP pipeline

Tools:

- TextBlob or
- Hugging Face pipeline

### Exercise 2: Intermediate Emotion Detection

Time: 45 minutes

Difficulty: Medium

Task:

- Fine-tune BERT
- 6-class emotions
- Analyze attention weights
- Visualize results

#### Learning Goal:

Work with transformers

Tools:

- Transformers library
- Pre-labeled dataset

### Exercise 3: Advanced Full Pipeline

Time: 90 minutes

Difficulty: Challenging

Task:

- Scrape real reviews
- Preprocess text
- Fine-tune model
- Generate personas
- Create dashboard

#### Learning Goal:

End-to-end implementation

Deliverable:

Emotion insights report

**Resources:** Notebooks at [github.com/ml-design-course/week3-nlp](https://github.com/ml-design-course/week3-nlp)

Progressive complexity builds competence - sequential skill development enables advanced capability through mastered fundamentals

### From Words to Design Insights

#### Technical Understanding:

- Why context matters in language
- How Bag of Words loses information
- What attention mechanisms do
- How BERT reads bidirectionally
- Why pre-training + fine-tuning works

#### Practical Skills:

- Identify emotion in text at scale
- Use pre-trained models effectively
- Fine-tune for specific domains
- Interpret attention visualizations
- Convert ML outputs to insights

#### Design Applications:

- Data-driven persona creation
- Emotion-based user journeys
- Priority matrices from ML analysis
- Scalable empathy systems
- Personalization strategies

#### Remember:

**The Goal:** Not to read faster, but to understand deeper  
**The Method:** Selective attention to what matters  
**The Result:** True user understanding at scale

### Next Week: Classification & Problem Definition

Technical capability transforms analytical scope - methods enable understanding previously inaccessible through manual approaches

## Key Takeaway

Understanding emotion at scale is not about reading faster—it's about attending smarter

**Next Week:** Classification & Problem Definition  
From emotions to actionable categories