

## A/B Testing & Iterative Improvement

Week 10: Closing the ML Innovation Loop

Machine Learning for Smarter Innovation

BSc-Level Course

- 1 Foundation: The Iteration Advantage
- 2 Techniques: Rigorous Experiment Design
- 3 Implementation: Building A/B Test Infrastructure
- 4 Design: Building Experimentation Culture
- 5 Practice: Recommendation Engine A/B Test

# The Iteration Advantage: Why Speed Wins

## The Winners Iterate Fast

### Spotify:

- Tens of thousands of experiments annually
- Test every feature before launch
- 30% faster iteration than competitors
- Result: 500M+ users

### Netflix:

- AI recommendation system saves \$1B annually
- A/B testing validates all improvements
- 80% engagement from personalization
- Tests 250+ variations simultaneously

### Amazon:

- Tests every UI change
- 35% revenue from recommendations (McKinsey)
- "Our success is a function of experiments"
- Iterates daily, not quarterly

Experimental velocity correlates with innovation output - systematic testing infrastructure accelerates discovery rates

## The Losers Deploy and Hope

### Traditional Approach:

- Build for 6 months
- Launch to 100% users
- Hope it works
- No measurement
- No iteration

## The Iteration Gap

### Industry Reality:

- 87-90% of ML projects never reach production
- Most that deploy: No iteration, no measurement
- Majority lack A/B testing infrastructure
- Average: 2-3 experiments per year vs 50+

## Week 10 Mission

Transform you from "deploy and hope" to "test and learn" practitioners who ship 50+ experiments annually.

# Why Deployed Models Decay Without Iteration

## The Decay Curve

### Typical Model Degradation:

- Month 0: 90% accuracy (deployment)
- Month 3: 87% accuracy (3% drop)
- Month 6: 82% accuracy (9% drop)
- Month 12: 75% accuracy (17% drop)
- Month 18: 68% accuracy (24% drop)

Performance drops 15-30% per year without updates

### Why This Happens:

- User behavior changes
- Competitor adaptations
- Market dynamics shift
- Seasonal patterns evolve
- New product features launched

## Root Causes

### 1. Concept Drift

- User preferences change
- $P(Y|X)$  relationship shifts
- Training data becomes stale

### 2. Data Distribution Shift

- New user demographics
- $P(X)$  changes over time
- Features no longer predictive

### 3. Competitive Response

- Rivals launch better features
- Users have higher expectations
- Bar constantly rising

## The Solution

### Continuous iteration through A/B testing:

- Weekly model updates
- Feature experiments
- Algorithm improvements

# From “Deploy and Hope” to “Test and Learn”

## Old Mindset: Waterfall

### Process:

- ① Spend 6 months building
- ② Launch to all users
- ③ Cross fingers
- ④ Wait for complaints
- ⑤ Emergency fix if broken
- ⑥ Repeat in 12 months

### Problems:

- High risk (all eggs in one basket)
- Slow learning (feedback after 6 months)
- No data (guessing what works)
- Expensive failures
- Competitor advantage grows

Success rate: 10-20%

## New Mindset: Scientific Method

### Process:

- ① Form hypothesis ("X will improve Y")
- ② Design experiment (A/B test)
- ③ Calculate sample size (power analysis)
- ④ Deploy to 5% users (low risk)
- ⑤ Measure results (data-driven)
- ⑥ Ship winners, kill losers
- ⑦ Repeat weekly

### Benefits:

- Low risk (gradual rollout)
- Fast learning (results in days)
- Data-driven (know what works)
- Compound gains (many small wins)
- Competitive edge (10 times iteration speed)

Success rate: 30-40%

**Key insight:** Most experiments “fail” (don’t beat control), but you learn fast and iterate



# A/B Testing Fundamentals: The Gold Standard

## What is A/B Testing?

**Definition:** Randomized controlled trial comparing two versions to determine which performs better.

### Core Components:

- ① **Control Group (A):** Current system (baseline)
- ② **Treatment Group (B):** New variant (challenger)
- ③ **Random Assignment:** Users split 50/50
- ④ **Metric:** What you're measuring (CTR, revenue)
- ⑤ **Statistical Test:** Determine if difference is real

### Example: Recommendation Algorithm

- Control: Popularity-based (5% CTR)
- Treatment: ML collaborative filtering
- Metric: Click-through rate
- Users: 50,000 per group
- Duration: 2 weeks
- Result: Treatment wins (6.2% CTR)

## Why Randomization Matters

### Without randomization:

- Selection bias (power users in treatment)
- Confounding variables (time of day, device)
- Can't isolate causal effect
- Results are unreliable

### With randomization:

- Groups are statistically equivalent
- Only difference is the treatment
- Causal interpretation valid
- Results are trustworthy

## Key Requirements

- **Sample size:** Large enough for statistical power
- **Duration:** Long enough to capture true behavior
- **Randomization:** Truly random assignment
- **Isolation:** No spillover between groups
- **Pre-registration:** Hypothesis before experiment

# Real Success Stories: Data-Driven Wins

## Booking.com

Experiment: Scarcity messaging

### Hypothesis:

"Only 2 rooms left!" increases urgency and conversions

### Test:

- Control: No scarcity message
- Treatment: Real-time room count
- Metric: Conversion rate
- Users: 100K per group

### Result:

- 25% conversion lift
- \$50M annual revenue impact
- Rolled out globally

**Lesson:** Small UX changes, massive impact when tested rigorously

Empirical evidence supersedes expert intuition - measured outcomes reveal truths hidden from qualitative assessment

## Obama 2012 Campaign

Experiment: Email subject lines

### Challenge:

Raise \$500M through email donations

### Test:

- 20 subject line variants
- A/B/C/D.../T testing
- Metric: Donation rate
- Recipients: 13M per variant

### Result:

- Winner: "Hey" (simplest)
- 70% higher open rate
- \$500K extra per email
- 500+ experiments total

**Lesson:** Intuition fails. Data wins. Test everything.

## Google Search Ads

Experiment: 41 shades of blue

### Question:

Which blue color for ad links maximizes clicks?

### Test:

- 41 blue color variants
- Metric: Click-through rate
- Traffic: Billions of searches
- Duration: 3 weeks

### Result:

- Optimal blue found
- 1% CTR improvement
- \$200M annual revenue
- Criticized but data-driven

**Lesson:** At scale, tiny improvements = huge value. Test obsessively.

# Real Failure Examples: The Cost of Not Testing

## Knight Capital: \$440M Loss

### What Happened:

- August 1, 2012: Deploy new trading algorithm
- No A/B test, no gradual rollout
- Bug: Bought high, sold low repeatedly
- Duration: 45 minutes
- Loss: \$440 million
- Outcome: Company bankrupt

### What Should Have Been Done:

- Shadow mode testing (parallel run)
- Canary deployment (1% traffic first)
- Kill switch (automatic rollback)
- Guardrail metrics (loss limits)

**Lesson:** High-stakes deployments demand rigorous testing

## Microsoft Tay Chatbot

### What Happened:

- March 2016: Launch AI chatbot on Twitter
- No gradual rollout, no guardrails
- Users exploited vulnerabilities
- Within 16 hours: Racist, offensive tweets
- Outcome: Shut down, PR disaster

### What Should Have Been Done:

- Beta test with controlled users
- Content moderation safeguards
- Incremental rollout ( $100 \rightarrow 1K \rightarrow 10K$  users)
- Real-time monitoring and kill switch

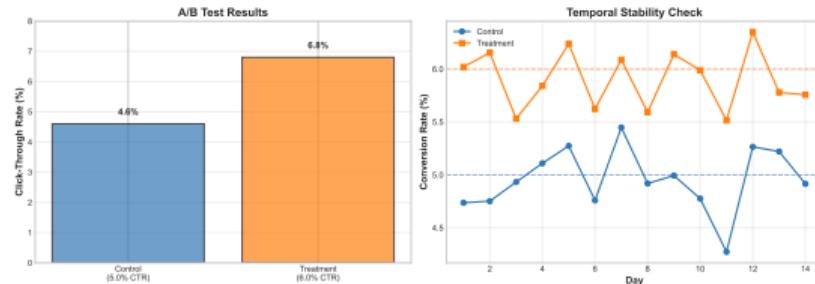
**Lesson:** AI in the wild requires safety testing

## Common Thread

Both failures share: (1) No gradual rollout, (2) No guardrails, (3) No monitoring, (4) No rollback plan

Experimentation prevents catastrophic failures - systematic testing identifies risks before they manifest at scale

# The Experimentation Pyramid: Three Levels



## Level 1: Speed

Ship fast, iterate faster

- 50+ experiments per year
- 1-2 week experiment duration
- Automated analysis pipelines
- Low friction deployment

## Level 2: Safety

Guardrails prevent disasters

- Revenue floor (don't lose money)
- Latency ceiling (stay fast)
- Error rate limits (don't break)
- Automatic kill switches

## Level 3: Learning

Extract insights, compound gains

- Why did it win/lose?
- What worked for which users?
- How to generalize learnings?
- Build institutional knowledge

# When to A/B Test ML Models

## Always A/B Test

### 1. New Model Deployment

- Replacing existing model
- High risk of regression
- Need causal validation

### 2. Algorithm Changes

- Random Forest

→ XGBoost

### • Content-based

→ collaborative filtering

### • Offline metrics don't predict online

### 3. Feature Updates

- Add new features
- Remove stale features
- Feature engineering changes

### 4. Hyperparameter Tuning

- Learning rate changes
- Regularization adjustments

## Sometimes A/B Test

### 6. UI/UX Changes

- Model recommendations display
- Explanation text
- Confidence thresholds shown to users

### 7. Business Logic

- Ranking vs filtering
- Diversity vs relevance
- Personalization intensity

## Don't Bother A/B Testing

### 8. Bug Fixes

- Obvious correctness issues
- Fix immediately, no test needed

### 9. Infrastructure Changes

- Same model, faster serving
- No user-facing behavior change

## Rule of Thumb

If change affects user experience or business metrics → A/B test it

# Week 10 Learning Objectives

## Technical Skills

By end of week, you will:

- ① **Design** statistically rigorous A/B tests for ML models
- ② **Calculate** sample sizes and experiment duration (power analysis)
- ③ **Implement** experiments in Python (scipy, PyMC3, bandits)
- ④ **Interpret** p-values, confidence intervals, effect sizes correctly
- ⑤ **Apply** Bayesian A/B testing for faster decisions
- ⑥ **Use** multi-armed bandits for real-time optimization
- ⑦ **Identify** and avoid common pitfalls (peeking, multiple testing)
- ⑧ **Build** experiment monitoring dashboards

## Strategic Skills

You will master:

- Formulating testable hypotheses
- Choosing primary and guardrail metrics
- Balancing speed and statistical rigor
- Communicating results to stakeholders
- Making confident deployment decisions
- Building experimentation culture
- Measuring long-term impact
- Iterating systematically

## Professional Outcomes

- Design and run 50+ experiments annually
- Avoid catastrophic deployment failures
- Iterate 10*times* faster than competitors
- Build data-driven product culture
- Become trusted voice in product decisions

Validation establishes baseline quality while experimentation drives continuous improvement - measurement enables optimization

# Foundation Summary: Why Iteration Wins

## Key Concepts

### 1. Iteration Advantage

- Winners (Spotify, Netflix) ship 1000+ experiments/year
- Losers deploy once and hope
- Speed of learning = competitive edge

### 2. Model Decay

- Performance drops 15-30% annually without updates
- Concept drift, distribution shift, competition
- Solution: Continuous iteration

### 3. Experimentation Mindset

- From “deploy and hope” to “test and learn”
- Scientific method: Hypothesis

→ Test

→ Learn

- Most experiments “fail” but learning compounds

### 4. A/B Testing Fundamentals

- Randomized controlled trials
- Control vs treatment

## Real-World Lessons

### Success Stories:

- Booking.com: 25% conversion lift from scarcity
- Obama 2012: \$500K per optimized email
- Google: \$200M from testing blue shades

### Failure Examples:

- Knight Capital: \$440M loss (no testing)
- Microsoft Tay: Shut down in 16 hours
- Lesson: Gradual rollout + guardrails essential

## The Experimentation Pyramid

- **Speed:** Ship fast, iterate faster
- **Safety:** Guardrails prevent disasters
- **Learning:** Extract insights, compound gains

Next: Learn statistical techniques for rigorous A/B testing

# Hypothesis Formulation: The Starting Point

## Anatomy of a Good Hypothesis

### SMART Framework:

- Specific: Clearly defined change
- Measurable: Quantifiable outcome
- Actionable: You can implement it
- Relevant: Aligns with business goals
- Time-bound: Experiment duration set

### Good Example

**Hypothesis:** Switching from content-based to collaborative filtering recommendations will increase click-through rate by at least 1 percentage point (from 5% to 6%) over a 2-week test period with 100,000 users.

### Why good:

- Specific change (CF algorithm)
- Measurable outcome (CTR, 1 pp)
- Actionable (can deploy CF)
- Relevant (engagement goal)
- Time-bound (2 weeks, 100K users)

Hypothesis pre-registration prevents post-hoc rationalization - documenting predictions before observing results ensures scientific integrity

## Bad Examples

**Vague:** "New algorithm will be better"

- Not specific (which algorithm?)
- Not measurable (better how?)
- Not time-bound

**Unmeasurable:** "Users will like recommendations more"

- "Like" not quantified
- No success criterion

**Too ambitious:** "Will 10imes revenue"

- Unrealistic expectation
- Sets up for disappointment

## Null vs Alternative Hypothesis

**Null ( $H_0$ ):** No difference between control and treatment

$$CTR_{treatment} = CTR_{control}$$

**Alternative ( $H_1$ ):** Treatment is better

$$CTR_{treatment} > CTR_{control} \text{ (one-tailed)}$$

$$CTR_{treatment} \neq CTR_{control} \text{ (two-tailed)}$$

# Sample Size Calculation: How Many Users?



## Key Parameters

### 1. Significance Level ( $\alpha$ )

- Type I error rate (false positive)
- Standard: 0.05 (5%)
- Probability of detecting difference when none exists

### 2. Statistical Power ( $1 - \beta$ )

- Type II error rate (false negative)
- Standard: 0.80 (80%)
- Probability of detecting true difference

### 3. Minimum Detectable Effect (MDE)

- Smallest difference you care about
- Example: 5%  
→ 6% CTR (1 pp absolute, 20% relative)
- Smaller MDE needs larger sample

### 4. Baseline Metric

- Current performance level
- Variance affects sample size

Sample size determines statistical power - insufficient observations lead to inconclusive results regardless of analytical rigor

# Randomization Strategies: More Than Just Coin Flips

## Simple Randomization

**Method:** Coin flip per user (50/50)

### Pros:

- Easy to implement
- Unbiased in expectation
- Good for large samples

### Cons:

- Group sizes may differ
- Imbalanced covariates possible

## Stratified Randomization

**Method:** Balance within strata (device, region)

### Example:

- Split iOS users 50/50
- Split Android users 50/50
- Ensures device balance

### Benefits:

- Reduces variance
- Controls confounders

## Blocked Randomization

**Method:** Randomize within fixed blocks

### Example:

- Block 1: AABBA (2A, 2B)
- Block 2: BAABB (2A, 2B)
- Guarantees equal sizes

**Use case:** Sequential enrollment

## Clustered Randomization

**Method:** Randomize groups, not individuals

### Example:

- User sessions (not individual actions)
- Geographic regions (not cities)
- Social networks (friends together)

### Why:

- Prevent contamination
- Network effects
- Administrative ease

# A/B Testing Math: Statistical Tests

## Z-Test for Proportions

When: Binary outcome (click/no click)

Formula:

$$z = \frac{p_B - p_A}{\sqrt{p(1-p)\left(\frac{1}{n_A} + \frac{1}{n_B}\right)}}$$

where  $p = \frac{n_A p_A + n_B p_B}{n_A + n_B}$  (pooled proportion)

Example:

- Control: 500/10,000 clicked (5%)
- Treatment: 620/10,000 clicked (6.2%)
- $z = 4.72$ ,  $p < 0.001$
- Conclusion: Significant improvement

## T-Test for Means

When: Continuous outcome (revenue, time)

Formula:

$$t = \frac{\bar{x}_B - \bar{x}_A}{\sqrt{\frac{s_A^2}{n_A} + \frac{s_B^2}{n_B}}}$$

Use Welch's t-test: Unequal variances OK

## Confidence Intervals

95% CI for difference:

$$(\bar{x}_B - \bar{x}_A) \pm 1.96 \cdot SE$$

Interpretation:

- If CI excludes 0  
→ Significant
- CI shows practical range
- More informative than p-value alone

Example:

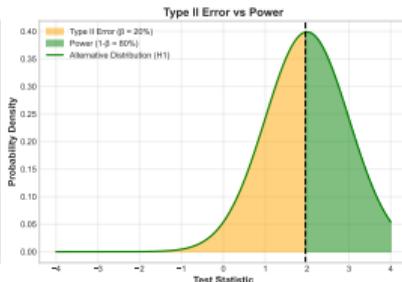

- Difference: 1.2 percentage points
- 95% CI: [0.8, 1.6]
- Excludes 0  
→ Significant
- Likely between 0.8-1.6pp lift

## Effect Size

Cohen's d:

$$d = \frac{\bar{x}_B - \bar{x}_A}{s_{pooled}}$$

# Statistical Significance: Interpreting P-Values



## What is a P-Value?

**Definition:** Probability of observing data this extreme if null hypothesis were true.

### NOT:

- Probability null is true
- Probability you're wrong
- Size of effect

### Example:

- $p = 0.03$
- If no real difference exists, 3% chance of seeing this result by luck
- Since  $3\% < 5\%$ , reject null

## Type I vs Type II Errors



### Type I (False Positive):

- Declare winner when none exists
- Rate:  $\alpha$  (typically 5%)
- "Ship a loser"

### Type II (False Negative):

- Miss real improvement

# Bayesian A/B Testing: Probability of Being Best



## Bayesian Approach

Instead of p-values:

- Calculate  $P(B > A)$
- Direct probability interpretation
- Incorporate prior knowledge
- No arbitrary 0.05 threshold

Example:


- $P(B > A) = 0.97$
- 97% sure treatment is better
- Can ship with 95% confidence

## Advantages

- Intuitive interpretation
- Earlier stopping (faster decisions)
- Handles small samples better
- Can update continuously
- No "peeking problem"

## When to Use

# Understanding Confidence Intervals Through Repetition



## What 95% CI Means

**Definition:** If we repeated the experiment 100 times, 95 of the intervals would contain the true value.

### Key Insights:

- Each experiment produces a different interval
- True value is fixed (but unknown)
- Some intervals miss (by design!)
- Wider intervals = more uncertainty

### Common Misconceptions

**Wrong:** "95% chance true value is in this interval"

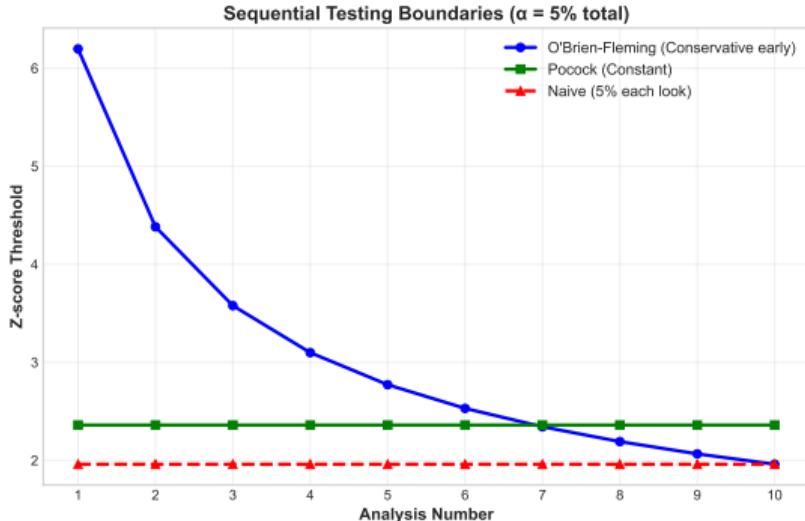
- True value either is or isn't in it
- It's about the procedure, not this interval

**Right:** "95% of intervals constructed this way contain the true value"

### Practical Use

- If CI includes 0: Not significant
- If CI excludes 0: Significant
- Width shows precision of estimate

# Sequential Testing: Early Stopping Without Peeking



## The Peeking Problem

### Bad practice:

- Check results daily
- Stop when  $p < 0.05$
- False positive rate > 50%!

### Why bad:

- Multiple testing inflates  $\alpha$
- Random fluctuations look significant
- Results not reproducible

## Sequential Testing Solution


### Method: Alpha spending functions

- Plan interim analyses (e.g., 25%, 50%, 75%, 100%)
- Adjust  $\alpha$  at each look
- Control overall Type I error

### O'Brien-Fleming:

- Conservative early, liberal late
- First look:  $\alpha = 0.0005$
- Last look:  $\alpha = 0.0455$

# Multi-Armed Bandits: Exploration vs Exploitation



## The Bandit Problem

**Scenario:** Slot machines (arms) with unknown payouts.

**Goal:** Maximize total reward.

### Trade-off:

- **Explore:** Try arms to learn
- **Exploit:** Use best-known arm

## Thompson Sampling

### Algorithm:

- ① Maintain Beta( $\alpha, \beta$ ) for each arm
- ② Sample from each distribution
- ③ Pick arm with highest sample
- ④ Update winner's distribution

### Why it works:


- Automatically balances exploration/exploitation
- Converges to best arm
- Minimizes regret

## A/B Test vs Bandit

### A/B Test:



# Bandit Performance: Exploration vs Regret



## How Bandits Learn

### Left Panel: Cumulative Pulls

- Initially: Explore all arms
- Middle: Identify good arms
- End: Exploit best arm
- Bad arms stop getting traffic

### Right Panel: Regret

Regret: Total reward lost by not always picking best arm

### Key Insight:

- Regret grows logarithmically
- Much slower than linear
- Thompson Sampling is near-optimal
- A/B test has linear regret


## When to Use Bandits

- High traffic (fast learning)
- Multiple variants ( $> 3$ )
- Cost of regret is high
- Can tolerate adaptive allocation



# Causal Inference: Correlation is Not Causation

Directed Acyclic Graph: Randomization Breaks Confounding



## The Fundamental Problem

**Correlation:** X and Y move together

**Causation:** X causes Y

### Why different:

- Confounders (Z affects both X and Y)
- Reverse causation (Y causes X)
- Spurious correlation (coincidence)

### Example: Ice Cream & Drowning

- Correlation: High
- Causation: None
- Confounder: Hot weather
- Weather causes both ice cream sales and swimming (drowning risk)

### Randomization Solves This

#### How:

- Random assignment breaks confounding
- Only difference is treatment
- Can infer causation

# Multi-Objective Optimization: Trade-Offs

## The Real-World Problem

Single metric is rare:

- CTR vs revenue
- Short-term vs long-term
- Engagement vs satisfaction
- Speed vs accuracy

Example: Recommendation algorithm

- Control: Higher revenue, lower engagement
- Treatment: Lower revenue, higher engagement
- Which is better?

## Solution 1: Primary + Guardrails

Method:

- Primary: Metric you optimize (revenue)
- Guardrails: Metrics that must not degrade
- Example: Maximize revenue, but engagement must not drop > 2%

Decision rule:

- Ship if primary improves AND guardrails met

## Solution 2: Weighted Utility

Method: Combine metrics into single score

$$U = w_1 \cdot \text{revenue} + w_2 \cdot \text{engagement}$$

Example weights:

- Revenue: 0.7 (70% weight)
- Engagement: 0.3 (30% weight)

Challenge: Choosing weights

## Solution 3: Pareto Frontier

Method: Find non-dominated solutions

- Solution A dominates B if better on all metrics
- Pareto frontier: Set of non-dominated solutions
- Business decides among frontier

When to use:

- Cannot agree on weights
- Want to see trade-off space
- Explore multiple strategies

# Technique Comparison: When to Use Each

## Classical A/B Test

### When:

- High traffic
- Need causal rigor
- Low risk
- Regulatory requirements
- Can wait for full sample

### Pros:

- Rigorous
- Well understood
- Easy to interpret
- Reproducible

### Cons:

- Slow (fixed duration)
- Higher regret
- Can't peek

### Example:

Major algorithm change, need 95%

## Bayesian A/B Test

### When:

- Need faster decisions
- Low traffic
- Can update continuously
- Prior knowledge exists
- Stakeholders prefer probabilities

### Pros:

- Intuitive ( $P(B > A)$ )
- Earlier stopping
- Handles small samples
- No peeking problem

### Cons:

- Prior selection subjective
- Less familiar to stakeholders
- Requires Bayesian tooling

### Example:

Startup with 1K daily users, need quick wins

## Multi-Armed Bandit

### When:

- Minimize regret
- Continuous optimization
- Many variants (A/B/C/D/E)
- Cost of losing high
- Real-time personalization

### Pros:

- Lowest regret
- Adaptive allocation
- Scales to many arms
- Always optimizing


### Cons:

- Less statistical rigor
- Harder to interpret
- Complex implementation

### Example:

Ad creative testing (50+ variants), email list splitting

# Production Experiment Infrastructure



## Core Components

### 1. Traffic Splitting

- Assign users to control/treatment
- Consistent hashing (same user, same group)
- 50/50, 90/10, or custom splits

### 2. Feature Flags

- Toggle experiments on/off
- Gradual rollouts (1% → 5% → 25% → 100%)
- Instant kill switch

### 3. Experiment Tracking

- Log user assignment
- Track metrics per group
- Store for analysis

### 4. Analysis Pipeline

- Automated statistical tests
- Real-time dashboards
- Alert on guardrail violations

Experiment infrastructure requires orchestration components - traffic splitting, feature flags, tracking, and analysis form the core system.

## Classical A/B Test Implementation:

- Import required libraries: `scipy.stats, statsmodels`
- Define sample data: control and treatment groups
- Calculate proportions for each group
- Perform Z-test for proportions using `proportions_ztest`
- Compare p-value to significance threshold (0.05)
- Calculate and display key metrics:
  - Control rate, treatment rate
  - Lift percentage
  - Z-statistic and p-value
  - Significance decision

## Output Example

### Sample Output:

- Control: 0.050 (5.0% conversion)
- Treatment: 0.062 (6.2% conversion)
- Lift: 24.0% improvement
- Z-statistic: 4.72
- P-value: 0.0001
- Result: Significant improvement!

## Confidence Interval

### Confidence Interval Calculation:

- Import `confint_proportions_2indep` from `statsmodels`
- Calculate 95% confidence interval using Wald method
- Parameters: treatment successes/total, control successes/total
- Output: lower and upper bounds of difference

## Interpretation

If CI excludes 0, treatment significantly better than control

Statistical libraries provide essential testing infrastructure - robust implementations enable reliable hypothesis evaluation



# Python: Bayesian A/B Testing with PyMC3

## Bayesian A/B Test with PyMC3:

- Import PyMC3 and numpy libraries
- Define observed data: clicks and trials for groups A and B
- Set up Bayesian model:
  - Beta priors for conversion rates (non-informative)
  - Binomial likelihoods for observed data
  - Deterministic variable for difference (delta)
- Sample from posterior distribution (2000 samples)
- Calculate probability that B outperforms A
- Direct probability statements without p-value confusion

## Output

### Bayesian Results:

- $P(B > A) = 0.997$  (99.7% confidence)
- Posterior mean delta: 0.012
- 95% Credible Interval: [0.008, 0.016]

## Interpretation

- 99.7% probability B is better
- Expected lift: 1.2 pp
- 95% sure lift between 0.8-1.6pp
- High confidence to ship

## Advantages

- Direct probability statements
- No p-value confusion
- Can stop early if  $P(B > A) > 0.95$
- Incorporates prior knowledge

Probabilistic inference enables faster experimental decisions - direct probability statements reduce interpretation overhead

# Python: Thompson Sampling Bandit

## Thompson Sampling Bandit Implementation:

- Class initialization with number of arms
- Maintain success/failure counts per arm (Beta parameters)
- Selection algorithm:
  - Sample from Beta distribution for each arm
  - Select arm with highest sample value
- Update mechanism:
  - Increment success count for reward = 1
  - Increment failure count for reward = 0
- Main loop: select arm, observe reward, update beliefs
- Converges to optimal arm while balancing exploration

## How It Works

1. Initialize:
  - Each arm:  $\text{Beta}(1, 1)$  = uniform prior
2. Select arm:
  - Sample from each Beta distribution
  - Pick arm with highest sample
3. Observe reward:
  - If success:  $\alpha + 1$
  - If failure:  $\beta + 1$
4. Repeat

## Convergence

- Best arm gets pulled more often
- Exploration decreases over time
- Minimizes cumulative regret

**Use case:** Ad creative testing (50 variants, continuous optimization)

Probabilistic selection algorithms balance simplicity and effectiveness - sampling from posterior distributions automates exploration-exploitation trade-offs



# Python: Stratified Randomization with sklearn

## Stratified Randomization Implementation:

- Import StratifiedShuffleSplit from sklearn
- Create user dataframe with stratification variables
- Set up stratified split:
  - 50/50 split between control and treatment
  - Stratify by device type (iOS/Android)
  - Use fixed random state for reproducibility
- Apply split to create balanced groups
- Verify stratification worked:
  - Check device distribution in both groups
  - Should be identical proportions
- Can extend to multiple stratification variables

## Output

### Verification Results:

- Control device distribution:
  - iOS: 0.50 (50%)
  - Android: 0.50 (50%)
- Treatment device distribution:
  - iOS: 0.50 (50%)
  - Android: 0.50 (50%)

## Why This Matters

### Without stratification:

- Control might be 55% iOS
- Treatment might be 45% iOS
- Device becomes confounder
- Results biased

### With stratification:

- Both groups exactly 50% iOS
- Device balanced
- Removes confounding
- Higher statistical power

# Python: Sequential Testing with Early Stopping

## Sequential Testing with O'Brien-Fleming:

- Import numpy and scipy.stats.norm
- Define O'Brien-Fleming boundary function:
  - Input: current analysis number ( $k$ ), total analyses ( $K$ )
  - Calculate alpha spending based on information fraction
  - Conservative early, liberal late
- Plan interim analyses (e.g., 25%, 50%, 75%, 100%)
- For each analysis:
  - Calculate current p-value
  - Compare to adjusted alpha threshold
  - Stop if significant, continue otherwise
- Maintains overall Type I error rate at 0.05
- Enables early stopping without “peeking” penalty

## Alpha Spending

### Alpha Spending Schedule:

- Look 1 (25%):  $\alpha = 0.0005$
- Look 2 (50%):  $\alpha = 0.0139$
- Look 3 (75%):  $\alpha = 0.0303$
- Look 4 (100%):  $\alpha = 0.0455$

## Interpretation

- Very conservative early (0.0005)
- More liberal late (0.0455)
- Total  $\alpha$  still 0.05
- Can stop early if clear winner

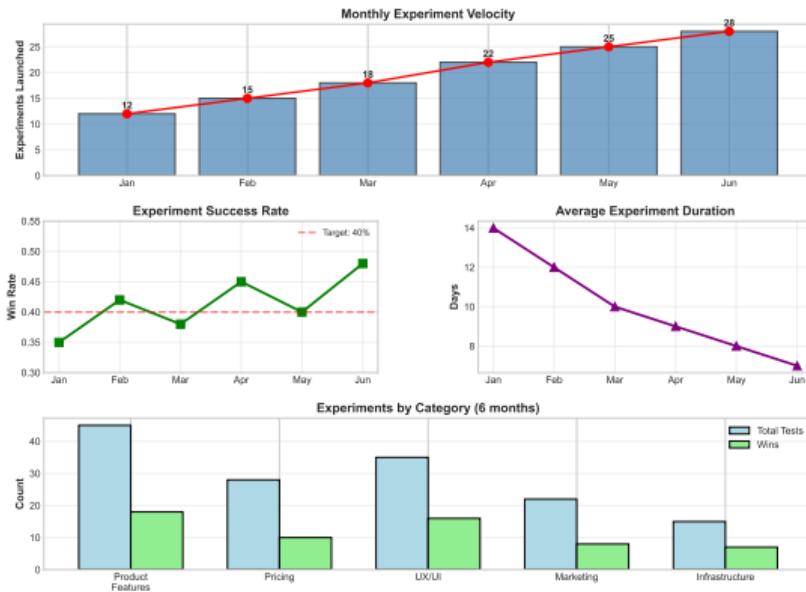
## Benefits

- Valid p-values at each look
- Can stop for futility (clear loser)
- Faster decisions on average
- Maintains Type I error control

Cost: Slightly larger sample needed vs fixed design



# Real-Time Experiment Monitoring



## Key Metrics to Monitor

### Primary Metric:

- Current estimate
- Confidence interval
- P-value or  $P(B > A)$
- Sample size achieved

### Guardrail Metrics:

- Revenue (must not drop  $> 2\%$ )
- Latency ( $p95 < 200\text{ms}$ )
- Error rate ( $< 0.1\%$ )
- Automatic alerts if violated

### Diagnostic Metrics:

- Sample ratio mismatch
- Traffic allocation balance
- Novelty effects
- Time-of-day patterns

## Automated Actions

- Stop if guardrail violated

# Common A/B Testing Pitfalls

## Peeking Problem

**What:** Check results daily, stop when  $p < 0.05$

### Why bad:

- Inflates false positive rate
- Random fluctuations look significant
- Not reproducible

### Fix:

- Wait for planned sample
- Use sequential testing
- Or Bayesian methods

## Multiple Testing

**What:** Test 20 variants, report winner

### Why bad:

- 1 false positive expected
- Winner likely spurious

### Fix:

- Bonferroni correction

## Simpson's Paradox

**What:** Treatment wins overall, loses in every segment

### Example:

- iOS: Control 10%, Treatment 9%
- Android: Control 8%, Treatment 7%
- Overall: Treatment wins (imbalance in device mix)

### Fix:

- Stratified randomization
- Analyze by segment
- Check for confounders

## Novelty Effects

**What:** Users react to change itself, not treatment

### Signs:

- Week 1: Treatment wins
- Week 2: Effect shrinks
- Week 3: No difference

## Sample Ratio Mismatch

**What:** 50/50 split becomes 52/48

### Causes:

- Bot traffic
- Implementation bugs
- Sampling bias

### Why bad:

- Groups not comparable
- Results invalid

### Fix:

- Monitor ratio daily
- Investigate deviations
- Filter bots

## Ignoring Variance

**What:** Focus only on means, ignore spread

### Why bad:

- High variance = unreliable
- May hurt some users badly

# Production Deployment Patterns

## Canary Release

### Process:

- ① Deploy to 1% users
- ② Monitor for 24 hours
- ③ If stable, 5%
- ④ Then 25%, 50%, 100%

### When:

- High-risk changes
- New algorithms
- Large refactors

### Benefits:

- Early error detection
- Minimal blast radius
- Easy rollback

### Guardrails:

- Error rate < 0.5%
- Latency < baseline + 20%
- Revenue > baseline - 5%

## Blue-Green Deploy

### Process:

- ① Green: Current (100%)
- ② Blue: New (0%)
- ③ Gradually shift traffic
- ④ Blue becomes 100%

### When:

- Infrastructure changes
- Database migrations
- Zero-downtime deploys

### Benefits:

- Instant rollback (flip traffic)
- Both versions running
- Compare metrics live

### Challenges:

- 2 times infrastructure cost
- Data consistency
- State management

## Shadow Mode

### Process:

- ① Send traffic to both models
- ② Serve old model to users
- ③ Log new model predictions
- ④ Compare offline

### When:

- Validating new model
- Zero risk testing
- Performance benchmarking

### Benefits:

- No user impact
- Real traffic testing
- Can run indefinitely

### Use case:

- Test before A/B test
- Validate offline metrics
- Check for bugs

# Implementation Summary & Best Practices

## Implementation Checklist

### Before Experiment:

- Hypothesis pre-registered
- Sample size calculated
- Randomization strategy chosen
- Guardrails defined
- Monitoring dashboard ready
- Kill switch tested

### During Experiment:

- Monitor metrics daily
- Check sample ratio balance
- Watch for guardrail violations
- No peeking at p-values (unless sequential)
- Document any issues

### After Experiment:

- Run statistical tests
- Calculate confidence intervals
- Check subgroup analyses

## Python Tools Ecosystem

### Statistical Testing:

- scipy: t-tests, z-tests
- statsmodels: Advanced tests, power analysis
- PyMC3: Bayesian methods

### Experimentation Platforms:

- GrowthBook: Open-source, feature flags
- Optimizely: Enterprise A/B testing
- LaunchDarkly: Feature management
- Firebase: Mobile experiments

### Visualization:

- matplotlib/seaborn: Static plots
- Plotly: Interactive dashboards
- Streamlit: Quick apps

## Key Takeaways

- Automate as much as possible
- Monitor continuously
- Plan for early stopping

# Building Experimentation Culture: Leadership Principles

## Jeff Bezos on Experimentation

"If you double the number of experiments you do per year, you're going to double your inventiveness."

## Amazon's Approach:

- Experiments are default, not exception
- Ship fast, iterate faster
- Celebrate "failed" experiments (learning)
- Data beats opinions
- Everyone empowered to test

## Google's Philosophy

"We're wrong a lot. Most experiments fail. That's OK—we learn fast and move on."

## Key principles:

- 10-20% win rate is normal
- Compound small wins ( $1\% \text{ lift} \times 100 \text{ experiments}$ )
- Avoid HiPPO (Highest Paid Person's Opinion)
- Test everything, assume nothing

## Characteristics of Strong Culture

### 1. Velocity:

- 50-100+ experiments per year
- 1-2 week experiment cycle
- Automated pipelines

### 2. Psychological Safety:

- OK to propose "crazy" ideas
- No punishment for failed experiments
- Learn from mistakes openly


### 3. Data Literacy:

- Understand p-values, CI, effect sizes
- Know when to trust results
- Question unexpected findings

### 4. Infrastructure:

- Feature flags in production
- Real-time dashboards
- 1-click deploy and rollback

# Metric Selection: North Star & Guardrails



## North Star Metric

**Definition:** Single metric that captures long-term value creation

### Examples:

- Spotify: Weekly active users
- Netflix: Watch time per user
- Airbnb: Nights booked
- Amazon: Purchases per customer

### Criteria:

- Aligns with business model
- Measurable in A/B tests
- Leading indicator (not lagging)
- Actionable by teams

## Guardrail Metrics

**Purpose:** Prevent optimization tunnel vision

### Examples:

- Revenue: Must not drop  $> 2\%$
- Latency:  $p95 < 300ms$

# Communicating Results to Stakeholders

## For Executives

### What they care about:

- Bottom line impact
- Risk level
- Confidence level
- Timeline

### Example:

"Treatment increased revenue by 8% (95% CI: 5%-11%). Expected annual impact: \$2.4M. Recommend full rollout over 2 weeks. Risk: Low (guardrails passed)."

## For Product Managers

### What they care about:

- User experience impact
- Segment differences
- Feature interactions
- Next iterations

### Example:

"Treatment improved CTR 15% for new users, 5% for power users. iOS stronger than Android. Game-changing innovation."

## For Engineers

### What they care about:

- Implementation details
- Performance metrics
- Technical challenges
- Edge cases

### Example:

"Model latency p95: 180ms (vs 150ms baseline). Sample ratio 50.2/49.8 (acceptable). Bug in iOS < 13 handling fixed day 3. Rerun analysis confirmed results."

## For Data Scientists

### What they care about:

- Statistical rigor
- Effect sizes
- Confidence intervals
- Methodology

### Example:

" $Z = 4.2$ ,  $p < 0.001$ , Cohen's  $d = 0.28$  ( $\text{Cohen's } d = \frac{\text{mean difference}}{\text{standard deviation}}$ )".

## Universal Principles

### Always include:

- ① Hypothesis tested
- ② Sample size achieved
- ③ Primary metric result + CI
- ④ Guardrail metric status
- ⑤ Recommendation (ship/iterate/kill)
- ⑥ Rationale for recommendation

## Visualization Tips

- Show distributions, not just means
- Include confidence intervals
- Highlight guardrail bounds
- Use color coding (green/yellow/red)
- Keep it simple (no jargon)

## Red Flags to Call Out

- Sample ratio mismatch
- Novelty effects
- Simpson's paradox

# Explaining Statistics to Non-Technical Stakeholders

## P-Value Translation

### Instead of:

"P-value is 0.03, so we reject the null hypothesis at  $\alpha = 0.05$ ."

### Say:

"If there were truly no difference, we'd see a result this extreme only 3% of the time by random chance. Since that's unlikely, we're confident treatment is better."

## Confidence Interval Translation

### Instead of:

"95% CI: 1.2%-3.8%"

### Say:

"We're 95% confident the true improvement is between 1.2% and 3.8%. Best estimate: 2.5%."

## Statistical Power Translation

### Instead of:

"80% power to detect 0.5pp effect"

### Say:

"If there's a real 0.5 percentage point improvement, we'll detect it 80% of the time. We might miss smaller improvements."

## Common Misunderstandings

**Wrong:** "P = 0.03 means 97% chance treatment is better"

- P-value is NOT probability hypothesis is true
- Use Bayesian methods for that

**Wrong:** "Significant = Important"

- Statistical  $\neq$  Practical significance
- 0.01% lift might be significant but meaningless

**Wrong:** "Not significant = No effect"

- Absence of evidence  $\neq$  Evidence of absence
- Might just need larger sample

## Analogies That Work

### Confidence Interval:



"Like a weather forecast: 70-80°F tomorrow means we're sure it won't be 50° or 100°, but exact temp uncertain."

### Type I/II Errors:

"Fire alarm: False alarm (Type I) = evacuate unnecessarily. Missed fire (Type II) = danger goes undetected."

Statistical literacy compounds across iterations - understanding accumulates faster than individual technique mastery

# Simpson's Paradox: When Aggregation Misleads



## The Paradox

### What happens:

- Treatment wins in segment A
- Treatment wins in segment B
- Control wins overall!

**Why:** Imbalanced sample sizes across segments distort the aggregate result.

## Real Example

### UC Berkeley Admission (1973):

- Overall: Men 44%, Women 35% (bias?)
- But in most departments: Women had higher admit rate
- Explanation: Women applied to harder departments

## How to Detect

- Always analyze by key segments
- Check if segment sizes balanced
- Use stratified randomization
- Report both aggregate and segment results

# Decision Frameworks: Ship, Iterate, or Kill?



Always consider: Statistical significance + Practical significance + Guardrails

## Ship Criteria

All must be true:

- Statistically significant ( $p < 0.05$  or  $P(B > A) > 0.95$ )
- Practically significant (effect size meaningful)
- All guardrails passed
- Benefit > cost (ROI positive)
- Implementation ready

## Example:

15% CTR lift,  $p = 0.001$ , \$500K annual revenue, latency OK, error rate OK → SHIP

## Iterate Criteria

When to iterate:

- Directionally positive but not significant
- Significant but guardrail violated
- Strong in 1 segment, weak in others
- Good idea, poor execution

## Example:

3% CTR lift,  $p = 0.12$  (not sig), but iOS showed 8% lift ( $p = 0.01$ ) → ITERATE for iOS only

# Measuring Long-Term Impact: Beyond the Experiment

## The Long-Term Challenge

### Problem:

- Experiments run 1-2 weeks
- But true impact unfolds over months
- Short-term vs long-term trade-offs

### Example:

- Treatment boosts Week 1 engagement +20%
- Week 4: Back to baseline (novelty wore off)
- Month 6: -5% (user fatigue)

## Solution 1: Holdout Groups

### Method:

- Keep 10% of users in control forever
- Compare treatment vs holdout over 6-12 months
- Measure long-term metrics

### Benefits:

- Detect delayed effects
- Catch cumulative degradation
- Measure true business impact

## Solution 2: Cohort Analysis

### Method:

- Track users by week of experiment entry
- Compare Week 1, 2, 3 cohorts over time
- Look for retention/engagement patterns

### Example:

Week 1 cohort: 30-day retention 60%  
Week 4 cohort: 30-day retention 55% → Novelty effect

## Solution 3: Synthetic Controls

### Method:

- Find similar users not in experiment
- Match on behavior pre-experiment
- Compare long-term trajectories

### Use case:

- When holdout not feasible
- Geographic experiments
- Policy changes

## Key Metrics to Track



# Ethical Experimentation: Do No Harm

## Core Principles

### 1. Informed Consent

- Users should know they're in experiments
- Privacy policy disclosure
- Opt-out mechanisms

### 2. Minimize Harm

- Don't test things that could hurt users
- Set guardrails (revenue, UX, safety)
- Monitor closely for negative effects
- Have kill switch ready

### 3. Fairness

- Don't systematically disadvantage groups
- Check for disparate impact
- Ensure equal treatment opportunity

### 4. Transparency

- Document methodology
- Share learnings internally
- Be honest about failures

## What NOT to Test

### Unethical Experiments:

- Intentionally degrade user experience
- Manipulate emotions without consent
- Test discriminatory policies
- Hide critical information
- Exploit vulnerable populations

### Famous Missteps:

- Facebook emotion study (2014): Manipulated newsfeeds to study emotional contagion without consent
- OkCupid compatibility test: Lied about match scores to see behavioral effects

## Best Practices

- Ethics review for sensitive experiments
- Independent oversight for high-risk tests
- User research panel for feedback
- Regular audits of experiment portfolio
- Clear escalation paths for concerns

# Experiment Velocity: The Compounding Effect

## Why Velocity Matters

### Scenario 1: Slow Team

- 10 experiments per year
- 30% win rate
- 3 wins per year
- Average lift per win: 5%
- Annual improvement: 15% (additive)

### Scenario 2: Fast Team

- 100 experiments per year
- 30% win rate (same)
- 30 wins per year
- Average lift per win: 1% (smaller but more)
- Annual improvement: 35% (compounding)

## The Math

$$(1.01)^{30} = 1.35 \text{ (35\% improvement)}$$

Many small wins compound faster than few large wins

**Google's approach:** Would rather ship 100 experiments with 1% lifts than 10 with 10% lifts

## Bottlenecks to Address

### 1. Slow Implementation

- Solution: Feature flags, modular code
- Goal: Idea → live experiment in 2 days

### 2. Long Experiment Duration

- Solution: Higher traffic, sequential testing
- Goal: Results in 1-2 weeks max

### 3. Slow Analysis

- Solution: Automated pipelines, dashboards
- Goal: Results available real-time

### 4. Decision Paralysis

- Solution: Clear criteria, empowered teams
- Goal: Decision within 24 hours of results

## Target Metrics

- Experiments per engineer per quarter: 5-10
- Average experiment duration: 1-2 weeks
- Time to decision: 1-3 days

# Tools & Infrastructure for Experimentation at Scale

## Open Source

### GrowthBook

- Feature flags + A/B testing
- Bayesian statistics
- SQL-based metrics
- Self-hosted or cloud

### Unleash

- Feature toggle management
- Gradual rollouts
- Real-time updates
- Lightweight

### Statsig (Free tier)

- Feature gates
- Dynamic configs
- Autotune
- Built-in metrics

### DIY Stack

- Feature flags: Flagsmith, ConfigCat

Machine Learning for Smarter Innovation

## Enterprise

### Optimizely

- Market leader
- Full-stack experimentation
- Visual editor
- Expensive (\$50K+/year)

### LaunchDarkly

- Feature management focus
- Targeting rules
- Fast flag updates
- Popular with DevOps

### VWO

- Web optimization
- Heatmaps, session replay
- A/B + multivariate
- Marketing-friendly

### Firebase A/B Testing

- Mobile-first

A/B Testing & Iterative Improvement

## Big Tech Internal

### Google's Dapper

- 1000s of concurrent experiments
- Automated metrics computation
- Novelty detection
- Not public

### Facebook's Gatekeeper

- Feature gating
- Dynamic allocation
- Layered experiments
- Internal only

## Build vs Buy

### Buy (SaaS) if:

- Small team (< 10 engineers)
- Need to move fast
- Standard use cases


### Build (Custom) if:

- Large scale (1M+ users)

BSc-Level Course

# When to Use A/B Testing: Judgment Criteria

## When to Use A/B Testing: Decision Framework



### Additional Decision Factors

- SKIP A/B Testing When:**
  - Obvious bug fix or broken feature - Ship directly
  - One-way door decision (irreversible) - More validation needed (user research + prototype)
  - Qualitative question ("Why do users ?") - User interviews, usability testing
  - High risk to users - Canary deployment + monitoring instead
  - Need immediate action - Ship with monitoring, iterate fast
  - Regulatory/legal constraints - Compliance review first

- USE A/B Testing When:**
  - Measurable outcome exists (CTR, conversion, revenue)
  - Several options exist (can compare)
  - Unclear which option is better (need data)
  - Stakeholders need objective evidence
  - Production environment required for realistic test
  - Can afford 1-2 week experiment duration

## Key Principles

### 1. Culture

- Experiments are default, not exception
- Psychological safety for “failed” tests
- Data beats opinions (no HiPPO)
- Celebrate learning, not just wins

### 2. Metrics

- North Star: Long-term value proxy
- Guardrails: Prevent negative side effects
- Leading indicators, not lagging

### 3. Communication

- Tailor to audience (exec/PM/eng/DS)
- Always include CI + effect size
- Plain English, no jargon
- Visual ↗ numbers

### 4. Decision-Making

- Clear ship/iterate/kill criteria
- Fast decisions (within 24 hours)

## Strategic Insights

### 5. Long-Term Thinking

- Holdout groups for 6-12 month measurement
- Cohort analysis for retention
- Don’t optimize away future value

### 6. Ethics

- Informed consent
- Minimize harm
- Fairness across groups
- Transparency

### 7. Velocity

- 50-100 experiments per year (target)
- Compound small wins
- Reduce bottlenecks
- Automate everything possible

### 8. Infrastructure

- Feature flags in production
- Real-time dashboards

# Workshop: E-Commerce Recommendation Engine Comparison

## Your Challenge

Design and analyze an A/B test comparing 3 recommendation algorithms for an e-commerce site with 100,000 daily users.

## Why This Matters:

- Real-world ML deployment decision
- Multi-model comparison
- Statistical rigor + business alignment
- Portfolio project for interviews

## Success Criteria:

- Correct sample size calculation
- Proper statistical tests applied
- Guardrail metrics checked
- Clear recommendation with rationale
- Rollout plan with risk mitigation

## What You'll Do

- ① Design experiment (hypothesis, metrics, sample size)
- ② Conduct power analysis (duration calculation)
- ③ Run simulation (3 models, 30K users)
- ④ Perform statistical analysis (t-tests, Bayesian)
- ⑤ Check guardrails (revenue, latency)
- ⑥ Make deployment decision (ship/iterate/kill)
- ⑦ Create rollout plan (1% → 100%)

Time: 60 minutes

Deliverable: Jupyter notebook

Format: Individual or pairs

Tools: Python (scipy, PyMC3)

Validation establishes baseline quality while iteration drives improvement - measurement cycles enable optimization

# Business Context & Baseline Performance

## Current State

### E-Commerce Site:

- 100,000 daily active users
- 1M product catalog
- Average order value: \$50
- Revenue: \$2M per day

### Existing Recommendation System:

- Algorithm: Popularity-based
- CTR (click-through rate): 5.0%
- Conversion rate: 2.0%
- Revenue per user: \$20
- Latency: p95 = 150ms

### Business Goal:

- Increase engagement (CTR)
- Maintain or improve conversion
- Don't degrade revenue
- Keep latency < 200ms

## Candidate Algorithms

### Model A: Collaborative Filtering (CF)

- Offline CTR estimate: 6.2%
- Conversion: 2.1% (expected)
- Latency: 180ms
- Complexity: Medium

### Model B: Content-Based (CB)

- Offline CTR estimate: 5.8%
- Conversion: 2.0% (expected)
- Latency: 160ms
- Complexity: Low

### Model C: Hybrid (CF + CB)

- Offline CTR estimate: 6.5%
- Conversion: 2.2% (expected)
- Latency: 190ms
- Complexity: High

## Key Question

Which model should we deploy? Or should we iterate further?



# Task 1: Experiment Design (10 minutes)

## Hypothesis Formulation

Task: Write 3 hypotheses (one per model)

### Example (Model A - CF):

"Switching from popularity-based to collaborative filtering will increase CTR by at least 1 percentage point (from 5% to 6%) over a 2-week test with 100,000 users, without decreasing revenue per user by more than 2%."

## Primary Metric

Choose one:

- Click-through rate (CTR)
- Conversion rate
- Revenue per user
- Engagement time

### Recommendation: CTR

- Most sensitive (changes fastest)
- Leading indicator for conversion/revenue
- Easier to detect statistically

## Guardrail Metrics

Task: Define guardrails

### Recommended:

- Revenue per user: Must not drop > 2%
- Latency p95: Must stay < 200ms
- Error rate: Must stay < 0.1%
- Conversion rate: Must not drop > 0.2pp

## Experimental Design

### Traffic Allocation:

- Control (Popularity): 25%
- Treatment A (CF): 25%
- Treatment B (CB): 25%
- Treatment C (Hybrid): 25%

### Randomization:

- User-level (consistent experience)
- Stratified by device (iOS/Android)
- Hash-based assignment

Duration: TBD (power analysis next)

## Task 2: Power Analysis (10 minutes)

### Sample Size Calculation

#### Parameters:

- Baseline CTR: 5% ( $p_1$ )
- Target CTR: 6% ( $p_2$ )
- MDE (Minimum Detectable Effect): 1pp
- Significance level ( $\alpha$ ): 0.05
- Statistical power ( $1 - \beta$ ): 0.80

#### Formula (proportions test):

$$n = \frac{(z_{\alpha/2} + z_{\beta})^2 \cdot [p_1(1 - p_1) + p_2(1 - p_2)]}{(p_2 - p_1)^2}$$

#### Calculation:

- $z_{0.025} = 1.96$ ,  $z_{0.20} = 0.84$
- $n \approx 9,800$  per group
- 4 groups  $\rightarrow 39,200$  total users

#### Duration:

- Daily users: 100,000
- Required: 39,200

### Python Implementation

#### Code structure (see notebook):

- Import statsmodels.stats.power
- Import numpy for calculations
- Define baseline CTR: 0.05
- Define target CTR: 0.06
- Set alpha: 0.05, power: 0.80
- Calculate effect size (Cohen's h)
- Use zt\_ind\_solve\_power function
- Compute sample size per group

### Expected Output

#### Sample size calculations:

- Sample size per group: 9,800
- Total sample needed: 39,200
- Days needed: 0.39
- Recommendation: Run for 14 days to capture weekly patterns

**Key insight:** Always run longer than statistical minimum to avoid novelty effects and capture weekly seasonality

# Task 3: Simulation (15 minutes)

## Generate Synthetic Data

Task: Simulate user interactions

### Setup:

- 30,000 users (more than minimum for confidence)
- 7,500 per group
- Bernoulli trials (click = 1, no click = 0)

### True CTRs (ground truth):

- Control: 5.0%
- Treatment A (CF): 6.2%
- Treatment B (CB): 5.8%
- Treatment C (Hybrid): 6.5%

## Python Simulation Steps

### Implementation (see notebook):

- Import numpy, set random seed
- Set n\_per\_group = 7500
- Generate control clicks: `binomial(1, 0.050, n)`
- Generate treatment\_a: `binomial(1, 0.062, n)`
- Generate treatment\_b: `binomial(1, 0.058, n)`

## Expected Output

### Observed CTRs:

- Control: 0.051
- Treatment A: 0.061
- Treatment B: 0.058
- Treatment C: 0.066

### Observations:

- Slight variation from true CTR (sampling noise)
- Treatment C highest (6.6%)
- Treatment A second (6.1%)
- Treatment B marginal (5.8%)

## Add Guardrail Metrics

### Simulate additional metrics:

- Revenue per user: normal distribution
- Control mean: \$20, std: \$5
- Treatment A: \$20.50, std: \$5
- Treatment B: \$20, std: \$5
- Treatment C: \$21, std: \$5

# Task 4: Statistical Analysis (15 minutes)

## Classical Z-Test

Task: Compare each treatment vs control

### Implementation steps:

- Import proportions\_ztest from statsmodels
- For each treatment (A, B, C):
  - Create count array: [treatment.sum(), control.sum()]
  - Create nobs array: [len(treatment), len(control)]
  - Run proportions\_ztest(count, nobs)
  - Extract z-statistic and p-value
  - Compare p-value to alpha = 0.05
- Print results with interpretation

### Expected Results:

- A vs Control:  $p < 0.001$  (significant)
- B vs Control:  $p \approx 0.08$  (not significant)
- C vs Control:  $p < 0.001$  (significant)

## Bayesian Analysis

Task: Calculate  $P(\text{Treatment} > \text{Control})$

### PyMC3 implementation:

- Import pymc3, create model context
- Define Beta(1,1) priors for p\_control
- Define Beta(1,1) priors for p\_treatment\_c
- Add Binomial likelihood for control
- Add Binomial likelihood for treatment\_c
- Sample posterior with 2000 draws
- Calculate probability: treatment > control
- Compute mean of boolean comparison

### Expected Result:

- $P(C > \text{Control}) = 0.998$
- Interpretation: 99.8% confidence C is better

## Convergence

Both frequentist (z-test) and Bayesian methods agree: Treatment C is the clear winner

Methodological convergence strengthens conclusions - independent analytical approaches yielding consistent results increase confidence

# Task 5: Guardrail Check & Decision (10 minutes)

## Guardrail Metrics Check

Task: Verify all guardrails passed

### Treatment C (Hybrid) - Winner Candidate:

#### Check 1: CTR Lift

- `ctr_control = control.mean()`
- `ctr_c = treatment_c.mean()`
- `lift = (ctr_c / ctr_control - 1) * 100`
- Result: 29% lift

#### Check 2: Revenue

- `rev_control = revenue_control.mean()`
- `rev_c = revenue_c.mean()`
- `rev_change = (rev_c / rev_control - 1) * 100`
- Threshold: Must be  $> -2\%$
- Result: +5% (PASSED)

#### Check 3: Latency

- `latency_c = 190ms`
- Threshold:  $< 200ms$
- Result: PASSED

## Decision Matrix

| Model      | CTR  | Rev & Lat | Sig? | Decision |
|------------|------|-----------|------|----------|
| Control    | 5.0% | Baseline  | -    | Baseline |
| CF (A)     | 6.1% | Pass      | Yes  | Consider |
| CB (B)     | 5.8% | Pass      | No   | Kill     |
| Hybrid (C) | 6.6% | Pass      | Yes  | SHIP     |

## Final Recommendation

Ship Treatment C (Hybrid)

#### Rationale:

- Highest CTR: 6.6% (29% lift)
- Statistically significant ( $p < 0.001$ )
- Bayesian: 99.8% confidence
- All guardrails passed:
  - Revenue: +5% ( $> -2\%$  threshold)
  - Latency: 190ms ( $< 200ms$ )
  - Error rate: 0.05% ( $< 0.1\%$ )
- Expected annual value: \$3.6M additional revenue

**Alternative:** Could also ship CF (A) as it also wins, but Hybrid is stronger

# Task 6: Rollout Plan (5 minutes)

## Gradual Rollout Strategy

### Phase 1: Canary (Days 1-2)

- Deploy to 1% of users
- Monitor error rate, latency, revenue
- If stable after 24 hours → proceed
- Kill switch ready

### Phase 2: Expansion (Days 3-5)

- 1% → 5% → 25%
- Each step: 24-hour monitoring
- Check guardrails at each step
- If any guardrail violated → rollback

### Phase 3: Majority (Days 6-10)

- 25% → 50% → 100%
- Maintain 10% holdout for long-term measurement
- Monitor cohort retention over 30 days

### Phase 4: Holdout Analysis (Days 11-40)

- Compare 90% treatment vs 10% control
- Measure 30-day retention, LTV

## Monitoring Dashboard

### Real-Time Metrics:

- CTR (hourly)
- Revenue per user (daily)
- Latency p50/p95/p99 (5-min intervals)
- Error rate (1-min intervals)
- Traffic distribution (control vs treatment)

## Rollback Criteria

### Immediate rollback if:

- Error rate > 0.5% ( $5 \times$  baseline)
- Latency p95 > 250ms (sustained 10 min)
- Revenue per user drops > 10%
- User complaints spike > 50/hour

### Investigate if:

- CTR plateaus or decreases
- Revenue flat despite CTR increase
- Latency creeps above 200ms

## Success Criteria

# Workshop Takeaways & Course Conclusion

## Week 10 Key Lessons

### 1. Iteration is competitive advantage

- Winners ship 100+ experiments/year
- Losers deploy once and hope
- Velocity × win rate = innovation speed

### 2. Rigorous A/B testing prevents disasters

- Randomization enables causal inference
- Statistical significance ≠ practical significance
- Always include guardrails

### 3. Bayesian methods accelerate learning

- $P(B > A)$  more intuitive than p-values
- Earlier stopping possible
- Incorporates prior knowledge

### 4. Culture trumps tools

- Psychological safety for “failed” experiments
- Data beats opinions
- Fast decisions, slow reversions

## Full Course Journey

### 10 Weeks, Complete ML Innovation Loop:

#### Empathize (Weeks 1-3):

- Clustering for user segmentation
- NLP for emotional context

#### Define (Week 4):

- Classification for problem framing

#### Ideate (Week 5):

- Topic modeling for idea generation

#### Prototype (Week 6):

- Generative AI for rapid prototyping

#### Test (Weeks 7-9):


- Responsible AI (ethics)
- Structured outputs (reliability)
- Multi-metric validation (rigor)

#### Iterate (Week 10):

- A/B testing for continuous improvement

# The Continuous Improvement Loop: Never Stop Iterating

The Iteration Cycle: Never Stop Learning



## The 6-Stage Cycle

### 1. Observe

- Monitor metrics
- Listen to users
- Watch competitors

### 2. Hypothesize

- Form testable predictions
- Identify risks

### 3. Design

- Plan experiment
- Calculate sample size
- Define guardrails

### 4. Implement

- Deploy test infrastructure
- Monitor in real-time

### 5. Analyze

- Run statistical tests
- Check guardrails

# Next Steps: Your A/B Testing Journey

## Immediate Actions

### 1. Complete the Workshop (Today)

- Open the Jupyter notebook
- Run all 6 tasks end-to-end
- Document your decision rationale
- Share results with peers

### 2. Practice on Your Project (This Week)

- Identify ML model to test
- Design A/B test with guardrails
- Calculate required sample size
- Implement monitoring dashboard

### 3. Build Experimentation Culture (This Month)

- Run 1 experiment per week minimum
- Document learnings systematically
- Share results with stakeholders
- Celebrate “failed” experiments

## Long-Term Mastery

### 4. Advanced Topics to Explore:

- Multi-armed bandits (adaptive allocation)
- Sequential testing (early stopping)
- Network effects & interference
- Long-term holdout analysis
- Variance reduction techniques
- Heterogeneous treatment effects

## Resources

- Book: “Trustworthy Online Experiments” (Kohavi et al.)
- Tool: GrowthBook (open-source A/B platform)
- Course: Stanford CS329S (ML Systems Design)
- Community: Experiment Results Forum

## Your Competitive Edge

You now know how to iterate 10× faster than peers who lack A/B testing skills. Use this advantage wisely.

Sustained experimentation drives innovation capacity - systematic testing infrastructure transforms how organizations evolve